

Towards A General Aggregation Framework in Chapel

Oliver Alvarado Rodriguez, Engin Kayraklıoglu, Bartosz Bryg, Mohammad Dindoost, David Bader, and Brad Chamberlain

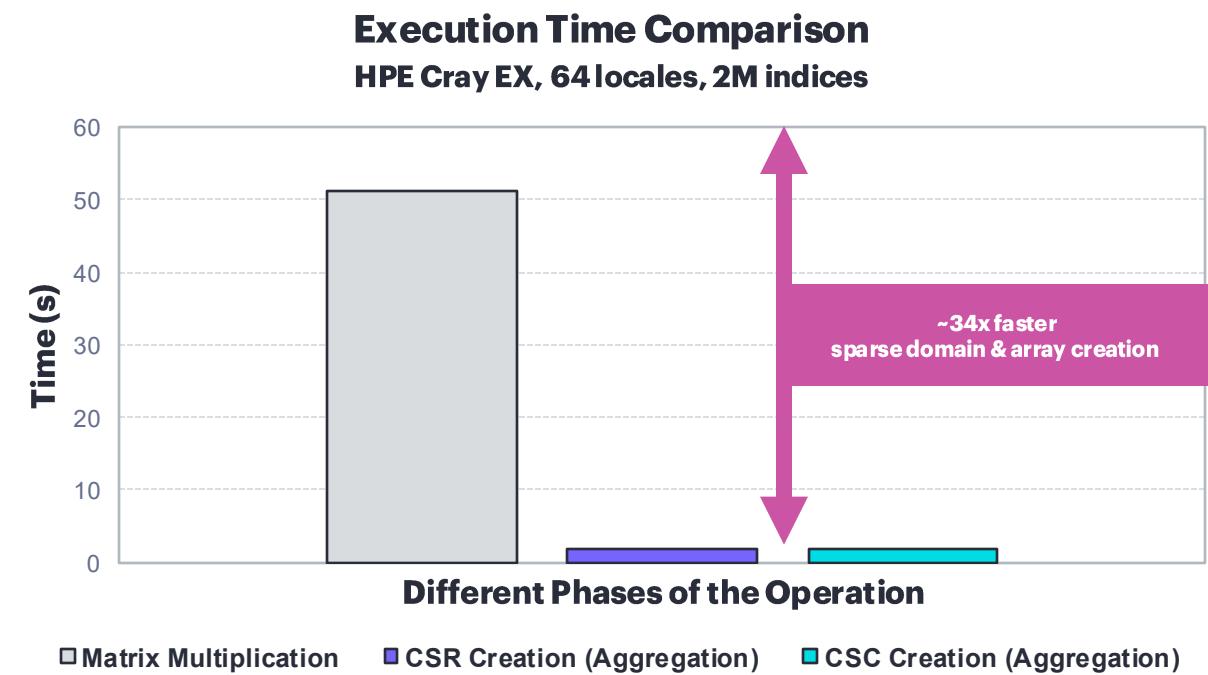
October 10, 2024

Introduction

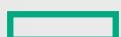
- In distributed-memory, parallel programming, one of the most common bottlenecks is the quantity of communications performed between remote, parallel tasks.
 - Increases in communication are especially noticeable in irregular workloads such as those that use sparse matrices and graphs.
 - For this talk the focus will be on optimizing sparse matrix creation but also taking a minor look at sparse matrix multiplication and RMAT matrix generators for background.
- Chapel has the CopyAggregation module to facilitate the batching of fine-grained communications for array-specific operations.
 - However, there is not a more general framework to support the aggregation of more user-specific operations.
 - We will introduce a framework prototype for more general aggregation.
 - The work for this talk did not only involve the framework prototype.
 - Compressed sparse layouts in Chapel were modified to add parallel safety.
 - The sparse domain buffer functionality, for faster adding of indices into sparse domains, was updated to pass uniqueness and sorted flags.

Motivating Use Case – Sparse Operations in Arkouda

```
1. n = 20                                     # Matrices will be of size (2**n X 2**n).  
2. rows, cols = rmat(n)                      # Create 1D pdarrays of row and column indices.  
3. vals = ak.randint(1, len(rows), len(rows))  # Generate random values.  
4. A = create_sparse_matrix(2**n, rows, cols, vals, "CSR") # Create sparse matrix with CSR layout.  
5. B = create_sparse_matrix(2**n, rows, cols, vals, "CSC") # Create sparse matrix with CSC layout.  
6. C = ak.sparse_matrix_matrix_mult(A, B)       # Do the sparse matrix multiplication.
```



Background



Background – Array Aggregation

```
1  use BlockDist, CopyAggregation;
2
3  const size = 10000;
4  const space = {0..size};
5  const D = space dmapped new blockDist(space);
6  var A, rA: [D] int = D;
7
8  forall (ra, i) in zip(rA, D) with (var agg = new SrcAggregator(int)) do
9      agg.copy(ra, A[size-i]);
```

- Forall loops spawn a source aggregator per-task, therefore there are per-task buffers, making aggregation memory-intensive.
- In this case the destination is ra, which is local, but the source of the data is remote, which is A[size-i].
 - This example shows source aggregation, but the rest of the talk focuses on destination aggregation.
- Whenever the buffer gets full, or all the iterations of the task are finished, then a flush gets issued, that moves all the saved values in the buffer to the memory location they belong in.

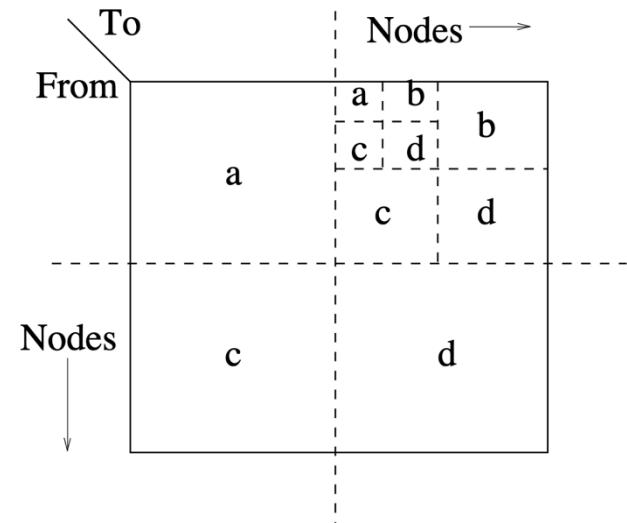
Background – Power-Law Matrices

What? Experiment graphs were **recursive matrix** (R-MAT) random graphs:

- $|V| \leq 2^{SCALE}$
- $|E| \leq (2^{SCALE} * eFACTOR)$

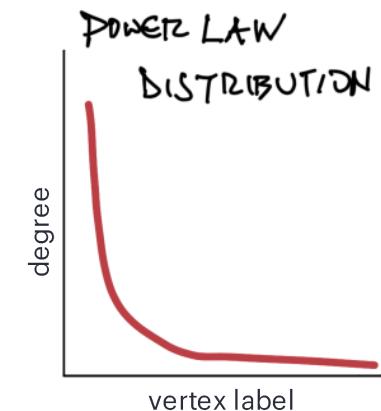
How?

1. Split adjacency matrix into four equal parts.
2. Choose part & subdivide again into four equal parts.
3. Once you reach a 1x1 cell, assign it 0 or 1 to keep that edge in the graph



Why?

- Probabilities $a=0.57$, $b=c=0.19$, and $d=0.05$ give a **Kronecker** graph.
- This type of graph exhibits a power-law vertex degree distribution.
- Real-world graphs exhibit power-law vertex degree distributions.
- This makes them an ideal random graph for benchmarking.



Background – Sparse Matrix Layouts

Compressed Sparse Row (CSR) format

(sorted within each row)

values: `1 2 3 4 5 6 7 8 9 1 2 3`

row starts: `1 2 3 6 8 10 11 12 13`

col inds: `2 8 3 4 6 1 2 1 5 6 6 4`

Space: $2 \cdot nnz + n$

1	n=8
0	1 0 0 0 0 0 0 0
0	0 0 0 0 0 0 0 2
0	0 3 4 0 5 0 0 0
6	7 0 0 0 0 0 0 0
8	0 0 0 9 0 0 0 0
0	0 0 0 0 0 1 0 0
0	0 0 0 0 0 2 0 0
0	0 0 0 3 0 0 0 0

Compressed Sparse Column (CSC) format

(sorted within each column)

values: `6 8 1 7 3 4 3 9 5 1 2 2`

row inds: `4 5 1 4 3 3 8 5 3 6 7 2`

col starts: `1 3 5 6 8 9 12 12 13`

Space: $2 \cdot nnz + n$

Coordinate (COO) format (sorted in row-major order)

values: `1 2 3 4 5 6 7 8 9 1 2 3`

row indices: `1 2 3 3 3 4 4 5 5 6 7 8`

column indices: `2 8 3 4 6 1 2 1 5 6 6 4`

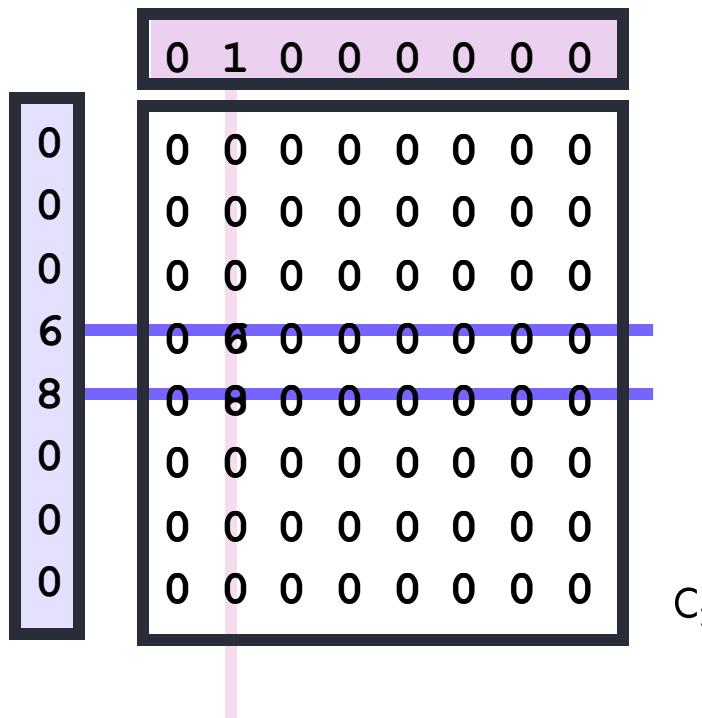
Space: $3 \cdot nnz$ (nnz = number of nonzeroes)



Background – Sparse Matrix Multiplication

LHS matrix (CSC storage)

1	0	1	0	0	0	0	0	0
1	0	0	0	0	0	0	0	2
1	0	0	3	4	0	5	0	0
1	6	7	0	0	0	0	0	0
1	8	0	0	0	9	0	0	0
1	0	0	0	0	0	1	0	0
1	0	0	0	0	0	2	0	0
1	8	0	0	0	3	0	0	0



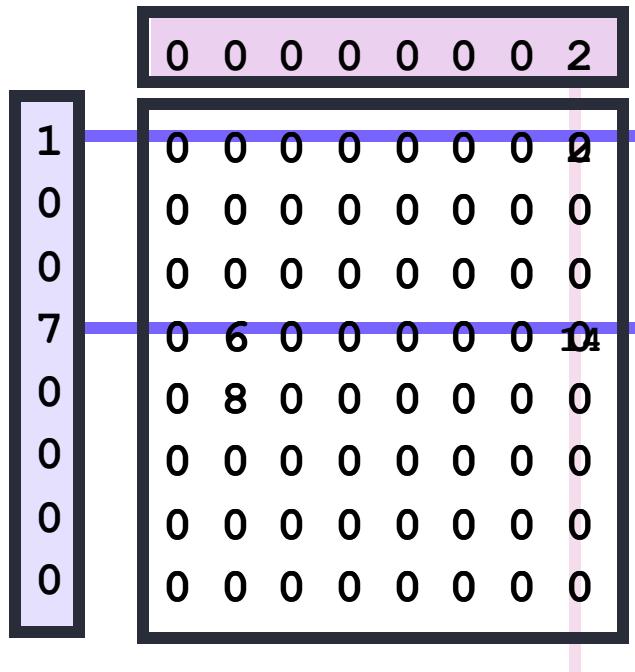
1	0	1	0	0	0	0	0	0
1	0	0	0	0	0	0	0	2
1	0	0	3	4	0	5	0	0
1	6	7	0	0	0	0	0	0
1	8	0	0	0	9	0	0	0
1	0	0	0	0	0	1	0	0
1	0	0	0	0	0	2	0	0
1	8	0	0	0	3	0	0	0

RHS matrix (CSR storage)

Background – Sparse Matrix Multiplication

LHS matrix (CSC storage)

1	0	1	0	0	0	0	0	0
1	0	0	0	0	0	0	0	2
0	0	3	4	0	5	0	0	0
6	7	0	0	0	0	0	0	0
8	0	0	0	9	0	0	0	0
0	0	0	0	0	1	0	0	0
0	0	0	0	0	2	0	0	0
8	0	0	0	3	0	0	0	0



1	0	1	0	0	0	0	0	0
1	0	0	0	0	0	0	0	2
0	0	3	4	0	5	0	0	0
6	7	0	0	0	0	0	0	0
8	0	0	0	9	0	0	0	0
0	0	0	0	0	1	0	0	0
0	0	0	0	0	2	0	0	0
8	0	0	0	3	0	0	0	0

etc.

Proposed Solution



Proposed Solution – High-Level Description

- We take the current CopyAggregation module and modify it to accept two new user-defined records.
 - The source handler to dictate where data is going to be transferred to and how it will be stored within the buffer.
 - The destination handler to perform a flushing operation to move the data from the buffers to their physical memory location.
- Gives more power to the user to let them specify what data structures they want to aggregate into and how should the data be treated.



Proposed Solution – Source Handling

```
1 class SourceHandler {
2     var dVal;
3     var aVal;
4     type elemType = (int, int, int);
5
6     proc init(D, A) {
7         // workaround as ref domain is not implemented
8         this.dVal = D._value;
9
10        // workaround as ref array is not implemented
11        this.aVal = A._value;
12    }
13
14    proc sourceCopy() {
15        return new unmanaged DestinationHandler(dVal, aVal);
16    }
17
18    proc getDestinationLocale(val: elemType) {
19        var (i, j, _) = val;
20        return dVal.parentDom.dist.dsiIndexToLocale((i, j));
21    }
22}
```

Line 4: we have the expected format for the source data, and eventually the destination data of type (i,j,v) where (i,j) is the sparse domain index and v is the data we are passing.

Lines 6-12: the initializer extracts the underlying record for domains and arrays; this is a workaround as current classes in Chapel do not let us take a ref (pointer) of a domain or array.

Lines 14-16: the aggregator has a backend handler that uses this function to instantiate a destination handler whenever it comes time to perform a flush.

Lines 18-21: gets the locale that owns the index (i,j) during a copy step that is used to add a full tuple (i,j,v) to the buffer.

Proposed Solution – Destination Handling

```
23  class DestinationHandler {
24    var domVal;
25    var arrVal;
26
27    proc init(domVal, arrVal) {
28      this.domVal = domVal;
29      this.arrVal = arrVal;
30    }
31
32    inline proc flush(ref rBuffer, const ref remBufferPtr, const ref myBufferIdx) {
33      const (_, locid) = this.domVal.dist.chpl_locToLocIdx(here);
34      var locIdxBuf = this.domVal.locDoms[locid]!.mySparseBlock._value.createIndexBuffer(bufSize);
35      for (dstAddr, srcVal) in rBuffer.localIter(remBufferPtr, myBufferIdx) {
36        assert(dstAddr == nil);
37        var (i,j,_) = srcVal;
38        locIdxBuf.add((i, j));
39      }
40      locIdxBuf.commit();
41      for (dstAddr, srcVal) in rBuffer.localIter(remBufferPtr, myBufferIdx) {
42        assert(dstAddr == nil);
43        var (i,j,v) = srcVal;
44        var (_,loc) = this.domVal.locDoms[locid]!.mySparseBlock._value.find((i,j));
45        this.arrVal.locArr[locid]!.myElems._value.data[loc] = v;
46      }
47    }
48  }
49
50  forall (i,j,v) in zip(rows, cols, vals) with (var agg = new CustomDstAggregator(new shared SourceHandler(SparseDom, SparseArr))) do
51    agg.copy((i,j,v));
```

Lines 32-45: we have the works for the flushing operation, which currently is a lot more scary-looking than intended, I just couldn't help myself and wanted to optimize as much as possible 😊.

Lines 33-34: create an index buffer for faster addition of indices into a sparse domain.

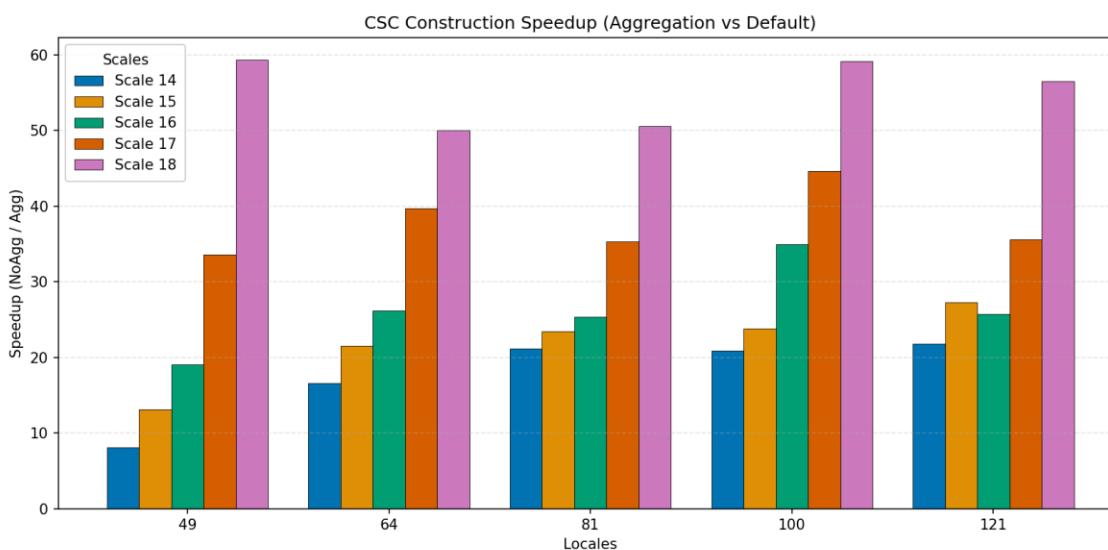
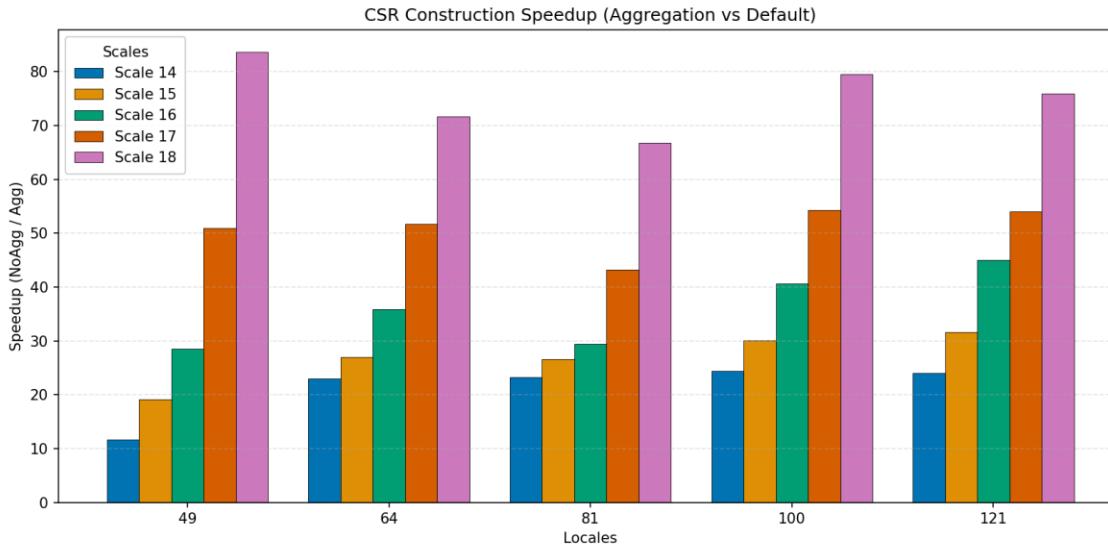
Lines 35-39: add a given index (i,j) into the buffer, once the buffer is full, or hits the commit() in line 40, the buffer gets flushed and those indices get added to the domain. **This is not the same as the remote buffer within the aggregator.**

Lines 41-45: once the indices have been added, we add in the actual data, which requires finding the index for (i,j) in the backend data array.

Lines 53-54: we see the aggregation prototype in action where we create a custom aggregator with the source handler shown in the previous slide.

Benchmarks

Aggregated vs. Non-Aggregated Sparse Matrix Creation Time



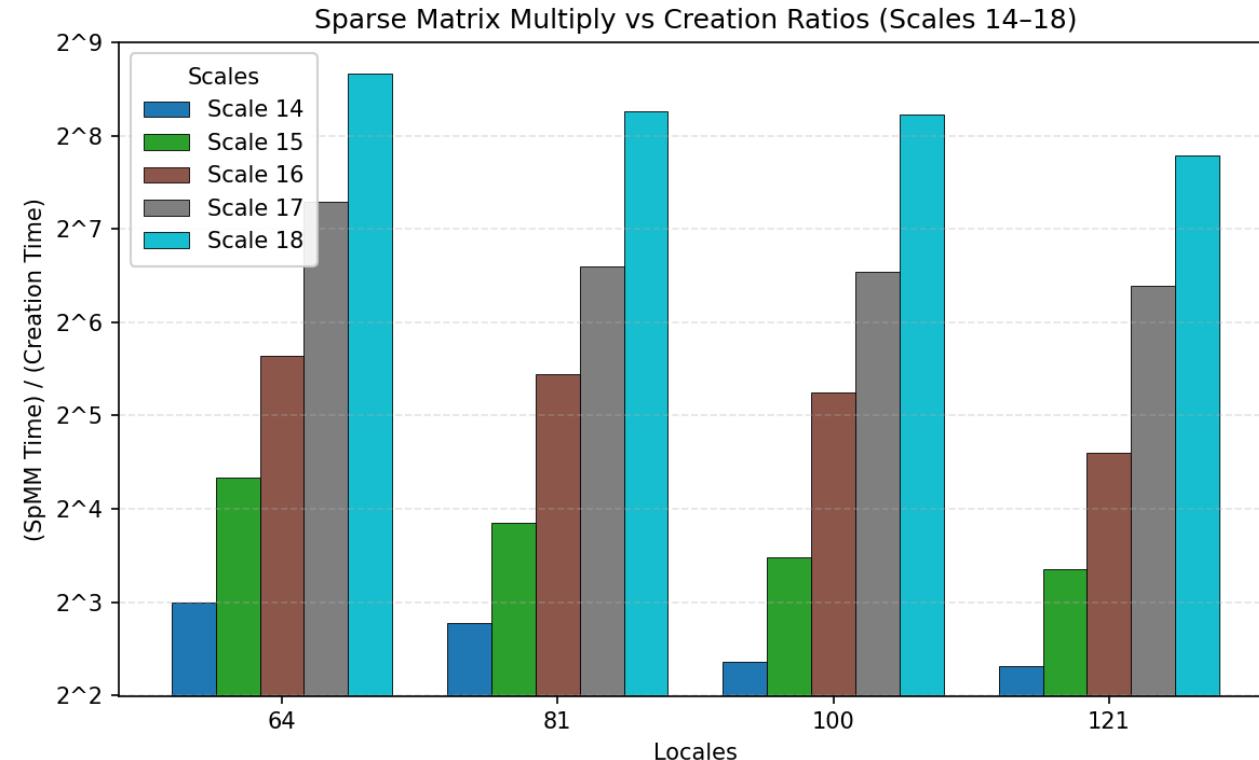
System: HPE Cray EX.

Slingshot-11 interconnect with communication managed through libfabric.

2 AMD EPYC 7763 processors with 256 cores total and 512GB memory per locale.

- We see consistent speed-ups across all scales for the aggregated construction times vs. the current Arkouda construction times.
- The speed-up for CSR is greater because CSC non-aggregated construction was generally slower.
 - I do not have a strong reason or sense of “why” and I do not want to speculate, but it would be a good study to figure out why.
- Generally, as the number of locales increased, the performance got better for aggregated construction whereas the non-aggregated construction slightly degraded in performance as locales increased.
- System-specific issues with compute nodes caused odd results like the ones for CSC construction at 121 locales, where the aggregated code did not perform as well as it should have.
 - Re-running the test with one trial showed much better performance, but I kept the “bad” result to showcase how system-specific issues can affect aggregation.

Sparse Matrix Multiplication vs. Aggregated Creation



System: HPE Cray EX.

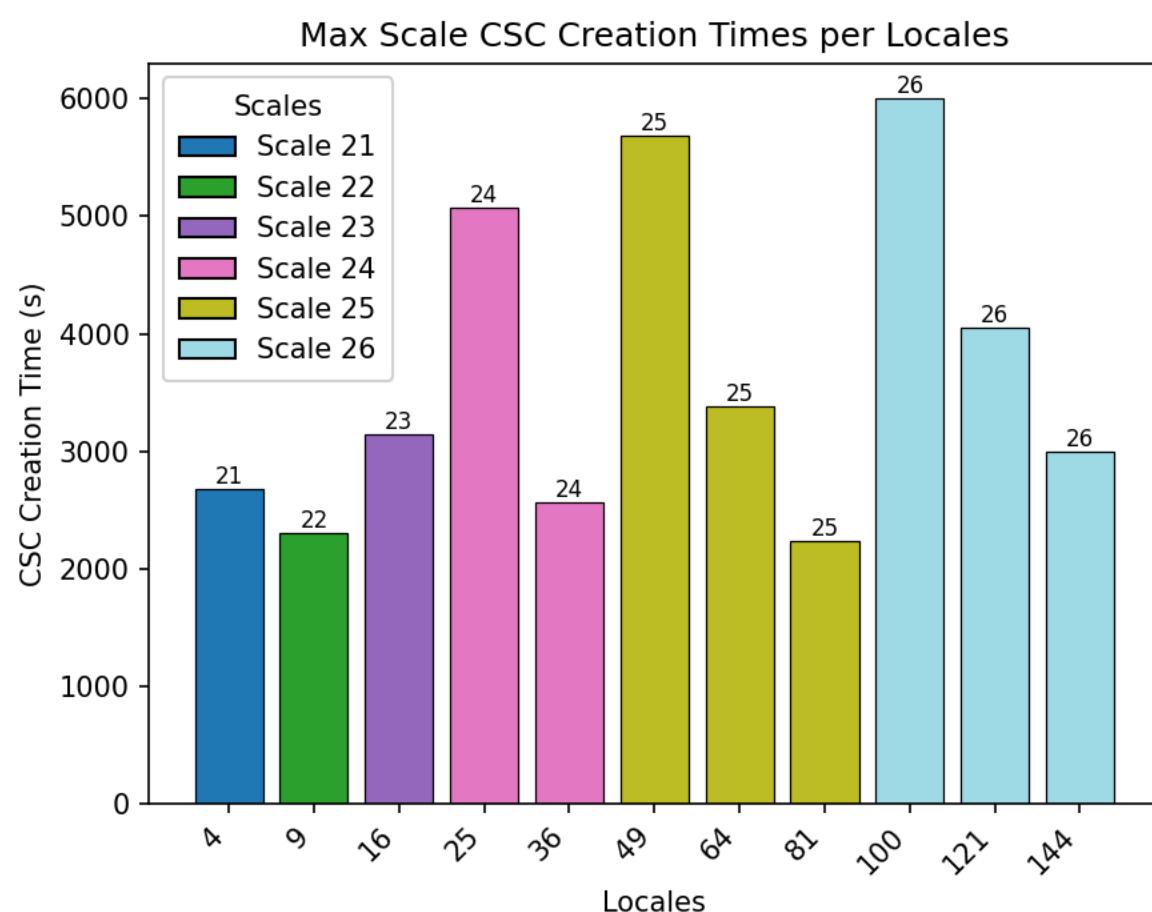
Slingshot-11 interconnect with communication managed through libfabric.

2 AMD EPYC 7763 processors with 256 cores total and 512GB memory per locale.

- Overall, the aggregated sparse matrix creation is now always significantly faster than sparse matrix multiplication.
 - This is good for Arkouda users to quickly load in data and immediately start their analyses without having to spend too much time waiting for data to load.
- I present the y-axis on a log₂ scale. There is no specific reason other than that without the log scale, the bars for Scale 14 looked almost non-existent, even though the creation code was consistently 4x or faster.



Testing the Limits of Aggregated Sparse Matrix Creation



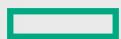
System: HPE Cray EX.

Slingshot-11 interconnect with communication managed through libfabric.

2 AMD EPYC 7763 processors with 256 cores total and 512GB memory per locale.

- Here, we see the maximum scale for the RMAT matrices constructed.
- As a reminder, the matrices occupy an area of $2^{\text{scale}} \times 2^{\text{scale}}$ and have NNZ of about $2^{\text{scale}} * 32$.
 - Give or take repeated edges getting parsed out.
 - If the scale is 20, then that will give a ratio of about $1/32768$ or 0.00003 of NNZ values to the area of the matrix.
- We can see for locales 25-36 the creation time significantly improving for scale 24, and then similarly for locales 49-81 and 100-144.
 - This exhibits strong scalability for the aggregated sparse matrix creation code.
- A graph of scale 26 is considered a “toy” size by the Graph500 benchmark (they run on supercomputers like Fugaku & Frontier).
 - Assuming 64 bits per edge, a graph of this scale takes up about ~20GB in memory.

In Conclusion...



Conclusion

- This work introduces early steps toward a general aggregation framework in Chapel beyond CopyAggregation to allow more general operations such as modifying sparse matrix domains and arrays.
- This prototype gives finer control of distributed communication while keeping Chapel's productive global namespace.
- Aggregated sparse matrix creation showcases benefits for irregular, power-law style workloads.
- Our performance results show aggregation alleviates communication bottlenecks in sparse workloads.
- We validate that explicit aggregation control yields predictable performance gains without sacrificing Chapel's high-level model.
- Going forward this work would benefit from the following.
 - A comparison against state-of-the-art methods like conveyors.
 - A more in-depth look at the framework prototype itself. How well does it support source aggregation? What other workloads can it be applied to?
 - "Hyperparameter" tuning. Aggregation has a ton of toggles like buffer sizes. Is there a Goldilocks-space for buffer sizes to number of aggregations that is most optimal?



Thank You!

Oliver Alvarado Rodriguez

oliver.alvarado-rodriguez@hpe.com