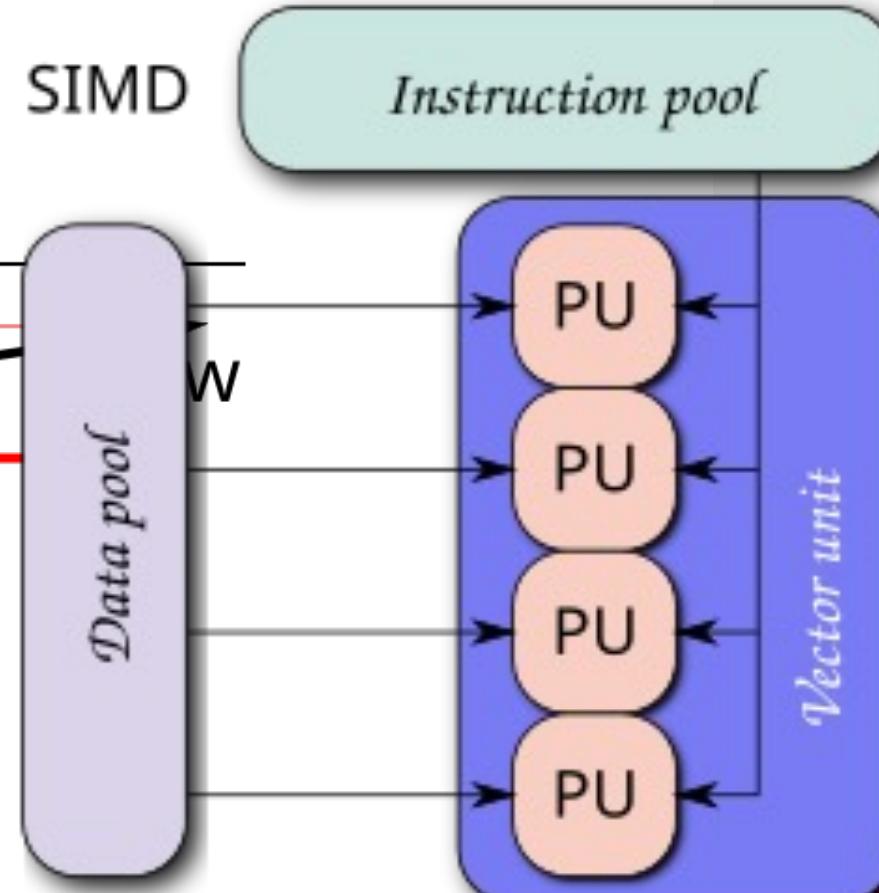
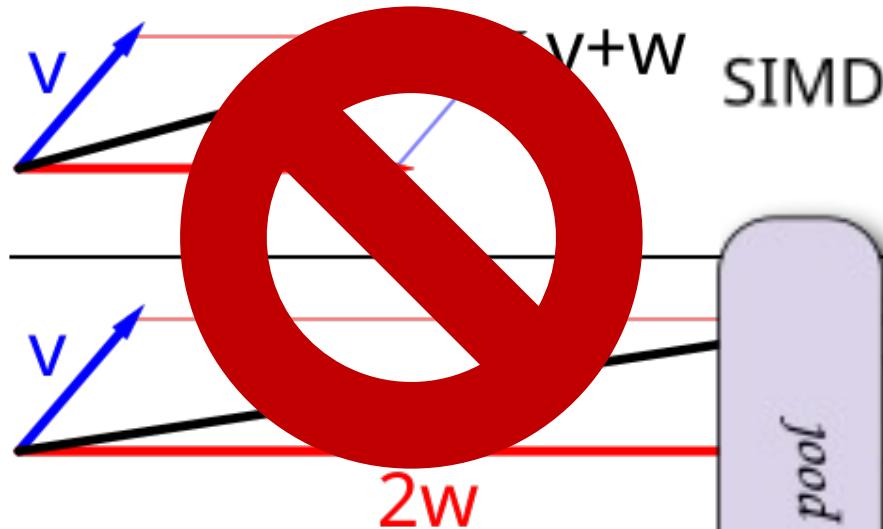


Unlocking Portable and Performant Vector Programming with chpl Vector Library

Jade Abraham

October 10, 2025

Vector?



std::vector
std::vector
Data in header
>

Vector Programming

- Vector programming is using SIMD execution units to process data in parallel within a single thread
 - This is instruction level parallelism
 - Why? More parallelism = more speed!
- For many applications, you don't have to explicitly use it or even know about it
 - Compilers are awesome!
 - Yay free performance!
- So we can end the talk here?



Thank you!



Vector Programming

- Vector programming is making using SIMD execution units to process data in parallel within a single thread
 - This is instruction level parallelism
 - Why? More parallelism = more speed!
- For many applications, you don't have to explicitly use it or even know about it
 - Compilers are awesome!
 - Yay free performance!
- ~~So we can end the talk here?~~

- What happens when the compiler can't do it for us?
 - The compiler may not know its safe or know how to make it safe (floating point error is a pain)
 - Did we write our code in an easy-to-read way for humans, but bad for SIMD?
 - ...et cetera...



Vector Programming - Chapel's missing piece

- Chapel's *multiresolution philosophy*
 - Both high- and low-level features
 - The high-level features are implemented in terms of the low-level features
- This works great for multi-core/distributed parallelism

```
forall a in Arr {  
    // ...something interesting  
}
```

```
coforall l in Arr.targetLocales() do on l do  
    coforall t in 0..#here.maxTaskPar {  
        const mySlice: range(?) = ...  
        for i in mySlice {  
            ref a = Arr[i];  
            // ...something interesting  
        }  
    }  
}
```

- What about instruction-level parallelism?

```
foreach a in Arr {  
    // ...something interesting  
}
```



What's my goal?

- I want to write explicit vector code...
 - ...without calling C/assembly
 - ...that is portable across architectures
 - ...that works orthogonally with existing Chapel features
 - ...that is fast
- I would like my code...
 - ...to not be a maintenance nightmare
 - ...to look nice



Introducing CVL

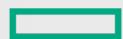


CVL – chpl Vector Library

- Provides a new portable ‘vector’ type which matches a hardware vector register
 - Supports 128-bit and 256-bit vectors with ‘int(?w)’ and ‘real(?w)’
 - Currently supports x86 SSE/AVX and Arm Neon
- Supports many common vector operations
 - Basic math, bit manipulation, and comparisons
 - Memory operations (load/store, limited support for ‘gather’ and load/store masks)
 - Shuffles/permutations/blends
 - Trigonometry (via Sleef - <https://github.com/shibatch/sleef>)
- Integrates seamlessly with Chapel
 - Works with many Chapel container types (arrays, c_ptr, tuples, and bytes)
 - Works with parallel and distributed code
 - Everything is written in pure-ish Chapel
- Open source: <https://github.com/jabraham17/cvl>



Examples, please?



The “Hello World” of HPC/Vector programming

```
proc stream(a: real, x: [?D] real, y: x.type, ref z: x.type) {  
    forall i in D {  
        z[i] = a * x[i] + y[i];  
    }  
}  
  
use CVL;
```

```
proc streamWithCVL(a: real, x: [?D] real, y: x.type, ref z: x.type) {  
    type vec = vector(real, 4); ← Specify the size of the vector,  
    const av = a: vec; ← It must match a hardware type!  
    forall i in D by vec.numElts { ← Create a vector from 'a'  
        const xv = vec.load(x, i); ← Adjust the iteration to be  
        const yv = vec.load(y, i); ← every 4th index  
        const zv = av * xv + yv;  
        zv.store(z, i); ← Load/store the memory  
    }  
}
```



The “Hello World” of HPC/Vector programming

```
use CVL;

proc streamWithCVL(a: real, x: [?D] real, y: x.type, ref z: x.type) {
    type vec = vector(real, 4);
    const av = a: vec;
    forall i in D by vec.numElts {
        const xv = vec.load(x, i);
        const yv = vec.load(y, i);
        const zv = av * xv + yv;
        zv.store(z, i);
    }
}
```

Specify the size of the vector,
It must match a hardware type!

Create a vector from 'a'

Adjust the iteration to be
every 4th index

Load/store the memory

- Explicit vector operations that are distributed and parallel!
- But it is overly verbose, hiding the actual computation
- We can do better



The “Hello World” of HPC/Vector programming

```
use CVL;

proc streamWithCVL(a: real, x: [?D] real, y: x.type, ref z: x.type) {
    type vec = vector(real, 4);

    forall (zv, xv, yv) in zip(vec.vectorsRef(z),
                                vec.vectors(x), vec.vectors(y)) {
        zv = a * xv + yv;
    }
}
```

The scalar is automatically made into a vector

The iterators handle all the load/store logic for us

```
proc stream(a: real, x: [?D] real, y: x.type, ref z: x.type) {
    forall i in D {
        z[i] = a * x[i] + y[i];
    }
}
```

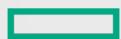


The “Hello World” of HPC/Vector programming

- Is the CVL version faster/better than the plain Chapel version?
 - Default Rectangular arrays: identical performance
 - Block distributed: CVL is ~2x slower
 - Block Cyclic distributed: CVL is A LOT slower
- The gap is likely Chapel specific optimizations that explicit SIMD thwarts
- Just because you can, doesn't mean you should



Something harder?



Kmeans Clustering

```
for cIdx in centroids.D {  
    const cX = centroids.x[cIdx], cY = centroids.y[cIdx];  
    forall pIdx in points.D with (ref points) {  
        const dist = distance(points, pIdx, centroids, cIdx); ← Compute the distance  
        if dist < points.minDist[pIdx] {  
            points.minDist[pIdx] = dist;  
            points.clusterId[pIdx] = cIdx; ← Conditionally update the minimum  
        }  
    }  
}
```

```
for cIdx in centroids.D {  
    const cIdxVec = new VT_IDX(cIdx);  
    const cVecX = new VT(centroids.x[cIdx]), cVecY = new VT(centroids.y[cIdx]);  
    forall pIdx in VT.indices(points.D) with (ref points) {  
        const dist = distance(VT, points, pIdx, cVecX, cVecY); ← Compute the distance  
        const minDist = VT.load(points.minDist, pIdx);  
        const oldClusterId = VT_IDX.load(points.clusterId, pIdx);  
  
        const mask = dist < minDist;  
        var newMinDist = bitSelect(mask, dist, minDist); ← Determine which value to use  
        var newClusterId = bitSelect(mask.transmute(VT_IDX), cIdxVec, oldClusterId);  
  
        newMinDist.store(points.minDist, pIdx);  
        newClusterId.store(points.clusterId, pIdx); ← Always update the minimum  
    }  
}
```

Kmeans Clustering

- Is the CVL version faster/better than the plain Chapel version?
 - At small problem sizes they are the same
 - At big problem sizes CVL beats plain Chapel
- What's the catch?

```
record pointsList {  
    type T;  
    const D: domain(1);  
    var x: [D] T;  
    var y: [D] T;  
    var clusterId: [D] int;  
    var minDist: [D] T;  
}
```

- If I use the wrong data structure
 - The plain Chapel code is slower
 - It is much harder to hand vectorize

	1 million points	10 million points	100 million points
Chapel	0.413s	8.723s	78.106s
Chapel + CVL	0.346s	3.004s	64.306s

```
record pointsList {  
    type T;  
    const D: domain(1);  
    var xy: [D] point(T);  
    var clusterId: [D] int;  
    var minDist: [D] T;  
}  
record point {  
    type T;  
    var x: T;  
    var y: T;  
}
```



How does it work?



A brief dive into the implementation

- The top-level 'vector' type is implemented by multiple layers of type abstractions
 - 'vector(eltType, numElts)' constructs an 'implType(eltType, numElts)'
 - 'implType' is implemented for each architecture/bit-width as a type-only type
- Each 'implType' has a set of operations and behaviors it must conform to
 - If the underlying hardware has a different behavior, shuffle the vector to match (e.g. pairwise adds)
 - Arbitrary shuffles/permutations/blends are not permitted
- At the lowest level, each operation on 'implType' is either
 - directly calling a compiler intrinsic
 - calling a C wrapper around a compiler intrinsic
- 'implType' is a fantastic example of Chapel metaprogramming
 - Compile-time dispatch greatly reduces boilerplate
 - Everything is done at compile-time, all you are left with in the generated code are the vector operations



How does it compare?



Who does vectorization the best?

- Nbody (50,000,000 iterations) from the Computer Language Benchmark Game

	M1 Arm (8 cores)	Intel Xeon E5-2690 v3 (24 cores)	AMD EPYC 7662 (128 cores)
Chapel	1.330s	3.490s	2.731s
Chapel + CVL	1.626s	2.621s	2.434s
Chapel + CVL (fma)	1.511s	2.437s	2.378s
C	2.730s	5.940s	4.150s
C (x86 Intrinsics)	N/A	1.911s	2.648s
Fortran	2.444s	4.025s	3.930s
Rust	1.449s	3.333s	3.268s

- Chapel! (kinda)



Is vector code faster?

- RGB -> Grayscale using integers (problem size scaled per platform)

	M1 Arm (8 cores)	Intel Xeon E5-2690 v3 (24 cores)	AMD EPYC 7662 (128 cores)
Chapel	1.009	6.505	1.524
Chapel + CVL	0.247	0.847	0.349

4x-8x improvements!

- RGB -> Grayscale using floating point (problem size scaled per platform)

	M1 Arm (8 cores)	Intel Xeon E5-2690 v3 (24 cores)	AMD EPYC 7662 (128 cores)
Chapel	1.024	8.760	1.700
Chapel + CVL	0.242	0.845	0.337

4x-10x improvements!

- Yes!



Conclusion

- CVL lets programmers fill a missing gap in Chapel's parallel story
 - Portable, performant, and pretty vector code
- CVL is ready for use!
 - <https://github.com/jabraham17/cvl>
- CVL is not a silver bullet for performance in Chapel, but it is another tool in the toolbox
- What's next?
 - Expanded 'vectorsRef()' support
 - Find a nice ergonomic story for tail loops
 - Leverage the Chapel compiler for more flexible shuffles
 - Even tighter integration with Chapel arrays
 - Close the distributed array performance gap
 - Support 2D arrays without 'reshape()'



Thank you!

