

# **Arkouda Bulletin**

## **A Year of Progress in Exploratory Data Analytics at Scale**

Amanda Potts, Engin Kayraklioglu

ChapelCon '25  
October 9, 2025



# Outline

- Arkouda Overview & Introduction
- Success Stories
- Alignment to NumPy/Pandas
- Random Module
- Multi-Dimensional Data
- Parquet I/O Support
- Sparse Computations
- Outlook/Conclusion



# Introduction

# Arkouda: NumPy for Supercomputers

- A **Pythonic interface to high-performance computing**
  - Open-source framework for **exploratory data analysis at scale**
  - Combines **NumPy-like syntax** with **Chapel's distributed performance**
  - Operates **interactively from Python** — no parallel programming required
  - Handles **billions of elements** across **many nodes**, reproducibly



# Arkouda: Where Python Meets Performance

- **Bridging Productivity and Performance**
  - **Familiar** — mirrors NumPy, pandas, and SciPy semantics
  - **Scalability** — parallel computation over **massive arrays**
  - **Reproducible** — deterministic RNG and shuffle operations
  - **Extensible** — easy to add new Chapel functions via message framework
  - **Accessible** — installable via **Spack; Docker/Kubernetes** support under exploration
  - **Evolving** — ak.numpy, ak.pandas, ak.scipy
- **Bottom Line**
  - **Empowers** researchers to move seamlessly from prototyping in Python to analyzing terabytes interactively, without rewriting code.




# Success Stories

# Telemetry Use Case 1

- How does energy-capping GPUs impact application performance?
- This work is a collaboration between our colleagues at HPE and ORNL
  - Using telemetry data from Frontier

## Experiment details:

- A pandas script has been transliterated to Arkouda
  - Achieved 3.5x better performance on a single node
  - Same script also demonstrates good weak scaling
  - Note that pandas can't be used on multiple nodes at all




# Telemetry Use Case 2

- What is the relationship between environment (e.g. temperature) and node failures?
- Imagine you have a very large server telemetry data, and information on failures, can you find any correlation?

## Experiment details:

- 4TB of data stored in Parquet files
- Operations include:
  - Histogram
  - Mean, max
  - Covariance
- All experiments were run on 64 nodes of HPE Cray EX




# Numpy & Pandas Alignment

# Numpy Alignment

## Strategy


- **Start with NumPy (foundation)** → pandas & SciPy depend on it; covering NumPy unlocks most downstream call paths.
- **Mirror NumPy APIs in ak.numpy** → **Users reuse NumPy muscle** memory (same names/args/semantics).
- **Use NumPy docs as the contract** → Solves the “web-search problem”: answers from NumPy docs apply to Arkouda.
- **Reorganize into ak.numpy, ak.pandas, ak.scipy** → Clear place for each function; easier for contributors to navigate.
- **Rank by real usage (analyzed pandas/SciPy calls)** → High-impact first; all but 2 NumPy funcs used  $\geq 10\times$  in pandas functions are now supported.
- **Next: verify per-function parity** → Audit dtype promotion, broadcasting/axis, NA/Inf, and error behavior for exact NumPy match.



```
>>>
>>> import arkouda.numpy as np
>>>
>>> x = np.arange(6, dtype="uint64").reshape(2, 3)
>>> y = np.array([10, 20, 30], dtype="int64")
>>>
>>> (x + y).dtype      # promotion
dtype('float64')
>>>
>>> (x + y).shape     # broadcasting
(2, 3)
>>>
>>> np.sum(x + y, axis=0)
array([23.0 45.0 67.0])
>>>
```

# Numpy Alignment


pandas & SciPy (what matters in practice)



Takeaway: Heavy hitters are covered; remaining work is a low-volume tail. Gaps by calls → pandas: 0.2%, SciPy: 3.1%.

# Numpy Alignment

pandas & SciPy (what matters in practice)



# Pandas ExtensionArray API (Arkouda)

## Experimental

- **What It Is**

- Arkouda-backed ExtensionArrays so pandas Series/DataFrame columns stay remote (numeric, bool, string, categorical).
- Supports zero-copy construction where possible.

- **Why**

- Enable scalable pandas workflows.
- Avoid rewriting all of pandas.

- **What works today**

- Column creation, indexing, equality, argsort, and common reductions.
- Clean fallback to NumPy dtypes when needed.

- **Caveats**

- Experimental — some pandas paths still call `.to_numpy()`.
- NA semantics and a few reductions incomplete in certain types.

```
>>>
>>> register_extension_dtype(ArkoudaInt64Dtype)
>>> register_extension_dtype(ArkoudaFloat64Dtype)
>>> register_extension_dtype(ArkoudaBoolDtype)
>>>
>>> x = pd.array([1, 2, 3], dtype="int64")
>>> y = pd.array([True, False, True], dtype="bool")
>>> z = pd.array([11, 22, 33], dtype="float64")
>>> x
ArkoudaArray([1 2 3])
>>>
>>> df1 = pd.DataFrame({"x":x, "y": y})
>>> df2 = pd.DataFrame({"x":x, "z": z})
>>>
>>> df3 = df1.merge(df2, on=["x"])
>>> df3
   x      y      z
0  1    True  11.0
1  2   False  22.0
2  3    True  33.0
>>>
>>> type(df["x"].values)
ArkoudaArray
>>>
```

# Pandas ExtensionArray API (Arkouda)

- **Pandas ↔ Arkouda Bridge (Joins/Merges/Groupby)**

- **Key point**

- Arkouda already has **distributed** merges/groupby — need a **pandas bridge** that dispatches to Arkouda (avoid .to\_numpy())

- **Preferred: pandas accessor**

- .ak on Series/DataFrame:
- df.ak.merge(...), df.ak.groupby(...).agg(...)

- **Other integration options**

- **Subclassing** (pd.DataFrame/Series + mixin): natural syntax (df.merge) — but brittle vs pandas internals & upgrades
- **Monkeypatching** (override DataFrame.merge, GroupBy.agg): fastest demo path — but risky, version-fragile; keep opt-in

- **Status**

- **Early experimental** EAs/dtypes exist; bridge layer TBD.


```
>>>
>>> df3 = df1.merge(df2, on=["x"])
>>> df3
   x      y      z
0 1  True  11.0
1 2 False  22.0
2 3  True  33.0
>>>
>>> df4 = df1.ak.merge(df2, on=["x"])
>>> df4
   x      y      z
0 1  True  11.0
1 2 False  22.0
2 3  True  33.0
>>>
```

# Random Module

# Random Module

## What's New

- **ak.random.Generator** via default\_rng(seed)
  - PCG64 backend, independent streams per dtype
- **Supported distributions**
  - integers, uniform, normal / lognormal
  - exponential / standard\_exponential, poisson, standard\_gamma
  - choice, permutation, shuffle
- **Method variants**
  - e.g. standard\_normal(method={"zig","box"})
  - standard\_exponential(method={"zig","inv"})
- **Legacy API (backward-compatible)**
  - ak.random.\* functions (ak.rand, ak.randint, ak.uniform)
  - Now thin wrappers over the new Generator




# Random Module

## Reproducibility

- **Seeded & deterministic**
  - Stable results if locale count is unchanged (most ops)
- **Locale sensitivity**
  - Changing locale count → different draws/permutations
- **Locale-invariant exception**
  - `shuffle(method="Feistel")`: keyed permutation over [0, N)
- **Shuffle methods**
  - Fisher-Yates: simple, single-locale (testing / small data)
  - MergeShuffle: scalable, fully distributed; reproducible only if locale count fixed
  - Feistel: distributed, keyed, reproducible (not cryptographic)
- **Looking ahead**
  - Exploring stateless RNGs (Philox, Threefry) for locale-independent draws and per-element determinism

Arkouda RNG Shuffle Visualization



# Multi-dim Support





# Multi-Dimensional Arrays

## What's New

### Multi-Dimensional Support

- **Creation:** array/zeros/ones/full/\_like accept tuple shapes
- **Shape Ops:** reshape/flatten/squeeze with negative-axis support
- **Broadcasting:** rules aligned; centralized axis validation
- **Elementwise:** abs/cos/clz/isinf
- **Manipulation:** repeat, tile, flip are axis-aware for N-D; concatenate, where
- **Reduction:** sum/prod/min/max/cumsum/cumprod/diff
- **Linear Algebra:** matmul/dot/vecdot (mixed-rank matmuls not supported)
- **Sorting:** argsort/coargsort/sort support axis on numeric arrays

Arkouda Multi-Dimensional Support: Shape • Broadcast • Axis Ops • Reduction



```
>>> A = ak.arange(6*8).reshape(6, 8)
>>> row = ak.arange(8)
```

# Multi-Dimensional Arrays

## Current Gaps/Next Up

| Feature Area                                  | 1-D | N-D     | Notes/Status                  |
|---|-----|---------|-------------------------------|
| Numeric Arrays                                | Yes | Yes     | Core Ops Complete             |
| Strings/Categorical                           | Yes | No      | Not Implemented               |
| Set Operations (intersect1d, union1d, etc...) | Yes | No      | Currently 1-D only            |
| Reductions (min/max etc...)                   | Yes | Partial | mink, argmink pending         |
| Other   | Yes | Partial | median, count_nonzero pending |
| Pandas Integration                            | Yes | No      | No DataFrame/Series support   |

# Parquet I/O Support


# Parquet I/O Support

Previous All-Column Read Implementation

**Logical Table:**

| Column 0 | Column 1 | Column 2 |
|----------|----------|----------|
|          |          |          |
|          |          |          |
|          |          |          |

**Simplified Representation of Parquet File:**



**Previous implementation**


Read a column

Jump back to read the next

**Arkouda Client's Dataframe:**

| Column 0 | Column 1 | Column 2 |
|----------|----------|----------|
|          |          |          |
|          |          |          |
|          |          |          |

**Arkouda Server's Symbols:**



# Parquet I/O Support

New All-Column Read Implementation

**Logical Table:**

| Column 0 | Column 1 | Column 2 |
|----------|----------|----------|
|          |          |          |
|          |          |          |
|          |          |          |

**Simplified Representation of Parquet File:**

|                       |
|-----------------------|
| RowGroup 0 / Column 0 |
| RowGroup 0 / Column 1 |
| RowGroup 0 / Column 2 |
| RowGroup 1 / Column 0 |
| RowGroup 1 / Column 1 |
| RowGroup 1 / Column 2 |
| RowGroup 2 / Column 0 |
| RowGroup 2 / Column 1 |
| RowGroup 2 / Column 2 |

**Current implementation**


Read the file linearly

Populate columns at the same time

**Arkouda Client's Dataframe:**

| Column 0 | Column 1 | Column 2 |
|----------|----------|----------|
|          |          |          |
|          |          |          |
|          |          |          |

**Arkouda Server's Symbols:**



# Parquet I/O Support

Performance Results in Synthetic Benchmarks

## Significantly Improved Read Performance, Especially with Multiple Columns

~400GBs of data, 5 columns, split into 128 files, read by 16 locales:

| before (s) | after (s) | speedup (x) |
|------------|-----------|-------------|
| 16.69      | 9.38      | <b>1.78</b> |

Noticeable improvement with smaller numbers of columns


~30GBs of data, 1000 columns, split into 128 files, read by 16 locales:

| before (s) | after (x) | speedup (x)  |
|------------|-----------|--------------|
| 335.97     | 9.98      | <b>33.67</b> |

Gets more significant as number of columns increase


# Parquet I/O Support

## Performance Results from the Telemetry Use Case 2



More than  
2x improvement  
at scale

Much better  
weak scaling  
behavior



# Sparse Computations

# Improving Arkouda's Sparse Linear Algebra Capabilities

**Challenge:** Can you create a distributed sparse domain & array pair,

- using 3 Arkouda pdarrays for rows, columns, and values,
- where the arrays are not necessarily sorted, nor contain unique data ?

**Potential Answer:** Well, of course! You can add indices to Chapel's sparse domains, just iterate over them and add to the domain using `+=`.

**Challenge:** Can you make it run fast at-scale?

**Likely Answer:** Hmm....

# A Quick Background on Copy Aggregation

- Copying random data into an ordered array is a common operation
  - sometimes called "gather"

```
forall (dst, idx) in zip(DstArr, SrcInds) do  
    d = SrcArr[idx];
```

Results in random remote access

```
forall (d, idx) in zip(DstArr, SrcInds) with (var agg = new DstAggregator(int)) do  
    agg.copy(d, SrcArr[idx]);
```

Random access is aggregated and data moved in bulk

# Improving Arkouda's Sparse Linear Algebra Capabilities

**Challenge:** Can you create a distributed sparse domain & array pair,

- using 3 Arkouda pdarrays for rows, columns, and values,
- where the arrays are not necessarily sorted, nor contain unique data ?

**Potential Answer:** Well, of course! You can add indices to Chapel's sparse domains, just iterate over them and add to the domain using `+=`.

**Challenge:** Can you make it run fast at-scale?

**Likely Answer:** Hmm....

# Improving Arkouda's Sparse Linear Algebra Capabilities

**Challenge:** Can you create a distributed sparse domain & array pair,

- using 3 Arkouda pdarrays for rows, columns, and values,
- where the arrays are not necessarily sorted, nor contain unique data ?

**Potential Answer:** Well, of course! You can add indices to Chapel's sparse domains, just iterate over them and add to the domain using `+=`.

**Challenge:** Can you make it run fast at-scale?

**Likely Answer:** Hmm.... A-ha! I am going to use copy aggregation!

**Challenge:** OK, can you copy the data in aggregate, and populate a sparse matrix during the operation?

# Enter Custom Aggregation


```
forall (i,j,v) in zip(rows,cols,vals) with (
    var agg = new CustomDstAggregator(
        new shared SourceHandler(SparseDom, SparseArr)
    )
) {
    agg.copy( (i,j,v) );
}
```

Custom Aggregator allows  
any operation to be done  
after the data is moved


User-defined type to  
describe the operation

We are still working on  
finishing this effort

# Performance w/o Custom Aggregation



# Performance w/ Custom Aggregation



# Honorable Mentions

## Checkpointing

- Arkouda server's state can now be checkpointed (for the most part, we are still closing gaps)

```
ak.save_checkpoint("cp_name") # arrays stored in server's symbol table are saved on the file system  
ak.load_checkpoint("cp_name") # and they are loaded back
```

- You can also opt-in for automatic checkpointing

```
> ./arkouda_server --checkpointMemPct=0.6 --checkpointIdleTime=300
```

Checkpoint after each operation if  
the used memory is  $\geq 60\%$  of available memory

Checkpoint if the server is idle for 300 seconds

## Python Interoperability

- Enables the user to run any simple Python function on Arkouda's pdarrays

```
arr = ak.array([1,2,3])  
res = ak.apply(arr, lambda x: x+1) # res is now [2, 3, 4]
```

# Conclusion & Outlook

# Outlook

- **What's Next**
  - Complete per-function alignment with NumPy semantics
  - Deepen pandas-style functionality and DataFrame operations
  - Advance benchmarking, diagnostics, and tooling for developers
- **Get Involved**
  - Open-source and community-driven — new contributors welcome!
  - **13 active contributors** over the past year.
  - <https://github.com/Bears-R-Us/arkouda>
  - Help shape Arkouda's next phase through **code, docs, testing**, and **new use cases**.



# Conclusion

- **Arkouda in 2025**
  - Mature, **NumPy-like** framework for distributed analytics
  - Stronger **alignment with NumPy 2.0** and **pandas semantics**
  - Expanded **multi-dimensional** and **Python interop** support
  - Enhanced **random utilities**, **sparse matrices**, and **checkpointing**
  - Faster **parquet I/O**
  - Easier deployment via **Spack** and **Docker**, and **Kubernetes**
- **Key Takeaway**
  - Arkouda brings **Python's productivity** to **HPC scale** — enabling reproducible, data-intensive computing at interactive speed.



# Thank You

## **Arkouda Contributors (Past Year)**

@ajpotts • @drculhane • @1RyanK • @jade-abraham • @vasslitvinov  
@ShreyasKhandekar • @e-kayrakli • @jeremiah-corrado • @jaketrookman  
@stress-tess • @john-hartman • @lydia-duncan • @alvaradoo