

Efficient Multi-GPU Communication with NVSHMEM in Chapel

Sosuke Hosokawa, Kenjiro Taura

The University of Tokyo

Limitation in inter-GPU Communication in Current Chapel

- In current Chapel, a GPU cannot access directly to the memory on another GPU.

```
on Locales[0] do on here.gpus[0] {
    // On GPU A
    var arrOnGpuA: [0..N] real;

    on Locales[1] do on here.gpus[0] {
        // On GPU B
        var arrOnGpuB: [0..N] real;

        // forall converted to a GPU kernel
        forall i in 0..N {
            arrOnGpuA[i] = i * i;
        }
    }
}
```

In current Chapel, this is NOT allowed operation:

- **A GPU kernel can only access the memory on the device it runs on.**

Limitation in inter-GPU Communication in Current Chapel

```
on Locales[0] do on here.gpus[0] {
    // On GPU A
    var arrOnGpuA: [0..N] real;

    on Locales[1] do on here.gpus[0] {
        // On GPU B
        var arrOnGpuB: [0..N] real;

        // Copy from arrOnGpuA to arrOnGpuB
        arrOnGpuB = arrOnGpuA;

        // forall converted to a GPU kernel
        forall i in 0..N do
            arrOnGpuB[i] = i;

        arrOnGpuA = arrOnGpuB;
    }
}
```

The only way to access the memory of another GPU in Chapel is **array-wise copy**.

This operation triggers a data transfer handled by **DMA** (intra-node) or **GASNet via CPU** (inter-node, without GPUDirect RDMA)

Disadvantages of this approach:

- High latency for small data transfer
- Poor inter-node bandwidth due to the lack of GPUDirect support.

NVSHMEM: PGAS-style communication for GPUs

NVSHMEM is a library for efficient PGAS-style inter-GPU communication.

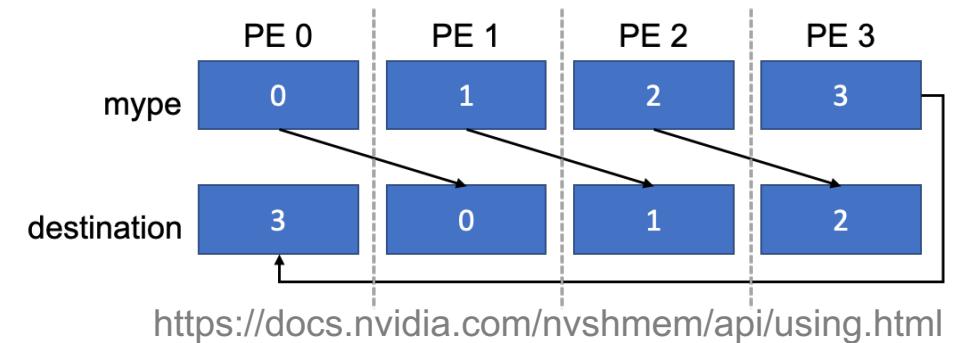
Major features of NVSHMEM are:

- **Device-initiated** remote memory accesses (RMA).
- Reduces RMA overhead by **leveraging some sophisticated technologies**, such as GPUDirect RDMA.
- SHMEM-like **symmetric memory style PGAS API**.

```
__global__ void simple_shifft(int *mem) {
    int mype = nvshmem_my_pe();
    int npes = nvshmem_n_pes();
    int nextpe = (mype + 1) % npes;

    // Memory access to another GPU
    nvshmem_int_p(mem, mype, nextpe);
}
```

Example kernel using NVSHMEM



`nvshmem_int_p` is inter-GPU
memory put operation inside kernel.

Idea: Integrate NVSHMEM into Chapel

- In this work, we are trying to build **more efficient inter-GPU communication library for Chapel** than Chapel's runtime on top of NVSHMEM.
- There are a lot of challenges both in its implementation and designing its API.

```
coforall loc in Locales do on loc {
    var pe = nvshmem_my_pe();
    var next = (pe + 1) % n_gpus;

    // Shift data to the next GPU
    forall i in 0..<1 {
        // Inside a GPU kernel
        nvshmem_int_p(sym_ptr, pe, next);
    }
}
```

What if we can use NVSHMEM with Chapel...?

Challenges in implementing NVSHMEM integration

There are some challenges in implementing NVSHMEM as a Chapel library binding.

The main challenges are:

- Build pipeline of FFI CUDA library for Chapel
- Runtime initialization of NVSHMEM

CUDA device function FFI for Chapel

- Currently, Chapel does not support CUDA's device function FFI.
- But, of course, NVSHMEM's device-side API is implemented as device function of CUDA.
- We overcome this issue by modifying CUDA code build pipeline inside Chapel compiler.
- Specifically, added **GPU code linking step** into compile pipeline.

```
__device__  
int cuda_dev_func(int arg) {  
    ...  
}
```



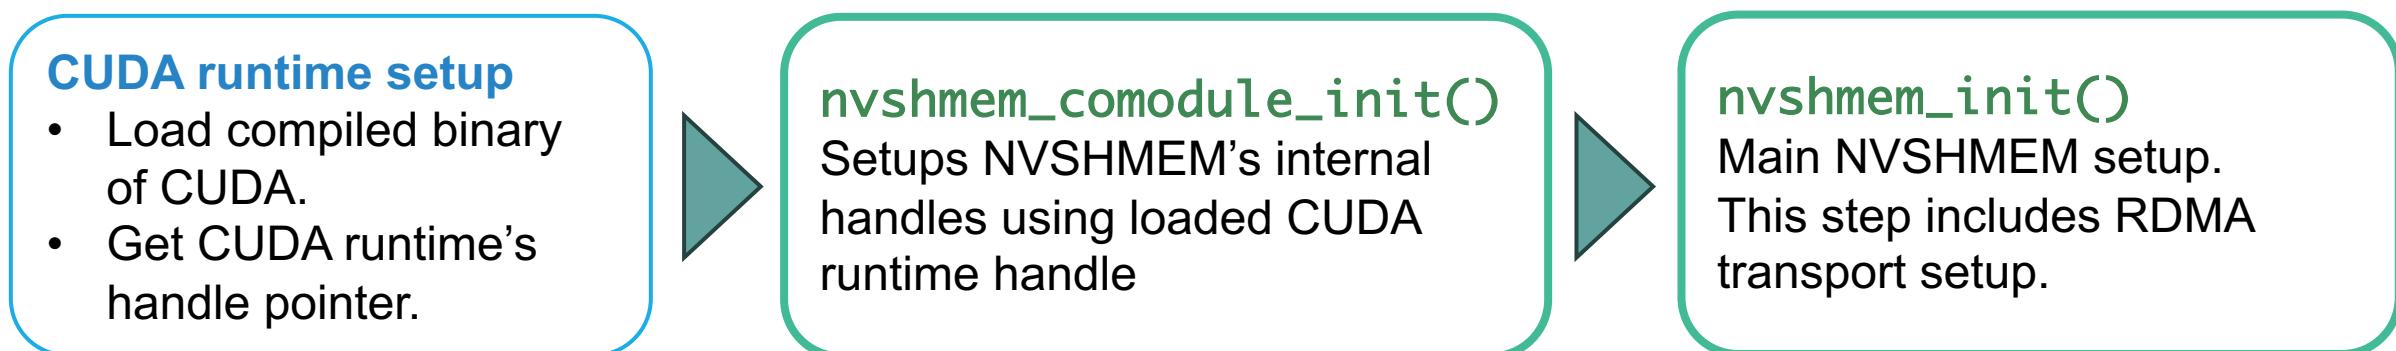
```
extern proc  
dev_func(arg: c_int): c_int
```

Device-callable function
implemented in CUDA.

extern function
in Chapel.

Runtime Initialization of NVSHMEM

- NVSHMEM requires extra step for initialization when it is called from outside of CUDA.
- We modified Chapel's GPU runtime initialization step to use NVSHMEM.



Chapel runtime's GPU initialization step with NVSHMEM



Steps from normal GPU initialization



Steps added for NVSHMEM initialization

API of NVSHMEM bindings

- In our current implementation of NVSHMEM bindings for Chapel, we only exposes **low-level NVSHMEM API**.
- However, this API is **not user-friendly for Chapel programmers** as it directly follows **SHMEM's symmetric-memory style PGAS**.
- We plan to improve this in the future, but for now, we **leave it as future work**.

```
on Locales[0] do on here.gpus[0] {
    // On GPU A
    var arrOnGpuA: [0..N] real;

    on Locales[1] do on here.gpus[0] {
        // On GPU B
        var arrOnGpuB: [0..N] real;

        // forall converted to a GPU kernel
        forall i in 0..N {
            arrOnGpuA[i] = i * i; // Inter-GPU access
        } } }
```

A **GPU-ready, Chapel-style PGAS** model like this is ideal, but it is challenging to implement on top of a symmetric-memory-based PGAS.

Evaluation: Environment

- We evaluated our NVSHMEM bindings on Utokyo's miyabi-g supercomputer.

Miyabi-g node spec:

- NVIDIA Grace CPU (72 cores/CPU, memory 72 GB)
- NVIDIA Hopper H100 GPU x1
- GPU memory: 96 GB
- Interconnect: Infiniband NDR (200 Gbps)

Evaluation: Bandwidth

| Data Size | Chapel Copy Bandwidth | NVSHMEM Put Bandwidth | NVSHMEM Get Bandwidth |
|-----------|-----------------------|-----------------------|-----------------------|
| 4 KiB | 0.0048 GB/s | 0.48 GB/s | 0.48 GB/s |
| 64 KiB | 0.075 GB/s | 5.87 GB/s | 5.87 GB/s |
| 1 MiB | 1.02 GB/s | 20.12 GB/s | 19.95 GB/s |
| 16 MiB | 8.97 GB/s | 24.35 GB/s | 24.31 GB/s |

```
forall i in 0..<1 {
  for iteration in 0..<numIters {
    nvshmem_int64_put(
      dstPtr, srcPtr, dataSize, dstPe);
  }
}

forall i in 0..<1 {
  for iteration in 0..<numIters {
    nvshmem_int64_get(
      dstPtr, srcPtr, dataSize, dstPe);
  }
}
```

- The memory access bandwidth of our runtime is better than original Chapel runtime copy
 - Both in small and large data size.
 - In small data size: ~ 100x
 - In large data size: ~ 3x
- Achieve theoretical bandwidth (25 GB/s = 200 Gbps) in large communication size.

Evaluation: latency

Results:

NVSHMEM on Chapel: 6.744 us
NVSHMEM on CUDA: 6.653 us
CPU OpenSHMEM: 2.140 us

- NVSHMEM on Chapel (ours)
result is as good as NVSHMEM
on CUDA.
- CPU OpenSHMEM ping-pong is
faster.

```
cobegin {
  on Locales[0] {
    var me = nvshmem_my_pe(); var peer = 1 - me;
    on here.gpus[0] {
      forall i in 0..<1 {
        for i in 0..<numIters {
          nvshmem_int_atomic_inc(sym_ptr, peer);
          nvshmem_int_wait_until(sym_ptr, CMP_EQ, i+1);
        }
      }
    }
  }
}

on Locales[1] {
  var me = nvshmem_my_pe(); var peer = 1 - me;
  on here.gpus[0] {
    forall i in 0..<1 {
      for i in 0..<numIters {
        nvshmem_int_wait_until(sym_ptr, CMP_EQ, i+1);
        nvshmem_int_atomic_inc(sym_ptr, peer);
      }
    }
  }
}
```

Conclusion

- In this work, we explored the integration of NVSHMEM with Chapel.
- Our implementation achieves good performance in microbenchmarks, but its API still needs improvement.
- We plan to further enhance our implementation to make Chapel's multi-GPU support more efficient and user-friendly.