

Using Formal Methods to Discover a Bug in the Chapel Compiler

Daniel Fedorin, HPE

The Story

I have a story in three parts.

1. I found it very hard to think through a section of compiler code.
 - Specifically, code that performed scope lookups for variables
2. I used the [Alloy Analyzer](#) to model the assumptions and behavior of the code.
 - I had little background in Alloy, but some background in formal logic
3. This led me to discover a bug in the compiler that I then fixed.
 - Re-creating the bug required some gymnastics that were unlikely to occur in practice.

Background: Chapel's Scope Lookups

The Humble `foo`

Imagine you see the following snippet of Chapel code:

```
foo();
```

Where do you look for `foo`? The answer is quite complicated, and depends strongly on the context of the call.

Moreover, the order of where to look matters: method calls are preferred over global functions, "nearer" functions are preferred over "farther" ones.

The Humble `foo` (example 1)

```
module M1 {
    record R {}

    proc R.foo() { writeln("R.foo"); }
    proc foo() { writeln("M1.foo"); }

    proc R.someMethod() {
        foo(); // which 'foo'?
    }
}
```

- Both `R.foo` and `M1.foo` would be valid candidates.
- We give priority to methods over global functions. So, the compiler would:
 - Search `R` and its scope (`M1`) for methods named `foo`.
 - If that fails, search `M1` for any symbols `foo`.
- We've had to look at `M1` twice! (once for methods, once for non-methods)

The Humble `foo` (example 2)

```
module M1 {
    record R {}

    proc foo() { writeln("R.foo"); }
}

module M2 {
    use M1;

    proc R.someMethod() {
        foo(); // which 'foo'?
    }
}
```

Here, we search the scope of `R` and `M1`, but **only for public symbols**.

How Chapel's Compiler Handles This

We want to:

- Respect the priority order
 - Including preferring methods over non-methods
 - As a result, we search the scopes multiple times
- Avoid any extra work
 - This includes redundant re-searches
 - Example redundant search: looked up "all symbols", then later "all public symbols"

Encoding Search Configuration

```
enum { PUBLIC = 1, NOT_PUBLIC = 2, METHOD_FIELD = 4, NOT_METHOD_FIELD = 8, /* ... */ };
```

1. For each scope, save the flags we've already searched with
2. When searching a scope again, exclude the flags we've already searched with

This was handled by two bitfields: `filter` and `excludeFilter`.

Populating `filter` and `excludeFilter`

```
if (skipPrivateVisibilities) { // depends on context of 'foo()'  
    filter |= IdAndFlags::PUBLIC;  
}  
  
else if (!includeMethods && receiverScopes.empty()) {  
    filter |= IdAndFlags::NOT_METHOD;  
}
```

```
excludeFilter = previousFilter;
```

```
// scary!  
previousFilter = filter & previousFilter;
```

Code notes `previousFilter` is an approximation.

Possible Problems

- `previousFilter` is an approximation.
- However, no case we knew of hit this combination of searches, or any like it.
 - All of our language tests passed.
 - Code seemed to work.
- If only there was a way I could get a computer to check whether such a combination could occur...

Formal Methods

Model Checking

Model checking involves formally describing the behavior of a system, then having a solver check whether other desired properties hold.

- Alloy is an example of a model checker.
- TLA is another famous example.

A Primer on Logic (example)

Model checkers like Alloy are rooted in temporal logic, which builds on first-order logic.

Example statement: "Bob has a son who likes all compilers".

$$\exists x. (\text{Son}(x, \text{Bob}) \wedge \forall y. (\text{Compiler}(y) \Rightarrow \text{Likes}(x, y)))$$

In Alloy:

```
some x { Son[x, Bob] and all y { Compiler[y] implies Likes[x, y] } }
```

A Primer on Temporal Logic

Temporal logic provides additional operators to reason about how properties change over time.

- $\Box p$ (in Alloy: `always p`): A statement that is always true.
 - Example: $\Box(\text{like charges repel})$
- $\Diamond p$ (in Alloy: `eventually p`): A statement that will be true at some point in the future.
 - Example: $\Diamond(\text{the sun is in the sky})$

In Alloy specifically, we can mention the next state of a variable using `'`.

```
// pseudocode:  
//   the next future value of previousFilter will be the intersection of filter  
//   and the current value  
previousFilter' = filter & previousFilter;
```

Modeling Possible Searches

Alloy isn't an imperative language. We can't mutate variables like we do in C++. Instead, we model how each statement changes the state, by relating the "current" state to the "next" state.

```
filter |= IdAndFlags::PUBLIC;
```

```
addBitfieldFlag[filterNow, filterNext, Public]
```

This might remind you of [Hoare Logic](#), where statements like:

$$\{P\} s \{Q\}$$

Read as:

“ If P is true before executing s , then Q will be true after executing s . ”

```
{filter = filterNow} filter |= PUBLIC {filter = filterNext}
```

Modeling Possible Searches

To combine several statements, we make it so that the "next" state of one statement is the "current" state of the next statement.

```
curFilter |= IdAndFlags::PUBLIC;  
curFilter |= IdAndFlags::METHOD_FIELD;
```

```
addBitfieldFlag[filterNow, filterNext1, Public]  
addBitfieldFlag[filterNext1, filterNext2, Method]
```

This is reminiscent of sequencing Hoare triples:

$$\{P\} s_1 \{Q\} s_2 \{R\}$$

Modeling Possible Searches

Finally, if C++ code has conditionals, we need to allow for the possibility of either branch being taken. We do this by using "or" on descriptions of the next state.

```
if (skipPrivateVisibilities) {
    curFilter |= IdAndFlags::PUBLIC;
}

if (!includeMethods && receiverScopes.empty()) {
    curFilter |= IdAndFlags::NOT_METHOD;
}
```

```
addBitfieldFlag[initialState.curFilter, bitfieldMiddle, Public] or
bitfieldEqual[initialState.curFilter, bitfieldMiddle]

// if it's not a receiver, filter to non-methods (could be overridden)
addBitfieldFlagNeg[bitfieldMiddle, filterState.curFilter, Method] or
bitfieldEqual[bitfieldMiddle, filterState.curFilter]
```

Putting this into a predicate, `possibleState`, we encode what searches the compiler can undertake.

Takeaway: We encoded the logic that configures possible searches in Alloy. This instructs the analyzer about possible cases to consider.

Are there any bugs?

Model checkers ensure that all properties we want to hold, do hold. To find a counter example, we ask it to prove the negation of what we want.

```
wontFindNeeded: run {
    all searchState: SearchState {
        eventually some props: Props, fs: FilterState, fsBroken: FilterState {
            // Some search (fs) will cause a transition / modification of the search state...
            configureState[fs]
            updateOrSet[searchState.previousFilter, fs]
            // Such that a later, valid search... (fsBroken)
            configureState[fsBroken]

            // Will allow for a set of properties...
            // ... that are left out of the original search...
            not bitfieldMatchesProperties[searchState.previousFilter, props]
            // ... and out of the current search
            not (bitfieldMatchesProperties[fs.include, props] and not bitfieldMatchesProperties[searchState.previousFilter, props])
            // But would be matched by the broken search...
            bitfieldMatchesProperties[fsBroken.include, props]
            // ... to not be matched by a search with the new state:
            not (bitfieldMatchesProperties[fsBroken.include, props] and not bitfieldOrNotSetMatchesProperties[searchState.previousFilter', props])
        }
    }
}
```

Uh-Oh!

Executing "Run wontFindNeeded"

Solver=sat4j Steps=1..10 Bitwidth=4 MaxSeq=4 SkolemDepth=1 Symmetry=20 Mode=batch
1..3 steps. 6453 vars. 261 primary vars. 14276 clauses. 493ms.

Instance found. Predicate is consistent. 61ms.


The Bug

We need some gymnastics to figure out what variables make this model possible.

Alloy has a nice visualizer, but it has a lot of information.

In the interest of time, I found:

- If the compiler searches a scope first for **PUBLIC** symbols, ...
- ...then for **METHOD_OR_FIELD**, ...
- ...then for any symbols, they will miss things!



The Reproducer

To trigger this sequence of searches, we needed a lot more gymnastics.

```
module TopLevel {
    module XContainerUser {
        public use TopLevel.XContainer;
    }
    module XContainer {
        private var x: int;
        record R {}
        module MethodHaver {
            use TopLevel.XContainerUser;
            use TopLevel.XContainer;
            proc R.foo() {
                var y = x;
            }
        }
    }
}
```

- the scope of `R` is searched with for methods
- The scope of `R`'s parent (`XContainer`) is searched for methods
- The scope of `XContainerUser` is searched for public symbols (via the `use`)
- The scope of `XContainer` is searched with public symbols (via the `public use`)
- The scope of `XContainer` searched for with no filters via the second use; but the stored filter is bad, so the lookup returns early, not finding `x`.

Thank you!



Read more

Extra Slides

Terms

- What are **formal methods**?
 - Techniques rooted in computer science and mathematics to specify and verify systems
- What part of the **Chapel compiler**?
 - The 'Dyno' compiler front-end, particularly its scope lookup phase
 - This piece is used by the production Chapel compiler.

The Humble `foo` (example 1)

```
module M1 {
    record R {
        proc foo() { writeln("R.foo"); }
    }
    proc foo() { writeln("M1.foo"); }

    foo(); // which 'foo'?
}
```

Here, things are pretty straightforward: we look in scope `M1.` `R.foo` is not in it, but `M1.foo` is. We return it.

The Humble `foo` (example 2)

```
module M1 {  
    record R {}  
  
    proc R.foo() { writeln("R.foo"); }  
    proc foo() { writeln("M1.foo"); }  
  
    foo(); // which 'foo'?  
}
```

Here, things are bit trickier. Since we are not inside a method, we know that `foo()` could not be a call to a method. Thus, we rule out `R.foo`, and find `M1.foo` in `M1`.

The Humble `foo` (example 3)

```
module M1 {
    record R {}

    proc R.foo() { writeln("R.foo"); }
    proc foo() { writeln("M1.foo"); }

    proc R.someMethod() {
        foo(); // which 'foo'?
    }
}
```

- Both `R.foo` and `M1.foo` would be valid candidates.
- We give priority to methods over global functions. So, the compiler would:
 - Search `R` and its scope (`M1`) for methods named `foo`.
 - If that fails, search `M1` for any symbols `foo`.
- We've had to look at `M1` twice! (once for methods, once for non-methods)

The Humble `foo` (example 4)

```
module M1 {
    record R {}

    proc foo() { writeln("R.foo"); }
}

module M2 {
    use M1;

    proc R.someMethod() {
        foo(); // which 'foo'?
    }
}
```

Here, we search the scope of `R` and `M1`, but **only for public symbols**.

A Primer on Logic

Model checkers like Alloy are rooted in temporal logic, which builds on first-order logic. This includes:

- Variables (e.g. x, y)
 - These represent any objects in the logical system.
- Predicates (e.g. $P(x), Q(x, y)$)
 - These represent properties of objects, or relationships between objects.
- Logical connectives (e.g. $\wedge, \vee, \neg, \Rightarrow$)
 - These are used to combine predicates into more complex statements.
 - \wedge is "and", \vee is "or", \neg is "not", \Rightarrow is "implies".
- Quantifiers (e.g. \forall, \exists)
 - $\forall x. P(x)$ (in Alloy: `all x { P(x) }`) means "for all x , $P(x)$ is true".
 - $\exists x. P(x)$ (in Alloy: `some x { P(x) }`) means "there exists an x such that $P(x)$ is true".

A Primer on Temporal Logic (example)

Some examples:

$$\square(\text{like charges repel})$$
$$\diamond(\text{the sun is in the sky})$$

In Alloy:

```
// likeChargesRepel and theSunIsInTheSky are predicates defined elsewhere
always likeChargesRepel

// good thing it's not `always`
eventually theSunIsInTheSky
```

Search Configuration in Alloy

Instead of duplicating `METHOD` and `NOT_METHOD`, use two sets of flags (the regular and the "not").

```
enum Flag {Method, MethodOrField, Public}

sig Bitfield {
    , positiveFlags: set Flag
    , negativeFlags: set Flag
}

sig FilterState { , curFilter: Bitfield }
```

Takeaway: We represent the search flags as a `Bitfield`, which encodes `PUBLIC`, `NOT_PUBLIC`, etc.

Constructing Bitfields

Alloy doesn't allow us to define functions that somehow combine bitfields. We might want to write:

```
proc addFlag(b: Bitfield, flag: Flag): Bitfield {
    return Bitfield(b.positiveFlags + flag, b.negativeFlags);
}
```

Instead, we can *relate* two bitfields using a predicate.

“ This bitfield is like that bitfield, but with this flag added. ”

```
pred addBitfieldFlag[b1: Bitfield, b2: Bitfield, flag: Flag] {
    b2.positiveFlags = b1.positiveFlags + flag
    b2.negativeFlags = b1.negativeFlags
}
```

Constructing Bitfields

Alloy doesn't allow us to define functions that somehow combine bitfields. We might want to write:

```
proc addFlag(b: Bitfield, flag: Flag): Bitfield {  
    return Bitfield(b.positiveFlags + flag, b.negativeFlags);  
}
```

Instead, we can *relate* two bitfields using a predicate.

“ This bitfield is exactly like that bitfield. ”

```
pred bitfieldEqual[b1: Bitfield, b2: Bitfield] {  
    b1.positiveFlags = b2.positiveFlags and b1.negativeFlags = b2.negativeFlags  
}
```

Modeling previousFilter

So far, all we've done is encoded what queries the compiler might make about a scope.

We still need to encode how we save the flags we've already searched with.

Model the search state with a "global" (really, unique) variable:

```
/* Initially, no search has happened for a scope, so its 'previousFilter' is not set to anything. */
one sig NotSet {}

one sig SearchState {
    , var previousFilter: Bitfield + NotSet
}
```

Above, `+` is used for union. `previousFilter` can either be a `Bitfield` or `NotSet`.

Types of Formal Methods

- Model checking involves formally describing the behavior of a system, then having a solver check whether other desired properties hold.
 - Alloy is an example of a model checker.
 - TLA is another famous example.
- Theorem proving is a heavier weight approach that involves building a formal proof of correctness.
 - Coq and Isabelle are examples of theorem provers.

Types of Formal Methods

- Model checking involves formally describing the behavior of a system, then having a solver check whether other desired properties hold.
 - Alloy is an example of a model checker.
 - TLA is another famous example.
- Theorem proving is a heavier weight approach that involves building a formal proof of correctness.
 - Coq and Isabelle are examples of theorem provers.

Reason: I was in the middle of developing compiler code. I wanted to sketch the assumptions I was making and see if they held up.

Putting it Together

We now have a model of what our C++ program is doing: it computes some set of filter flags, then runs a search, excluding the previous flags. It then updates the previous flags with the current search. We can encode this as follows:

```
fact step {
    always {
        // Model that a new doLookupInScope could've occurred, with any combination of flags.
        all searchState: SearchState {
            some fs: FilterState {
                // This is a possible combination of lookup flags
                possibleState[fs]

                // If a search has been performed before, take the intersection; otherwise,
                // just insert the current filter flags.
                updateOrSet[searchState.previousFilter, fs]
            }
        }
    }
}
```

Modeling previousFilter

If no previous search has happened, we set `previousFilter` to the current `filter`.

Otherwise, we set `previousFilter` to the intersection of `filter` and `previousFilter`, as mentioned before.

```
if (hasPrevious) {  
    previousFilter = filter & previousFilter;  
} else {  
    previousFilter = filter;  
}
```

```
pred update[toSet: Bitfield + NotSet, setTo: FilterState] {  
    toSet' != NotSet and bitfieldIntersection[toSet, setTo.include, toSet']  
}  
  
pred updateOrSet[toSet: Bitfield + NotSet, setTo: FilterState] {  
    (toSet = NotSet and toSet' = setTo.include) or  
    (toSet != NotSet and update[toSet, setTo])  
}
```