

# HiPerMotif: Novel Parallel Subgraph Isomorphism in Large-Scale Property Graphs

Mohammad Dindoost, Oliver Alvarado Rodriguez, Bartosz Bryg, Ioannis Koutis, and David A. Bader

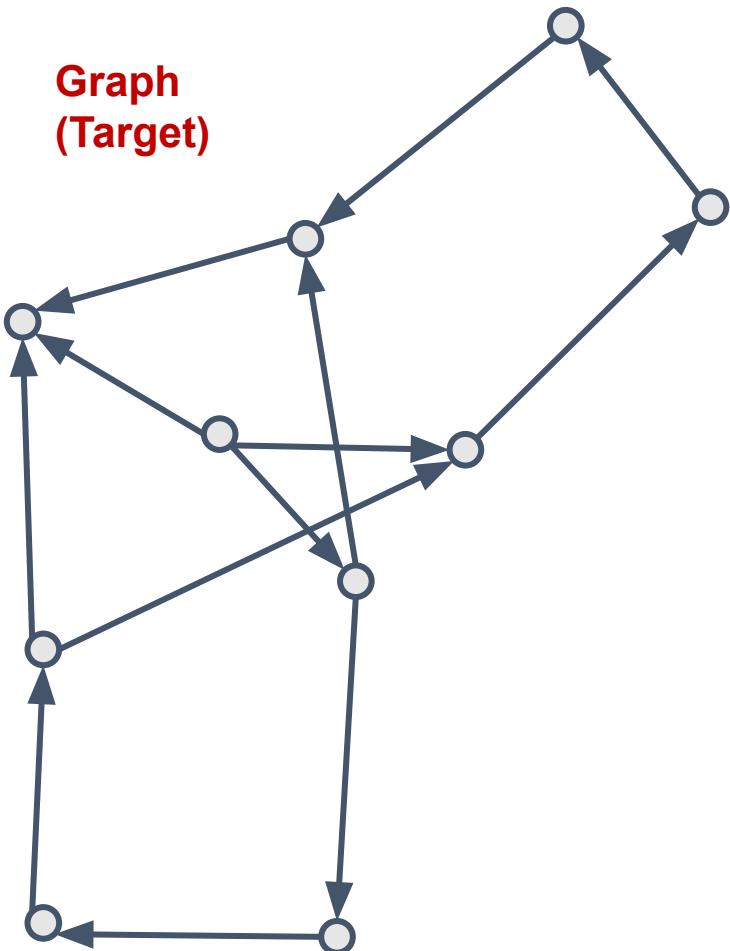


ChapelCon '25  
October 9, 2025

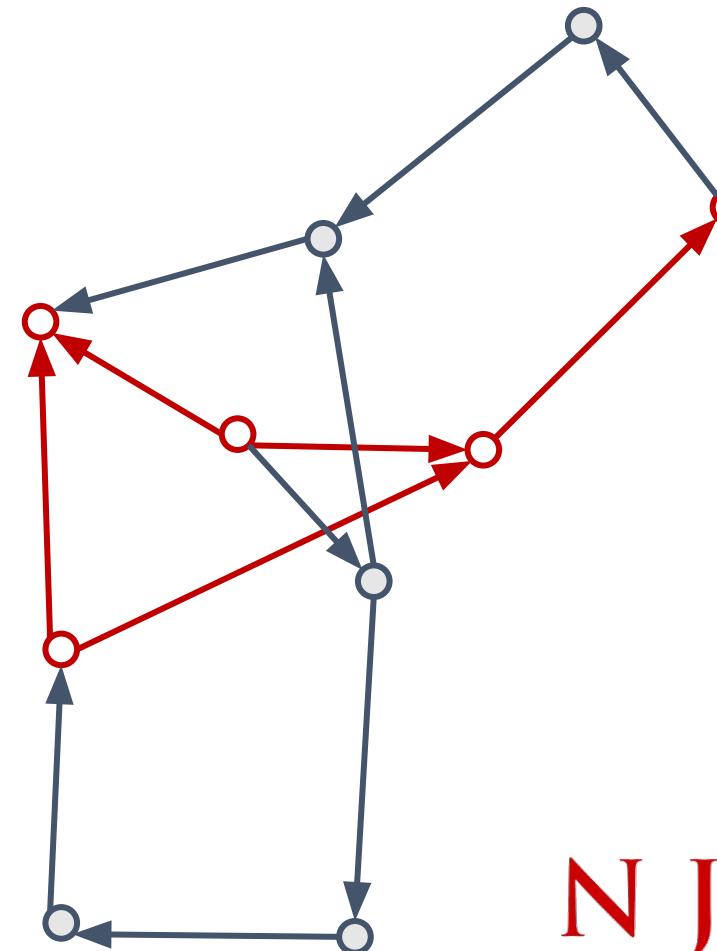
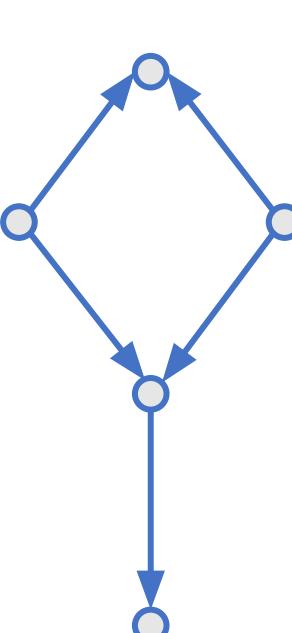


# Problem Statement

Graph  
(Target)



Pattern  
(Motif)



# Motivation

- **Subgraph Isomorphism Challenge**

Identifying small pattern graphs within larger graphs is a complex and computationally heavy problem in graph theory.(NP-Complete)

- **Applications Across Domains**

Subgraph isomorphism impacts neuroscience, biology, social networks, cybersecurity, and fraud detection.

- **Scalability Issues**

Traditional algorithms struggle with large graphs due to exhaustive search causing slow runtimes and memory failures.

- **HiPerMotif Solution**

HiPerMotif introduces a hybrid parallel algorithm improving initialization and scaling for large graph analysis. (Chapel-Arachne)

# Research Background

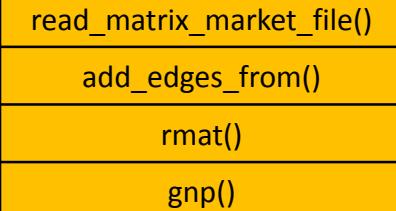


- Ullmann's algorithm, LAD (Labeled Distance), RI / RI-Plus / RI-DS, TurboISO, Glasgow Subgraph Solver, VF2, ...
- NetworkX, DotMotif, iGraph, ...
- VF2 is widely used.
- Great potential to be parallel.
- In neuroscience there are some tools already adapted it.

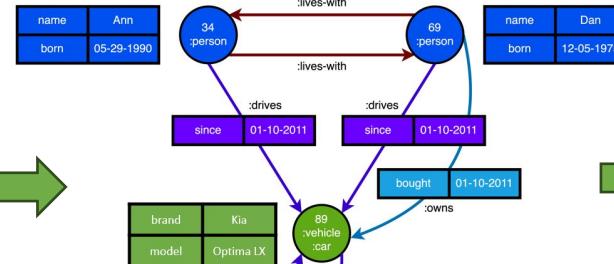
# A Bird's-Eye View of Arachne

id	label	name	born	brand	model
34	person	Ann	1990	NULL	NULL
69	src id	dst id	relationship	since	bought
89	34	69	lives-with	NULL	NULL
89	69	34	lives-with	NULL	NULL
34	89	drives	2011	NULL	
69	89	drives	2011	NULL	
69	89	owns	NULL	2011	
89	89	drives	NULL	NULL	

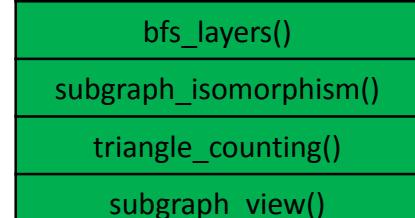
Load in large CSVs, HDF5s, Parquets, matrix market files, etc.



Convert dataframes to graphs or generate your own synthetic graphs.

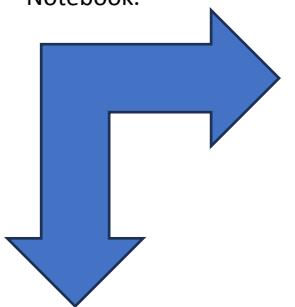


Work with your data as a graph.



Perform analysis or filter for NetworkX, iGraph, or graph-tool.

User edits a Python script or a Jupyter Notebook.

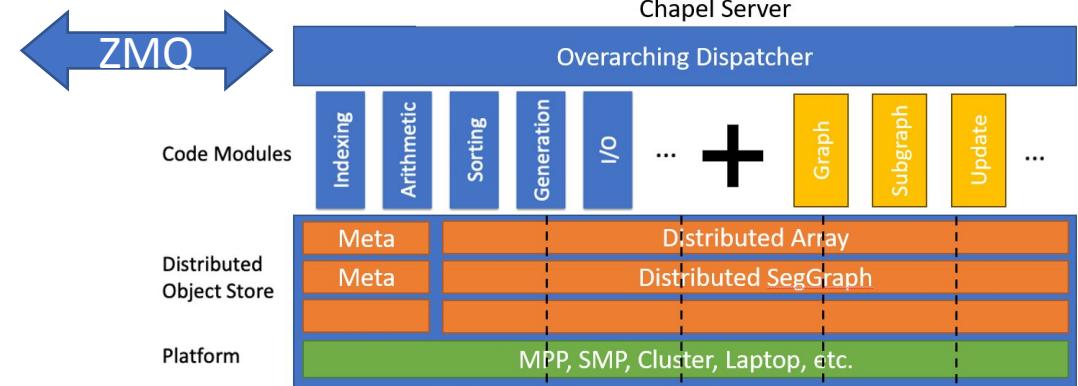


User

```
1. import arkouda as ak
2. import arachne as ar
3.
4. ## Get src and dst from input file.
5.
6. graph = ar.PropGraph()
7.
8. ## Generate label_df and relationships_df from input
9. ## file.
10. graph.load_edge_attributes(relationships_df)
11. graph.load_node_attributes(label_df)
12.
13. ## User generates labels_to_find and
14. ## relationships_to_find:
15. returned_nodes = graph.node_attributes[“column”] == 1
16. returned_edges = graph.edge_attributes[“column”] == 2
17.
18. subgraph_src = ak.in1d(returned_edges[0],
19. returned_nodes),
20. subgraph_dst = ak.in1d(returned_edges[1],
21. returned_nodes)
22.
23. kept_edges = subgraph_src & subgraph_dst
24. subgraph_src = subgraph_src[kept_edges]
25. subgraph_dst = subgraph_dst[kept_edges]
26.
27. subgraph = ar.Graph()
28. subgraph.add_edges_from(subgraph_src, subgraph_dst)
29. ## Run some other operations on subgraph!
```

Easily usable through NetworkX-like API.

OPEN SOURCE: <https://github.com/Bears-R-Us/arkouda-njit> & <https://github.com/Bears-R-Us/arkouda>  
PAPERS & TALKS AT: IEEE HPEC, IEEE HiPC, IEEE IPDPS, ACM PPoPP, SIAM PP, Nature, & MDPI Algorithms



Runs on any hardware, data stays in the back-end, user calls API through Python: powerful and productive. (Image credit: [Reus 2020])

# HiPerMotif

- **Edge-Centric Initialization**

HiPerMotif begins by identifying and validating all first-edge mappings, skipping empty initial mappings.

- **Pattern Graph Reordering**

Structural reordering prioritizes vertices with high connectivity to optimize the matching process.

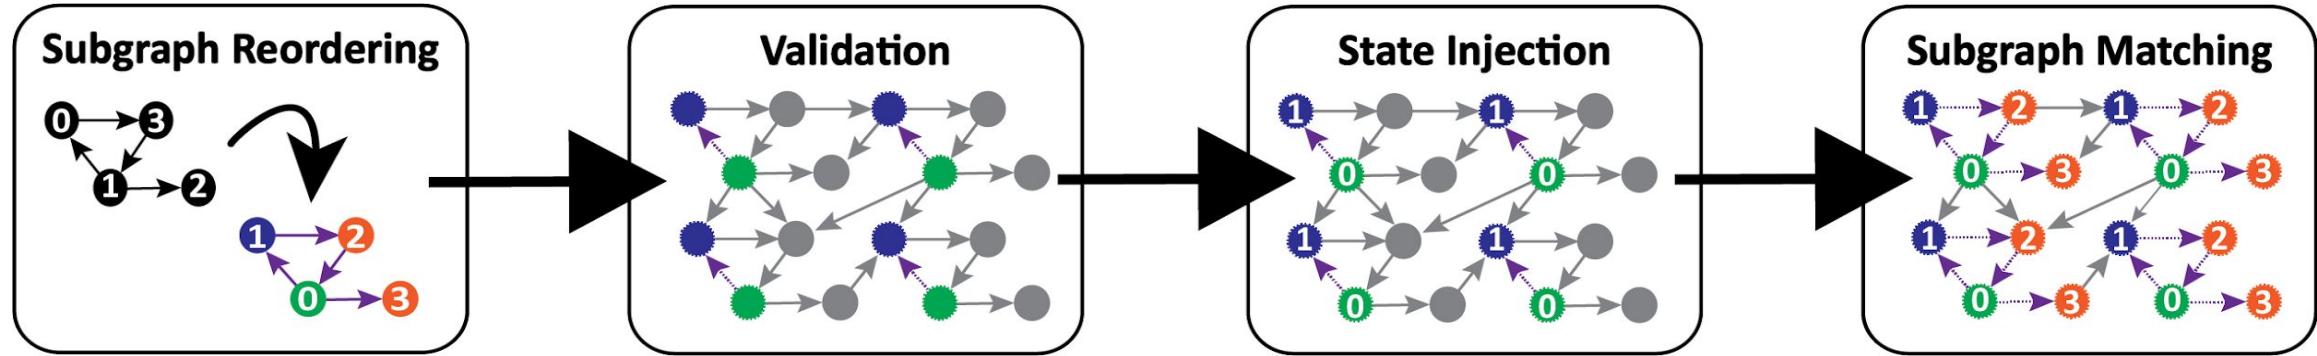
- **Parallel Edge Validation**

Each edge mapping is validated independently, enabling natural parallelization for efficiency.

- **Framework Implementation**

Implemented within Arkouda/Arachne, HiPerMotif scales efficiently for massive dataset analysis.

# HiPerMotif



Allows us to reduce the search space by starting off with vertices that are structurally significant by number of out-neighbors.

Allows us to reduce the search space further by allowing the algorithm to dynamically pick initial vertices and edges that are semantically valuable.

We force subgraph searching to “time travel” and skip unneeded states generated by vertices 0 and 1 of the subgraph and rather start at depth 2.

Match the full subgraphs.

# Performance (Synthetic and Real-World Graphs)

- **Evaluation on Synthetic Graphs**

HiPerMotif was tested on Erdős-Rényi, Barabási-Albert, and Watts-Strogatz graph models with varied densities and sizes.

- **Testing on Real-World Datasets**

Datasets included neuroscience connectomes, communication and social networks, plus a massive human cortex graph.

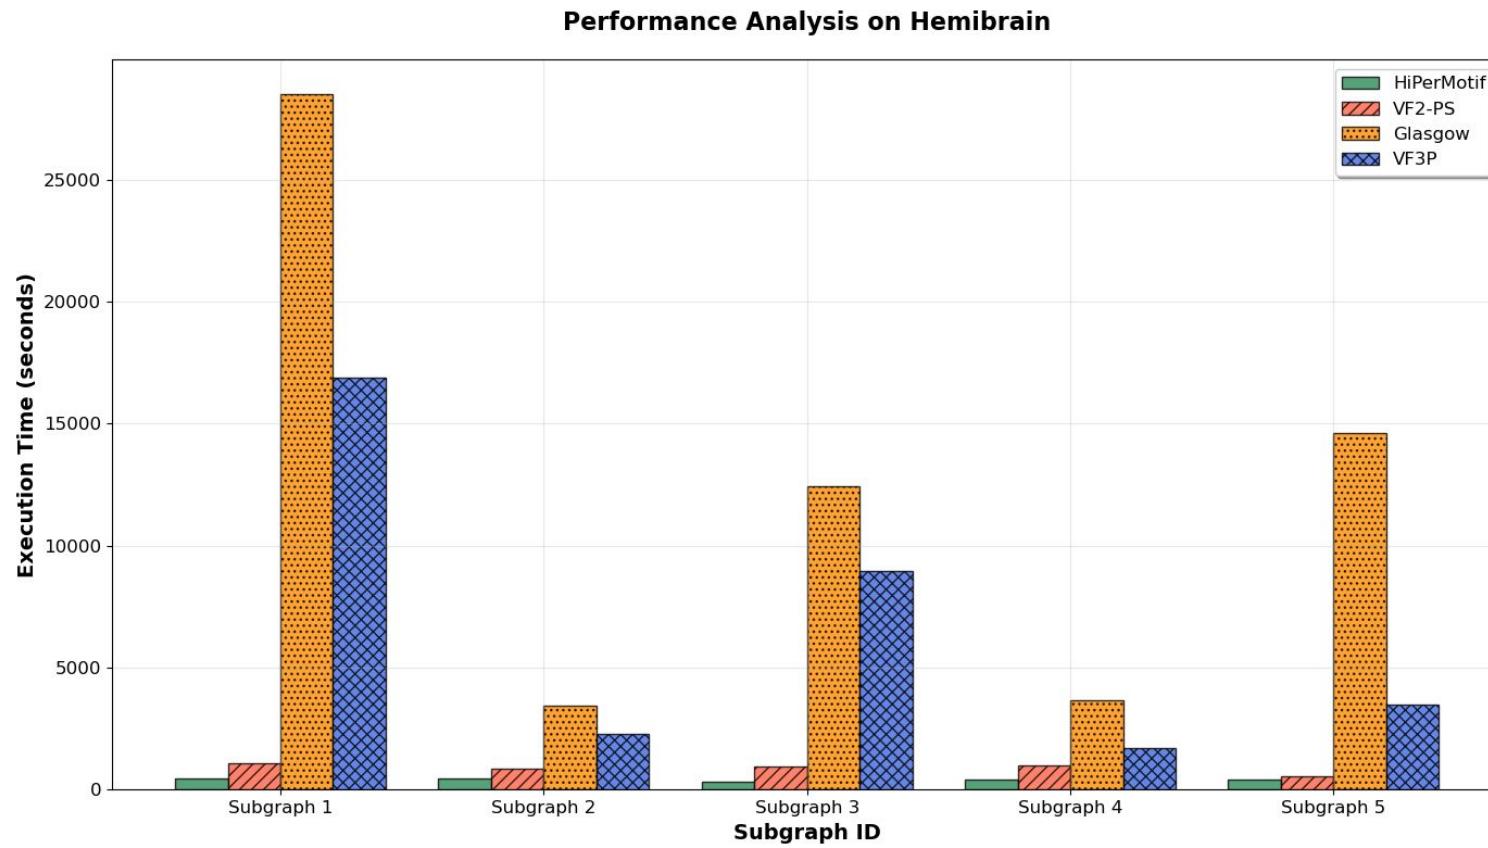
- **Superior Performance Metrics**

HiPerMotif achieved up to **66×** speedup and processed large graphs where baselines failed due to memory limits.

- **Impact of Structural Reordering**

Structural reordering strategy alone contributed up to **5×** speedup, enhancing HiPerMotif's efficiency.

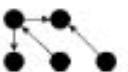
# Neuroscience (Hemibrain Dataset)



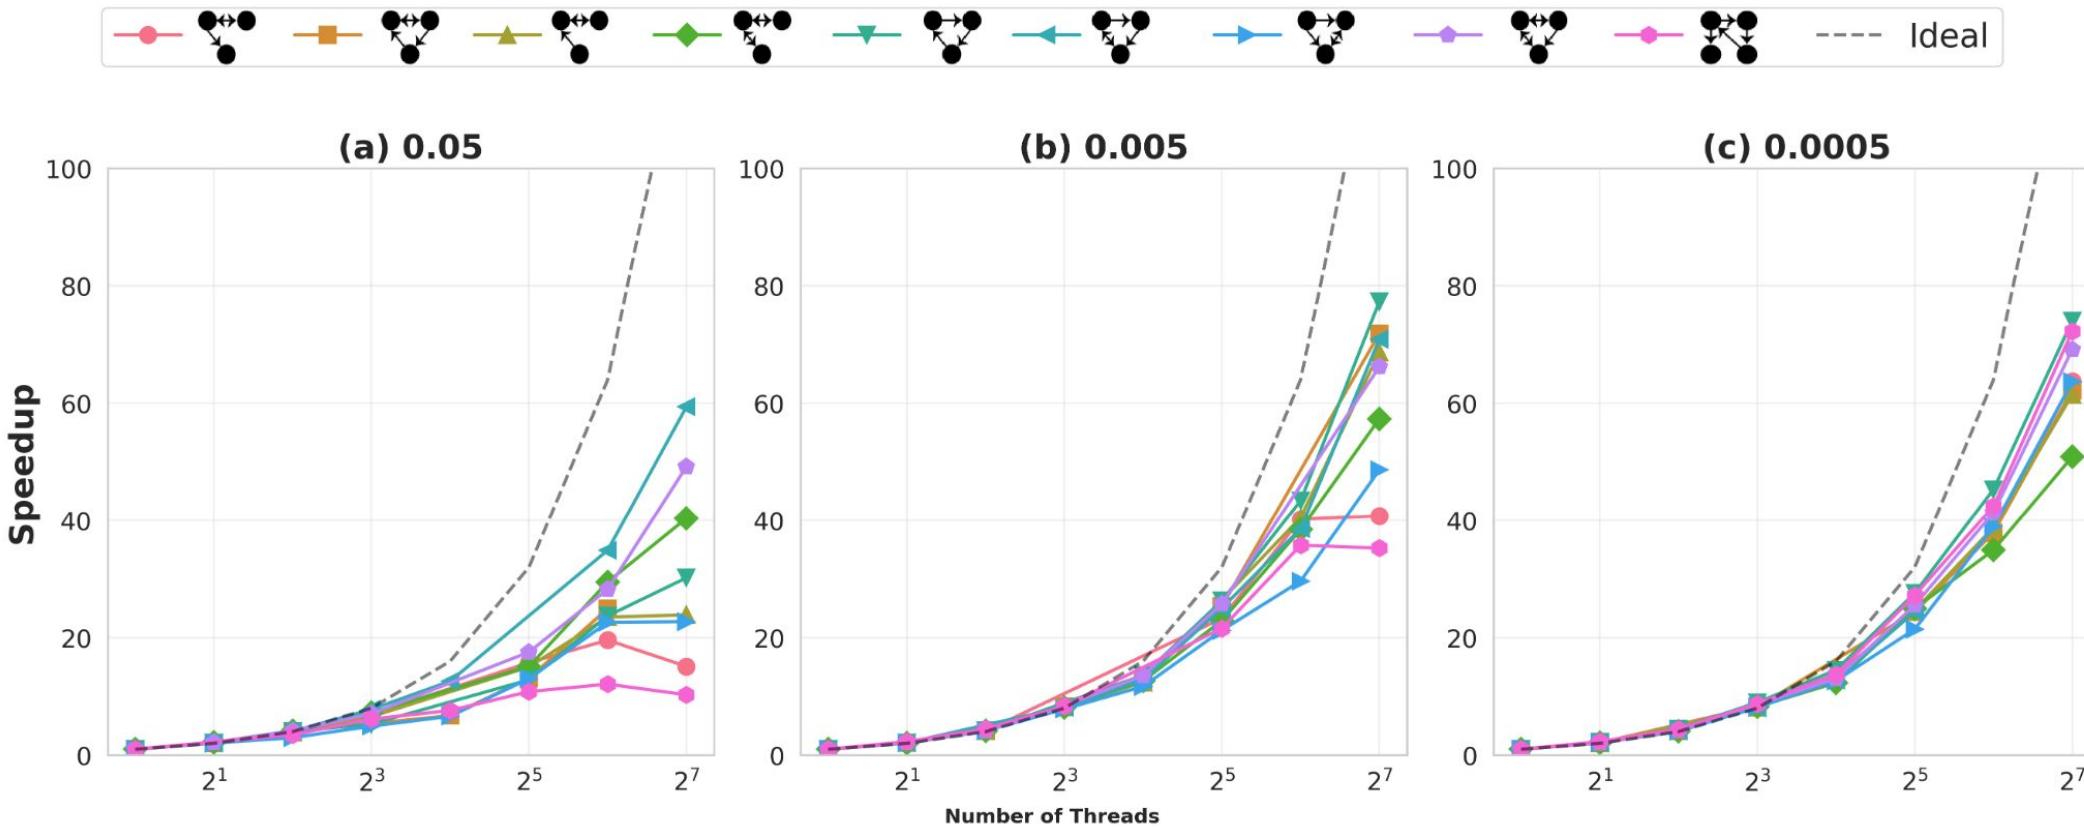
- All Motifs created randomly from 3 to 20 nodes
- Up to 66X speedups
- McCreesh et al, The Glasgow subgraph solver: using constraint programming to tackle hard subgraph isomorphism problem variants
- Carletti et al, A parallel algorithm for subgraph isomorphism

# H01 Dataset (Large-Scale Network)

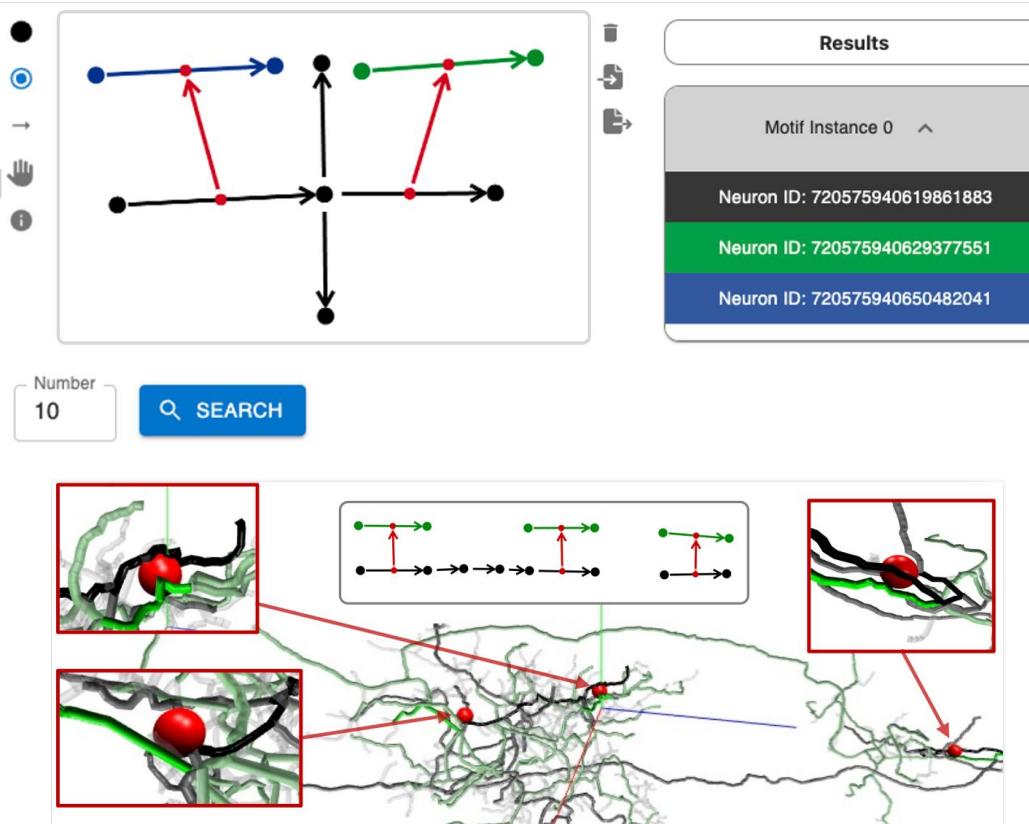
50K vertices and 150 Millions edges  
representing  
a cubic millimeter of human cortex.

Motif	H01 (seconds)
	571.94
	1011.62
	21.23
	363.54
	1209.82

# Parallel speedup



# MOMO: Use Case



M. Shewarega, J. Troidl, et al. / ToMo

## Our Collaboration with Harvard

Subgraphs	Arachne (s)	NetworkX (s)
	<b>2.48</b>	336.45
	<b>3.62</b>	173.75
	<b>2.88</b>	5,980.54
	<b>339.46</b>	16,436.85
	<b>1.56</b>	435.07
	<b>78.77</b>	810.23
	<b>4.10</b>	1,018.23
	<b>38.06</b>	>12,000

Dataset: 13,000 neurons with over 500,000 synaptic connections

using  
Arachne-HiPerMotif  
vs NetworkX VF2:  
Up to **2,000 X**  
faster!



**Thank you all for your attention.**

**HiPerMotif is open-source and available on GitHub, so feel free to explore the code, try it out, and reach out with any feedback.**

**I'd be happy to take any questions you have.**

# VF Family (VF, VF2, VF2+, VF2++, VF3P)

[V. Carletti, P. Foggia, M. Vento, A. Juttner, P. Madarasi, A. Saggese, C. Sansone, at al]

## Algorithm 1 A high level description of VF2

```

1: procedure VF2(Mapping  $m$ , ProblemType  $PT$ )
2:   if  $m$  covers  $V_1$  then
3:     Output( $m$ )
4:   else
5:     Compute the set  $P_m$  of the candidate pairs for extending  $m$ 
6:     for all  $p \in P_m$  do
7:       if  $\text{Cons}_{PT}(p, m) \wedge \neg\text{Cut}_{PT}(p, m)$  then
8:         call VF2(extend( $m, p$ ),  $PT$ )

```

```

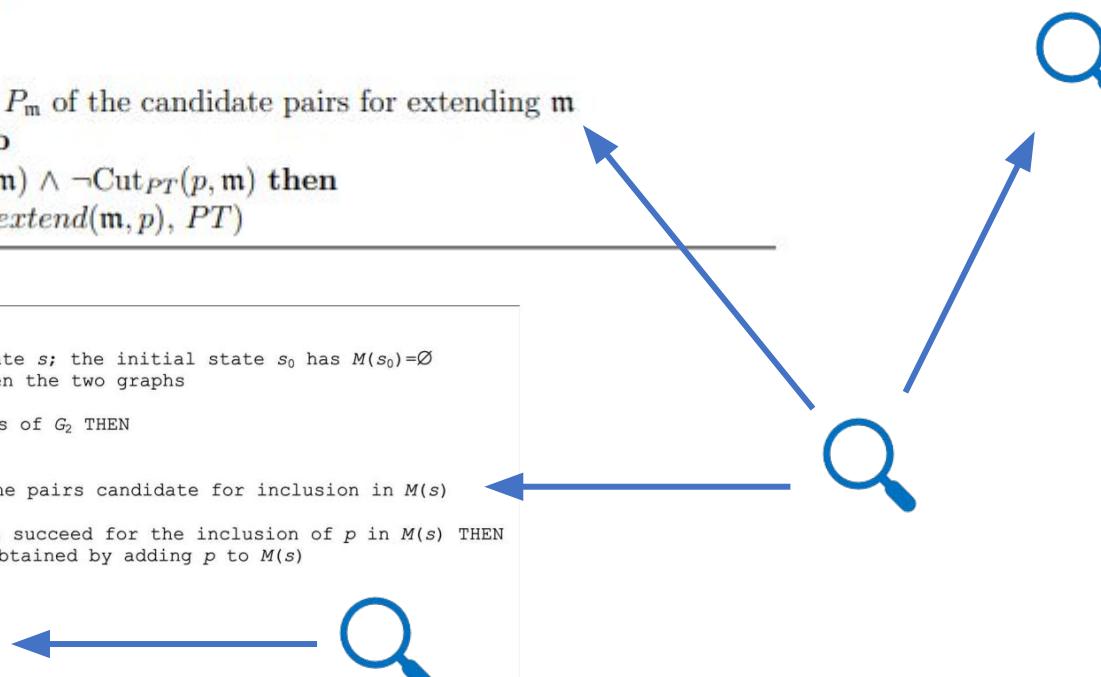
PROCEDURE Match( $s$ )
  INPUT: an intermediate state  $s$ ; the initial state  $s_0$  has  $M(s_0) = \emptyset$ 
  OUTPUT: the mappings between the two graphs

  IF  $M(s)$  covers all the nodes of  $G_2$  THEN
    OUTPUT  $M(s)$ 
  ELSE
    Compute the set  $P(s)$  of the pairs candidate for inclusion in  $M(s)$ 
    FOREACH  $p$  in  $P(s)$ 
      IF the feasibility rules succeed for the inclusion of  $p$  in  $M(s)$  THEN
        Compute the state  $s'$  obtained by adding  $p$  to  $M(s)$ 
        CALL Match( $s'$ )
      END IF
    END FOREACH
    Restore data structures
  END IF
END PROCEDURE Match

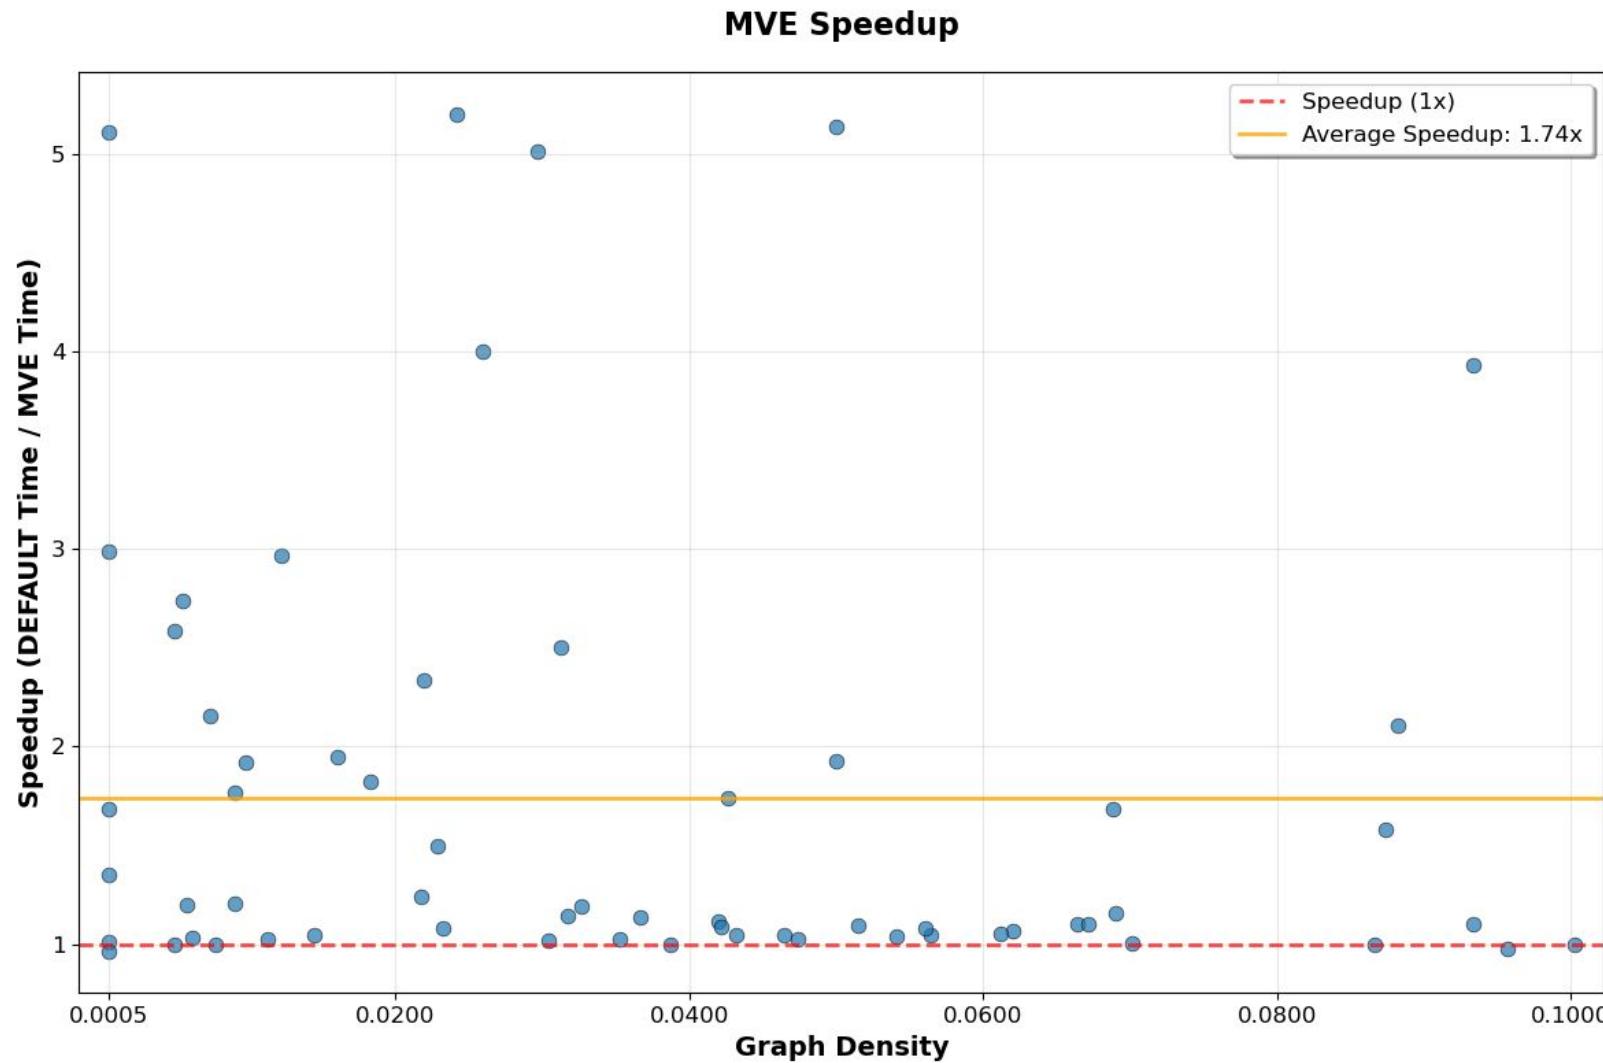
```

- Two vectors,  $\text{core\_1}$  and  $\text{core\_2}$ , whose dimensions correspond to the number of nodes in  $G_1$  and  $G_2$ , respectively, containing the current mapping; in particular,  $\text{core\_1}[n]$  contains the index of the node paired with  $n$ , if  $n$  is in  $M_1(s)$ , and the distinguished value  $\text{NULL\_NODE}$  otherwise. The same encoding is used for  $\text{core\_2}$ .
- Four vectors,  $\text{in\_1}$ ,  $\text{out\_1}$ ,  $\text{in\_2}$ ,  $\text{out\_2}$ , whose dimensions are equal to the number of nodes in the corresponding graphs, describing the membership of the terminal sets. In particular,  $\text{in\_1}[n]$  is nonzero if  $n$  is either in  $M_1(s)$  or in  $T_1^n(s)$ ; similar definitions hold for the other three vectors. The actual value stored in the vectors is the depth in the SSR tree of the state in which the node entered the corresponding set.

Core_1	$ G_1 $
Core_2	$ G_2 $
T_in_1	$ G_1 $
T_out_1	$ G_1 $
T_in_2	$ G_2 $
T_out_2	$ G_2 $



# Structural Reordering(Helps us to prune faster)



# Challenges in Traditional Algorithms

- **Inefficient Candidate Generation**

Traditional algorithms generate large search spaces due to inefficient candidate selection, increasing computation.

- **Rigid Vertex-Ordering Heuristics**

Fixed vertex-ordering heuristics fail to adapt dynamically to varying graph structures and patterns.

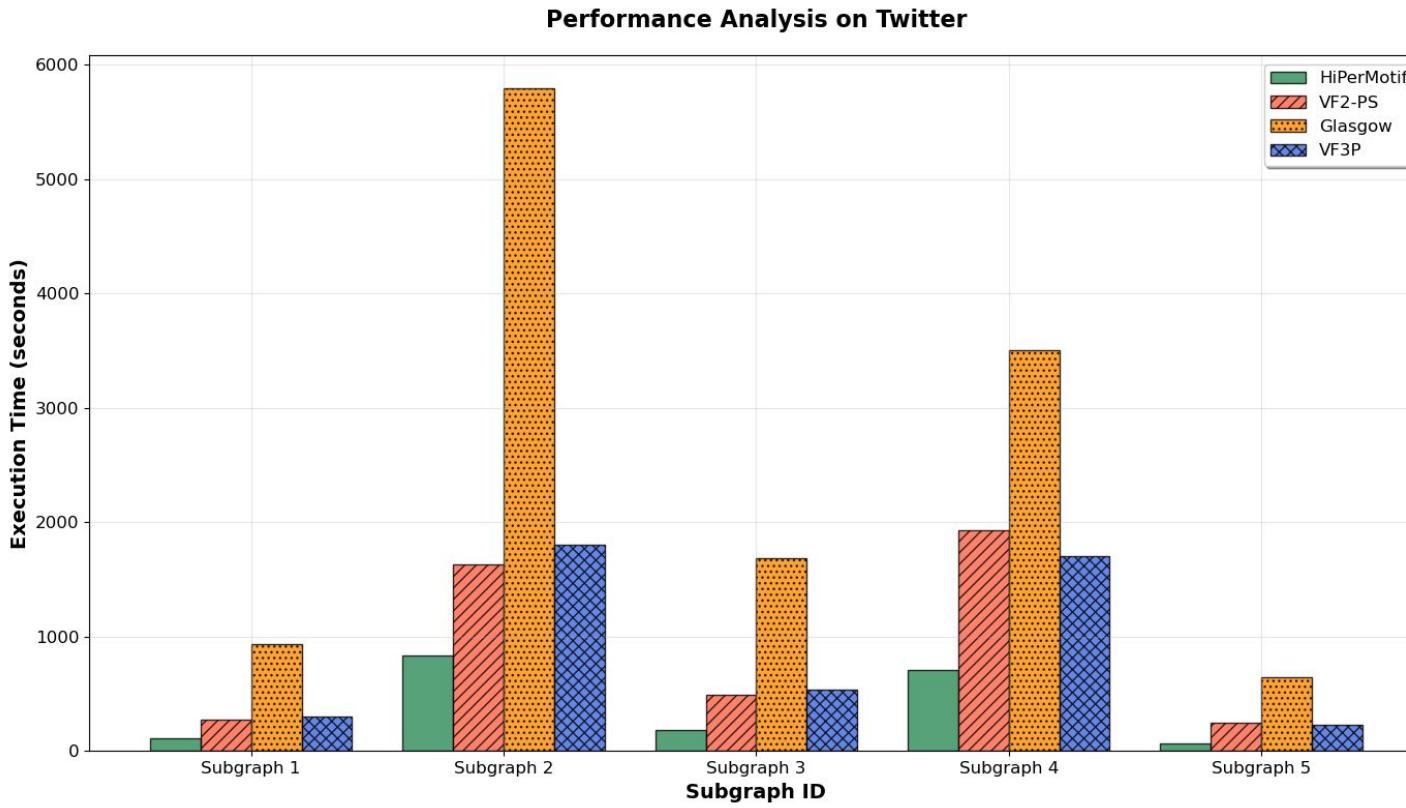
- **High Memory Overhead**

Tracking numerous partial states leads to significant memory consumption in traditional approaches.

- **Limited Parallelization**

Early search stages lack effective parallelization, reducing algorithm scalability on large graphs.

# Twitter



# VF2-PS



The **optimal point** to spawn the tasks  
is immediately after generating  
candidates



we leverage the highly efficient and  
dynamic parallelization capabilities of  
Chapel, which automatically generates  
parallel tasks and assigns them to  
available threads based on the current  
system load

---

**Algorithm 2** Parallel VF2-PS algorithm that generates the mappings of vertices  $u$  from the host graph that are mapped to vertices  $v$  from the subgraph.

---

**Input:** A state  $S_{current}$  with the current mapping information for a given recursive depth  $d$ .

**Output:** Mappings  $M$  of all host graph and subgraph pairs that induce a monomorphism.

```
1:  $M = list(int)$                                 ▷ Parallel-safe list.  
2: if  $d == n_2$  then      ▷  $n_2$  is the size of the subgraph.  
3:   for  $v \in S_{current}.core$  do  
4:      $M.pushBack(v)$   
5:   end for  
6:   return  $M$   
7: end if  
8:  $candidates = getCandidatePairs(S_{current})$   
9: for all  $(u, v) \in candidates$  do  
10:   if  $isFeasible(u, v, S_{current})$  then  
11:      $S_{clone} = S_{current}.clone()$   
12:      $addToTinTout(u, v, S_{clone})$   
13:      $M_{new} = VF2(S_{clone}, d + 1)$   
14:     for  $m \in M_{new}$  do  
15:        $M.pushBack(m)$   
16:     end for  
17:   end if  
18: end for all  
19: return  $M$ 
```

---

## Changes:

- States
- Fast start
- Early Termination
- getCandidatePairs
- getUnmappedNodes
- Support of Properties
- ONLY count
- ONLY time

# Algorithms

---

**Algorithm 1** Structural Reordering

---

```
1: procedure STRUCTURALREORDER(src, dst)
2:   Compute deg( $v$ ) for all  $v \in V$ 
3:    $\mathcal{R} \leftarrow \emptyset$ 
4:    $v^* \leftarrow \arg \max_{v \in V} \sigma(v)$ 
5:   SWAP( $v^*$ ,  $\pi(1)$ )
6:    $\mathcal{R} \leftarrow \mathcal{R} \cup \{v^*\}$ 
7:   while  $|\mathcal{R}| < |V|$  do
8:      $u \leftarrow \pi(|\mathcal{R}|)$ 
9:     Update indeg, outdeg, totaldeg
10:     $N^+(u) \leftarrow \{w \notin \mathcal{R} \mid (u \rightarrow w) \in E\}$ 
11:    if  $N^+(u) \neq \emptyset$  then
12:       $w^* \leftarrow \arg \max_{w \in N^+(u)} \sigma(w)$ 
13:    else
14:       $w^* \leftarrow \arg \max_{w \in V \setminus \mathcal{R}} \sigma(w)$ 
15:    end if
16:    SWAP( $w^*$ ,  $\pi(|\mathcal{R}|+1)$ )
17:     $\mathcal{R} \leftarrow \mathcal{R} \cup \{w^*\}$ 
18:  end while
19:  return ( $src, dst, \pi$ )
20: end procedure
```

---

---

**Algorithm 2** Vertex Validator

---

```
1: procedure VV( $G_1, G_2$ )
2:   vertexFlag =  $[a_1, \dots, a_n]$ 
3:    $T_{\text{in}}^0 \leftarrow |N_{G_1}^{\text{in}}(0)|, T_{\text{out}}^0 \leftarrow |N_{G_1}^{\text{out}}(0)|$ 
4:   for all  $v \in V(G_2)$  do
5:      $T_{\text{in}}^v \leftarrow |N_{G_2}^{\text{in}}(v)|, T_{\text{out}}^v \leftarrow |N_{G_2}^{\text{out}}(v)|$ 
6:     if checkAttributes( $v, 0$ )  $\wedge T_{\text{in}}^v \geq T_{\text{in}}^0 \wedge T_{\text{out}}^v \geq T_{\text{out}}^0$  then
7:       vertexFlag[ $v$ ]  $\leftarrow$  true
8:     end if
9:   end for
10:  return vertexFlag
11: end procedure
```

---

# Algorithms

---

**Algorithm 3** Edge Validator
 

---

```

1: procedure EV( $u, v, s$ )
2:    $T_{\text{in/out}}^{u,v} \leftarrow N_{G_2}^{\pm}(u, v); T_{\text{in/out}}^{0,1} \leftarrow N_{G_1}^{\pm}(0, 1)$ 
3:   if  $\neg \text{match}(v, 1)$  then return false
4:   end if
5:    $e_1 \leftarrow \text{getEdgeId}(u, v); e_1^r \leftarrow \text{getEdgeId}(v, u)$ 
6:    $e_2 \leftarrow \text{getEdgeId}(0, 1); e_2^r \leftarrow \text{getEdgeId}(1, 0)$ 
7:   if  $\neg \text{match}(e_1, e_2) \vee (e_2^r \neq -1 \wedge e_1^r = -1)$  then
8:     return false
9:   end if
10:  if  $e_1^r \neq -1 \wedge e_2^r \neq -1 \wedge \neg \text{checkAttributes}(e_1^r, e_2^r)$ 
11:    then
12:      return false
13:    end if
14:    if  $|T_{\text{in}}^v| < |T_{\text{in}}^1| \vee |T_{\text{out}}^v| < |T_{\text{out}}^1|$  then return false
15:    end if
16:     $N_{u,v} \leftarrow T_{\text{in}}^u \cup T_{\text{out}}^u \cup T_{\text{in}}^v \cup T_{\text{out}}^v$ 
17:     $N_{0,1} \leftarrow T_{\text{in}}^0 \cup T_{\text{out}}^0 \cup T_{\text{in}}^1 \cup T_{\text{out}}^1$ 
18:    if  $|N_u \cap N_v| < |N_0 \cap N_1|$  then return false
19:    end if
20:     $s.T_{\text{in/out}}^{G_2} \leftarrow T_{\text{in/out}}^u \cup T_{\text{in/out}}^v \setminus \{u, v\}$ 
21:     $s.T_{\text{in/out}}^{G_1} \leftarrow T_{\text{in/out}}^0 \cup T_{\text{in/out}}^1 \setminus \{0, 1\}$ 
22:     $s.\text{depth} \leftarrow s.\text{depth} + 2$ 
23:     $s.\text{core}[0] \leftarrow u; s.\text{core}[1] \leftarrow v$ 
24:    return true
25:  end procedure
  
```

---



---

**Algorithm 4** HiPerMotif Algorithm
 

---

```

1: procedure HiPERMOTIF( $G_1, G_2$ )
2:    $M \leftarrow \text{new list(int)}$ 
3:   for all  $e \in E_2$  do            $\triangleright$  Parallel over edges
4:      $u \leftarrow \text{src}(e), v \leftarrow \text{dst}(e)$ 
5:     if  $\text{vertexFlag}[u] \wedge u \neq v$  then
6:        $s \leftarrow \text{new State}(|V_2|, |V_1|)$ 
7:       if EV( $u, v, s$ ) then
8:          $M_{\text{new}} \leftarrow \text{VF2-PS}(s, 2)$ 
9:          $M \leftarrow M \cup M_{\text{new}}$ 
10:        end if
11:      end if
12:    end for
13:    return  $M$ 
14:  end procedure
  
```

---