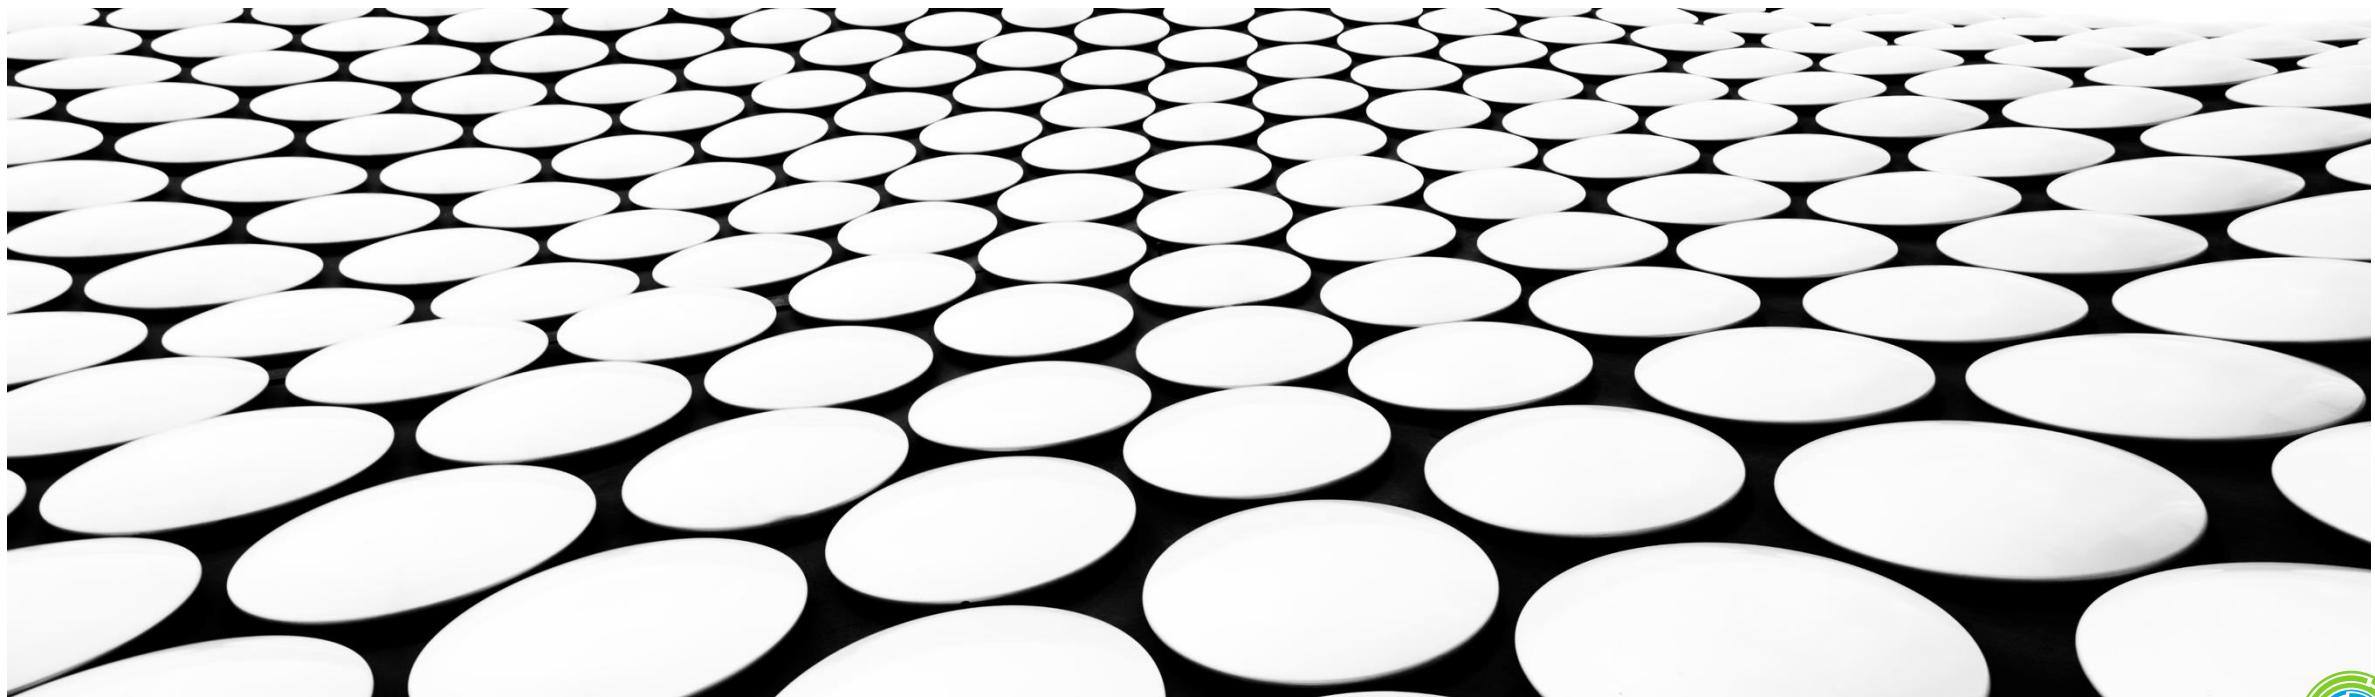


---

---

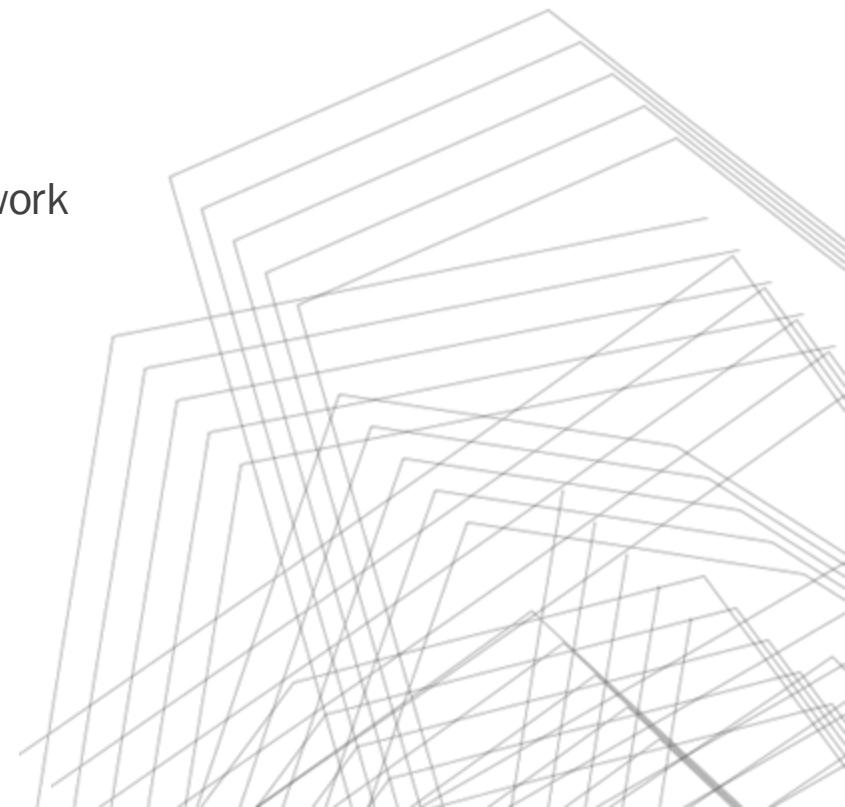
# REPARTITIONING FOR PERFORMANCE: FLEXIBLE DATA MOVEMENT IN CHAPEL

RYAN KECK

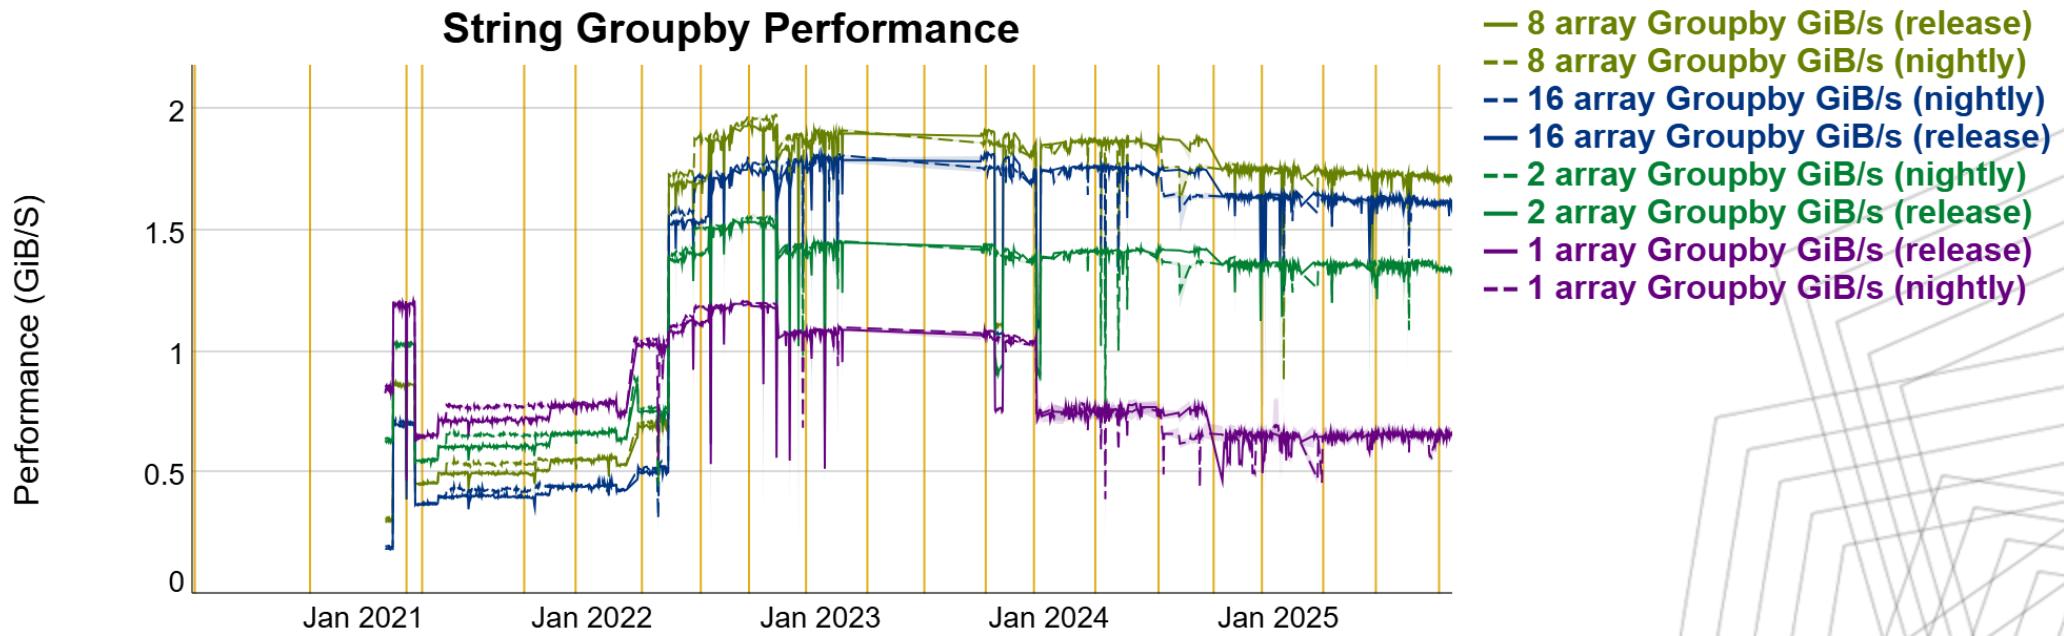


# INTRODUCTION

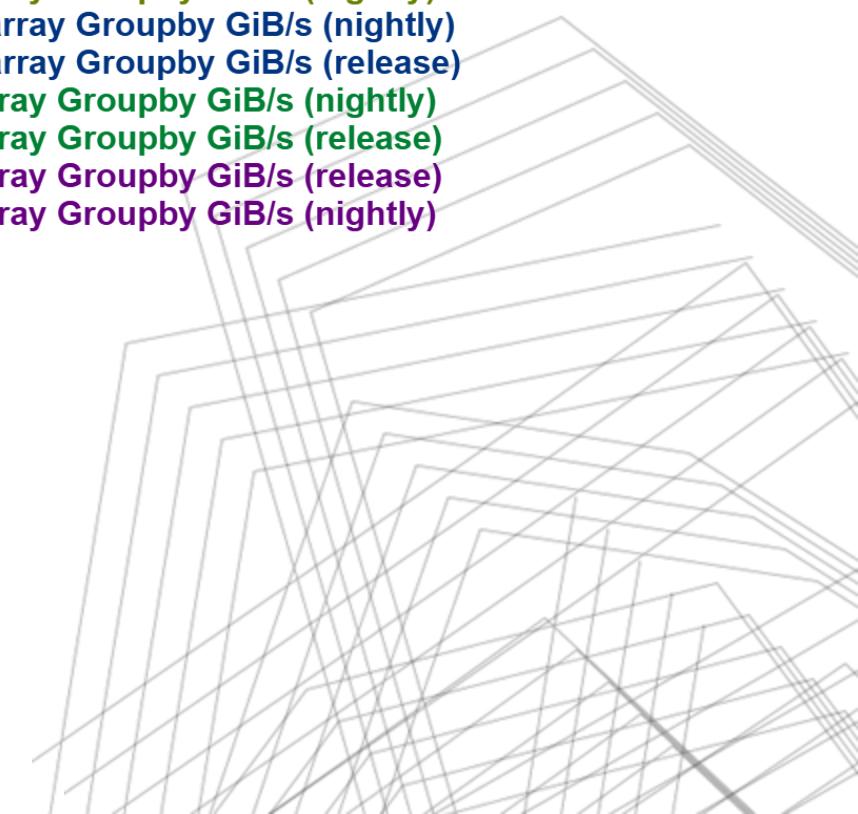
- Software engineer working on Arkouda, a Python+Chapel data analytics framework
- Focus area: performance optimization



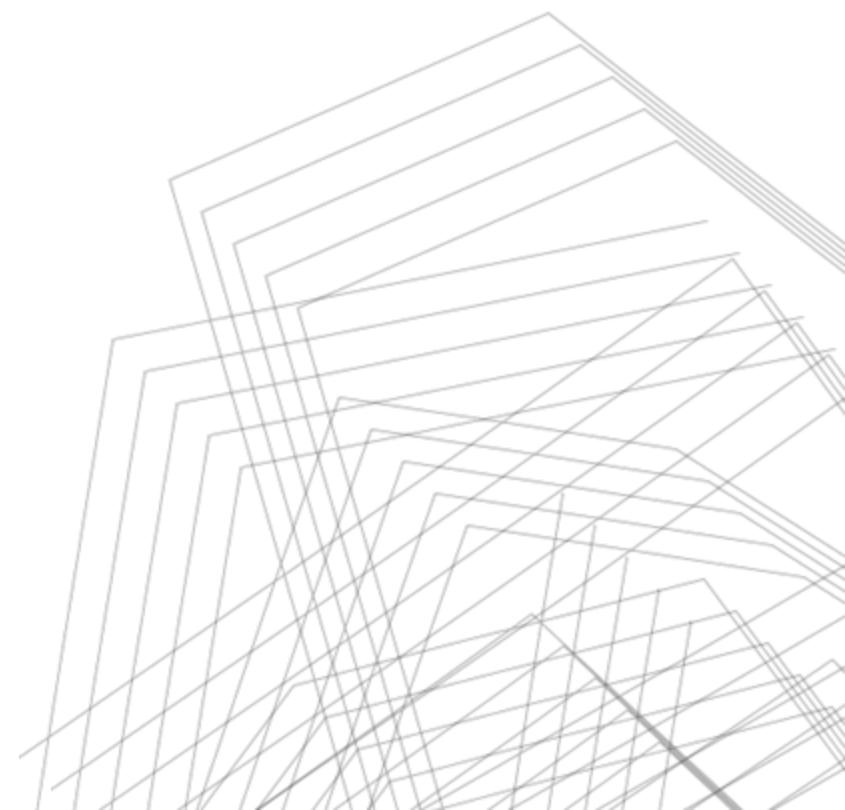
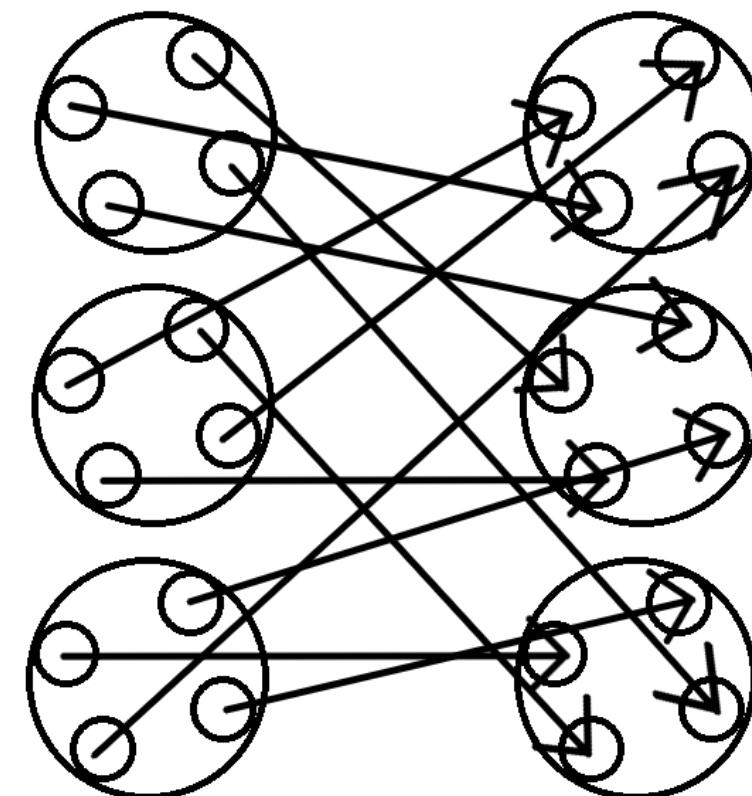
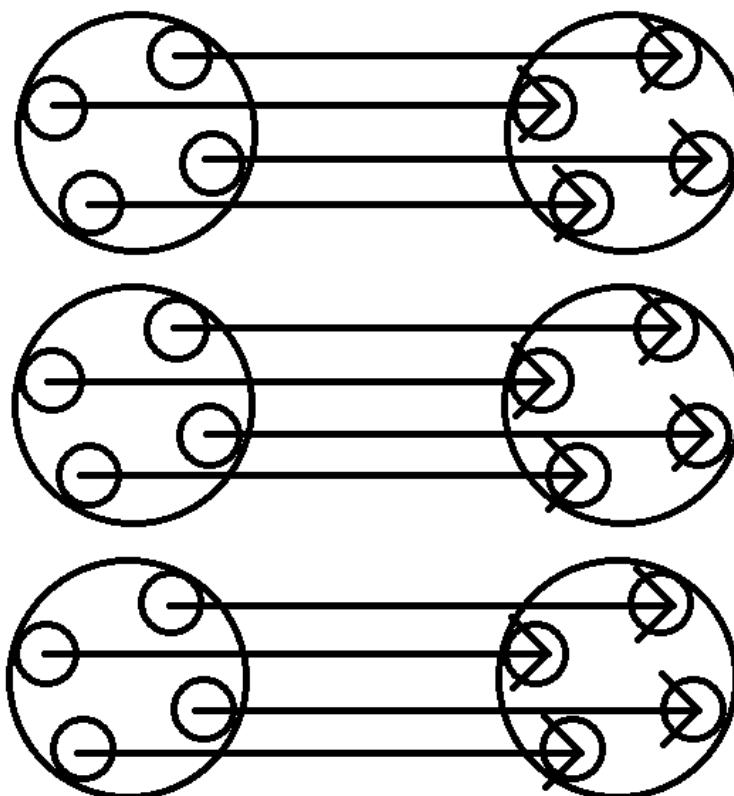
# AN UNEXPECTED BOTTLENECK



- 8 array Groupby GiB/s (release)
- 8 array Groupby GiB/s (nightly)
- 16 array Groupby GiB/s (nightly)
- 16 array Groupby GiB/s (release)
- 2 array Groupby GiB/s (nightly)
- 2 array Groupby GiB/s (release)
- 1 array Groupby GiB/s (release)
- 1 array Groupby GiB/s (nightly)

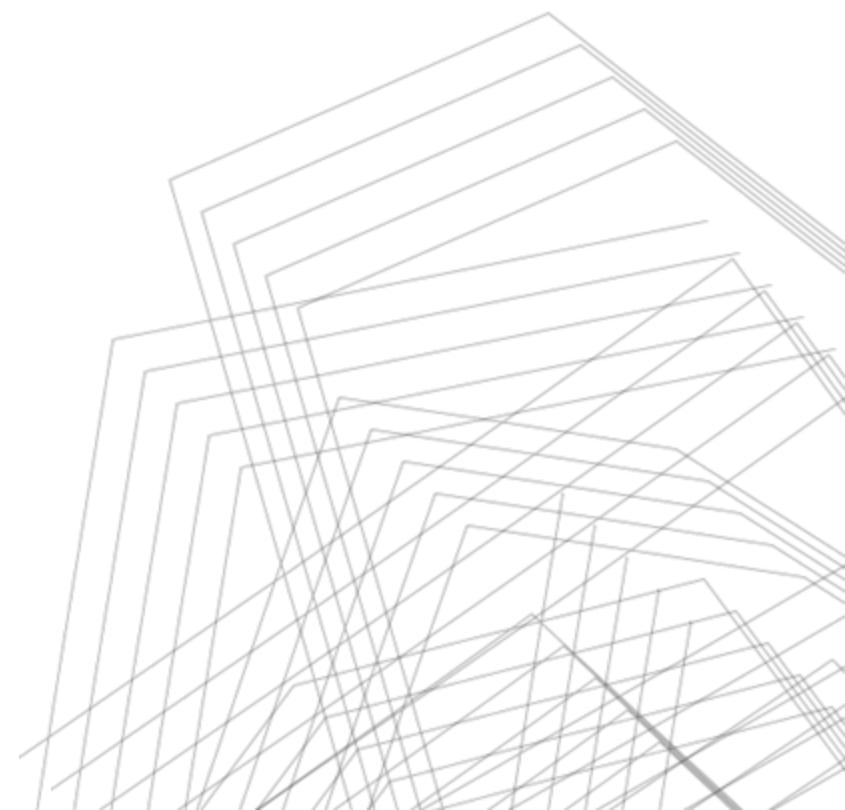


# PARALLELISM



# HASH BASED SHARDING

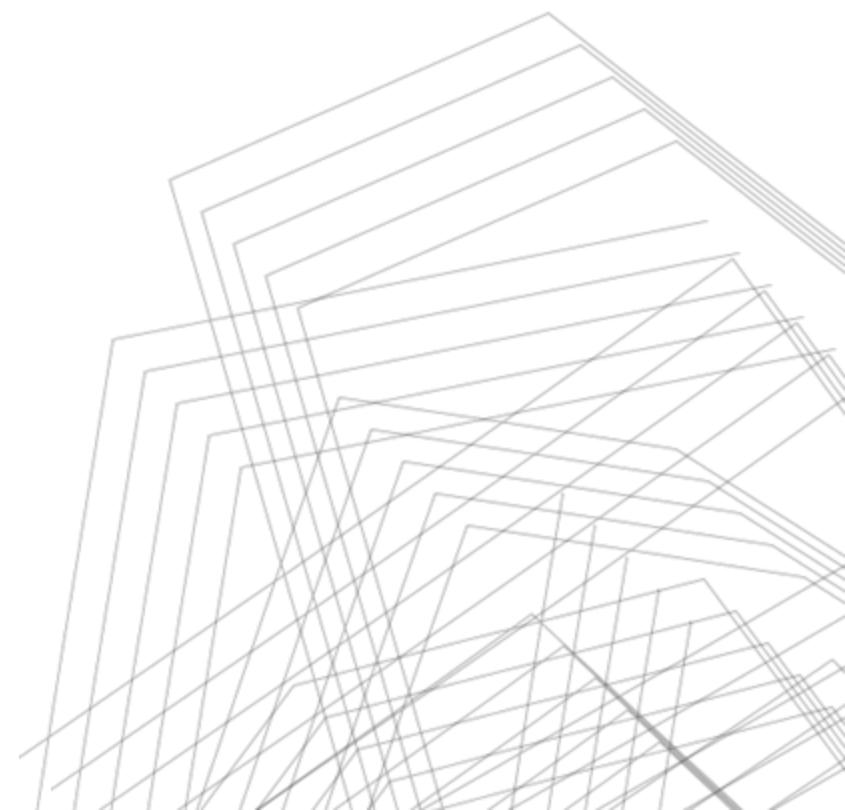
- Compute `hash(value) % numLocales`
- Potential drawbacks
  - Easy to generate data that would all get sent to the same locale
  - An array could be entirely composed of one value



# REPARTITIONING MODULE

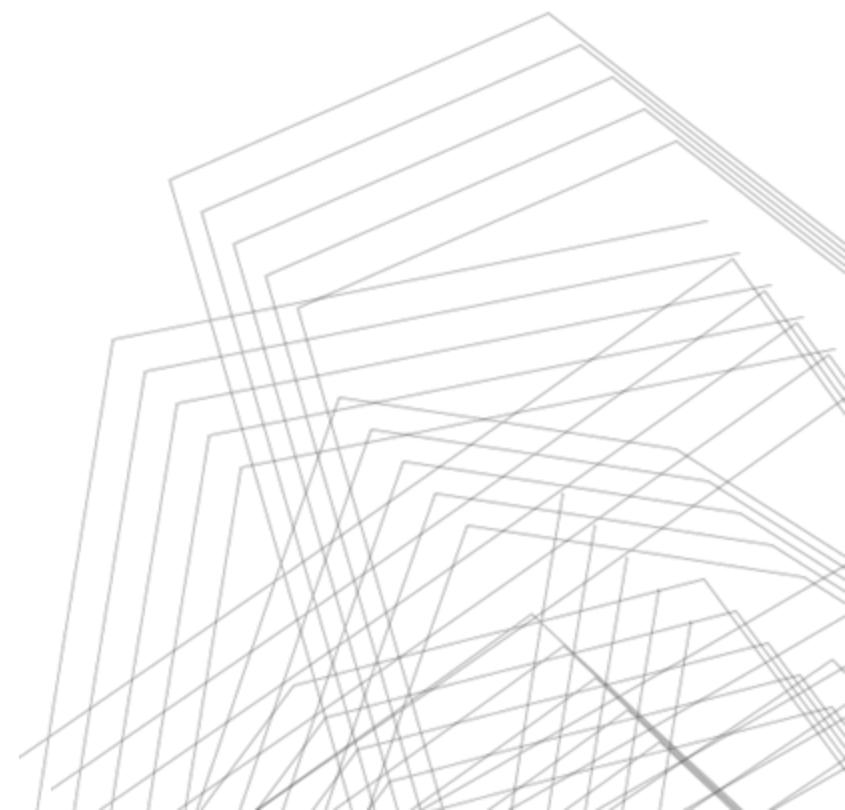


- Redistributions data across locales based on a target mapping
- Includes the option to hash the data for you
- Performs the data movement in a single coordinated step
- Returns a new object with data placed on the right locales



# IMPLEMENTATION DETAILS

- Data movement: Aggregators vs. bulk transfer
  - Decided on bulk transfer because it is simpler when destination indices are unknown
- Data structure for transferring: Lists vs. Arrays
  - Lists: No padding but many copies
  - Arrays: Better performance but some padding

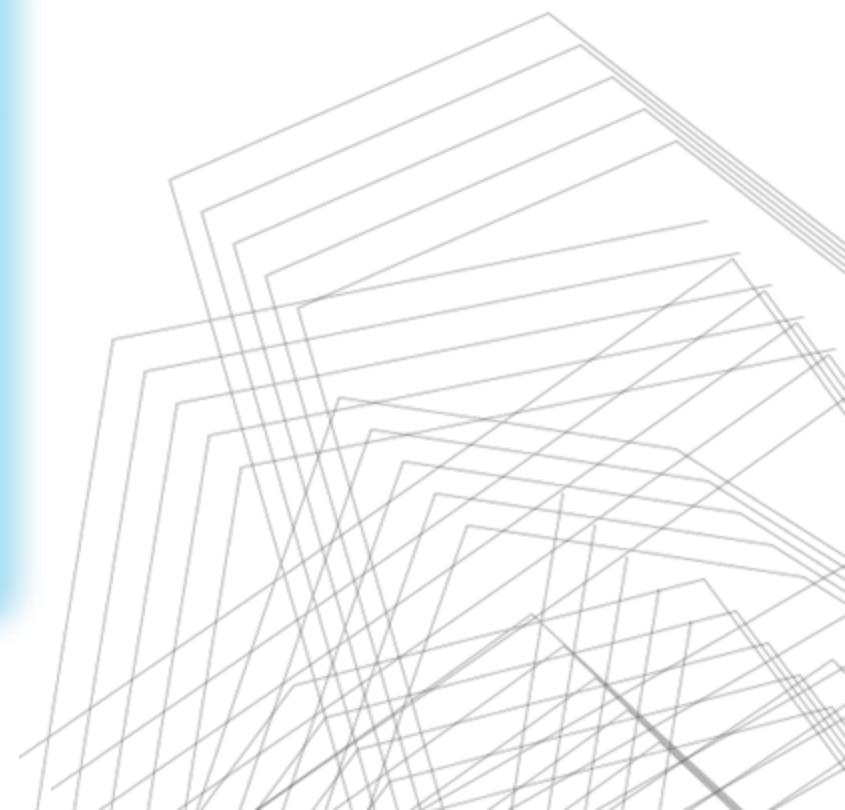


# INNERARRAY

```
record innerArray {
    type t;
    var Dom: domain(1);
    var Arr: [Dom] t;

    proc init(in Dom: domain(1), type t) {
        this.t = t;
        this.Dom = Dom;
    }

    proc init(type t) {
        this.t = t;
    }
}
```



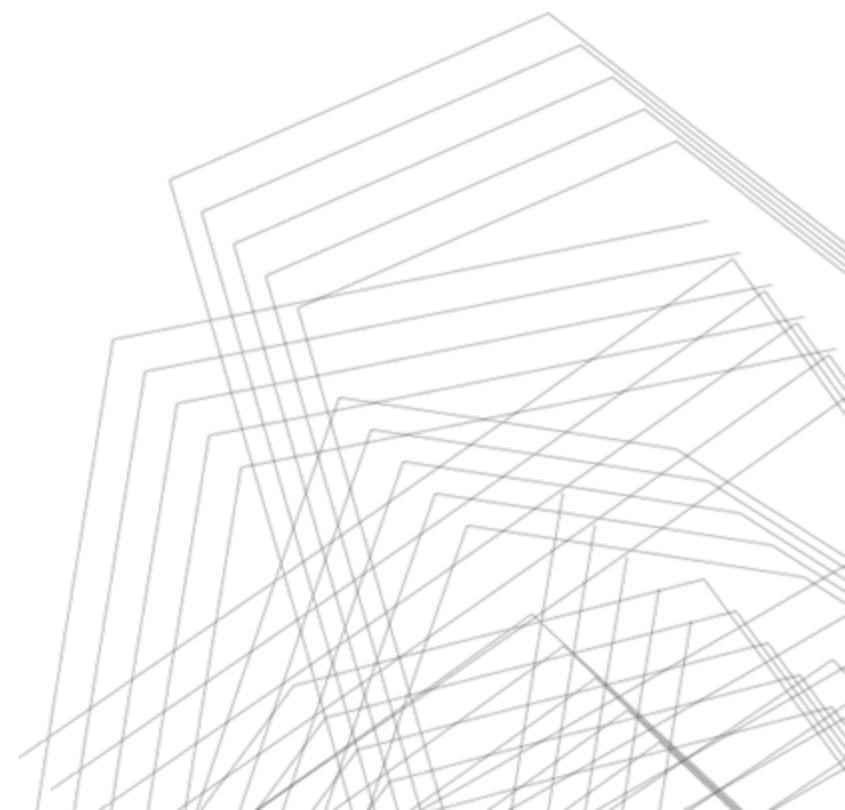
Thanks to Engin Kayraklıoğlu!

# RESULTS

N	Branch/Main	Locale	get	put	execute_on	execute_on_nb
1,000,000	Main	0	295	253	380	721
1,000,000	Main	Avg of others	315	401	383	0
1,000,000	Branch	0	239	928	84	1282
1,000,000	Branch	Avg of others	116	1209	116	155
10,000,000	Main	0	967	925	1052	721
10,000,000	Main	Avg of others	987	1072	1055	0
10,000,000	Branch	0	239	928	84	1282
10,000,000	Branch	Avg of others	116	1209	116	155
100,000,000	Main	0	7544	7502	7629	721
100,000,000	Main	Avg of others	7562	7648	7630	0
100,000,000	Branch	0	238	925	84	1283
100,000,000	Branch	Avg of others	117	1210	116	155

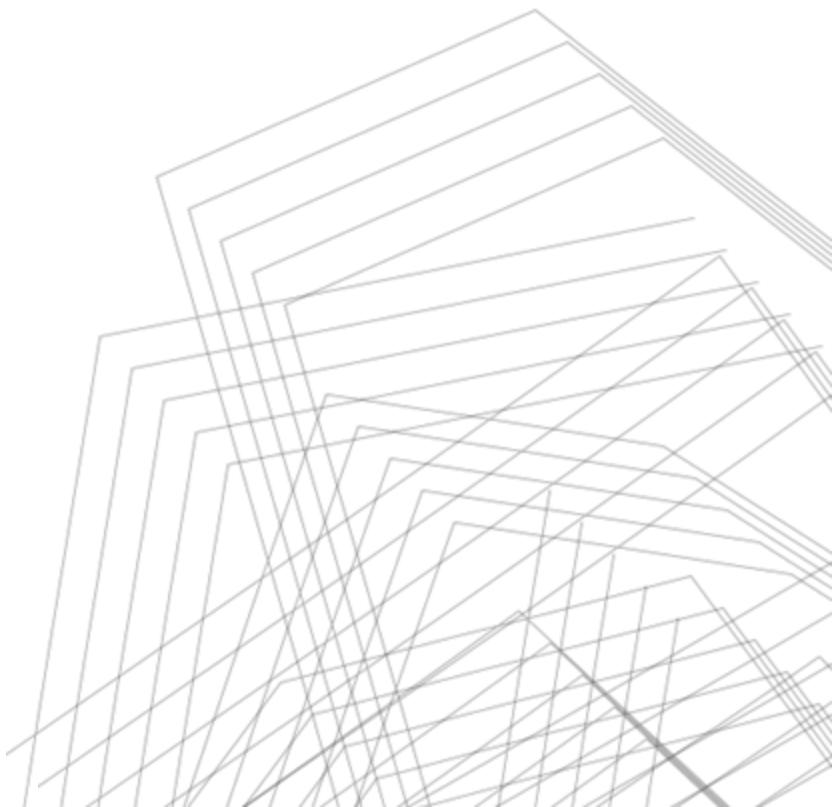
# CONCLUSION

- Flexible mechanism to reorganize data across locales
- Reduces communication overhead and improves scalability
- Provides a foundation for future optimization (e.g., GroupBy)



---

**THANK YOU**



---

# QUESTIONS?

