

# Transformer From Scratch

Thirin Sastarasadhit: [thirin.sastarasadhit@gmail.com](mailto:thirin.sastarasadhit@gmail.com)  
 Prof. Kenjiro Taura: [tau@eidos.ic.i.u-tokyo.ac.jp](mailto:tau@eidos.ic.i.u-tokyo.ac.jp)

## Methodology

### Models:


- C++
  - Chapel
  - Pytorch A
  - PyTorch B
- } Implemented from scratch
- } From Transformer-from-scratch
- } PyTorch A with the transformer layer replaced with `torch.nn.Transformer`

PyTorch A was the original implementation from Transformer-from-scratch


## Small-Size Model on Single Thread

Tested on Machine A with Small-Size configuration


Time spent on each layer (us) for one training iteration  
 Slowness mainly came from forward-pass



Time spent on each layer (us) during single forward pass iteration  
 Softmax and Dropout layer are slow, other layer just as good as others.



Time spent on each layer (us) during single backward pass iteration  
 Chapel do as good as others in backward pass



## Encountered performance issues and tricky solutions

### Matrix Representation

- 1D array is used
- Multi-dimensional array is slow
- Nested arrays cause non-continuous array



### Matrix Multiplication

- Block tiling with 64x64 block-size
- Chapel performs better/worse than C++ at some specific matrix size

### Matrix Operation

- Passing an array view may prevent certain optimizations
- Passing start, end, and array separately influence more optimizations

### Softmax

- Major cause of the model slowness
- No exponential vectorization is used.

### Dropout

- Use the integer version of `randomStream.fill()`
- Slower random number generator
- Another main cause of the model slowness

### Multihead Attention

- Mysteriously requires changing `param` to `var` in the config file to make the optimization to occur in certain place

### ReLU

- The backward pass must be divided into two separate sections for optimization to take place
- The compiler generated code for the forward pass differs from C++ version which is slower when tested on full-size model

## Test Environment:

Property	Small-Size Model on Single Thread	Full-Size Model on Single and Multiple Threads
Machine	Machine A	Machine B
Model Size	Small	Full
Iteration	500	40

- Models were run on English-Italian machine translation task
- C++ and Chapel code structure are identical
- The Degree of parallelism is the same in both C++ and Chapel
- Timers were inserted in each layer




Detailed Result

Code

## Full-Size Model on Single and Multiple Threads


Tested on Machine B with Full-Size configuration



Single Thread

Multiple Threads


Time spent on each layer (us) for one training iteration  
 PyTorch is fast mainly because of its optimized linear algebra library.  
 Chapel performs most layers as well as C++, with a better forward pass in some cases



Single Thread

Multiple Threads

Time spent on each layer (us) during single forward pass iteration  
 Chapel performs better on the linear layer in a single-threaded setting. This is mainly the reason it is faster than C++ in this case.



Single Thread

Multiple Threads

Time spent on each layer (us) during single backward pass iteration  
 In the backward pass, most layers perform as well as in C++, but achieve much greater speedup.

