

Transformer From Scratch

Thitrin Sastarasadhit¹, Prof. Kenjiro Taura²

¹Chulalongkorn University, Bangkok, Thailand

²University of Tokyo, Tokyo, Japan

Agenda

- Methodology
- Small-Size Model on Single Thread
- Full-Size Model on Single and Multiple Threads
- Productivity
- Conclusion

Methodology



Models

There are 4 models:

- C++ } Implemented from scratch
- Chapel } From [Transformer-from-Scratch](#)
- PyTorch A } PyTorch A with the transformer layer replaced with `torch.nn.Transformer`
- PyTorch B }

**This project does performance tests on CPU, single thread, and multiple threads*

**The Chapel and C++ implementations were very similar; all variables could be mapped from one to the other.*

Test Environments

Property	Machine A	Machine B
CPU	AMD Ryzen 7 4800H with Radeon Graphics	Intel(R) Xeon Phi(TM) CPU 7250 @ 1.40GHz
RAM	6.67 GB	204.45 GB
Clang	Ubuntu clang version 19.1.1 (1ubuntu1) Target: x86_64-pc-linux-gnu Thread model: posix	clang version 19.1.3 Target: x86_64-unknown-linux-gnu Thread model: posix
Chapel	chpl version 2.4.0 built with LLVM version 19.1.1 available LLVM targets: xtensa, m68k, xcore, x86-64, x86, wasm64, wasm32, ve, systemz, sparcel, sparcv9, sparc, riscv64, riscv32, ppc64le, ppc64, ppc32le, ppc32, nvptx64, nvptx, msp430, mips64el, mips64, mipsel, mips, loongarch64, loongarch32, lanai, hexagon, bpfeb, bpfel, bpf, avr, thumbeb, thumb, armeb, arm, amdgcn, r600, aarch64_32, aarch64_be, aarch64, arm64_32, arm64, x86-64, x86	chpl version 2.4.0 built with LLVM version 19.1.3 available LLVM targets: amdgcn, r600, nvptx64, nvptx, aarch64_32, aarch64_be, aarch64, arm64_32, arm64, x86-64, x86
Python	Python 3.11.13 PyTorch : 2.3.0 Numpy : 2.3.0	Python 3.11.13 PyTorch : 2.5.1 Numpy : 2.0.1

Model Configuration

Parameter	Machine A	Machine B	Description
dModel	32	512	Dimension of embedding layer of the encoder and decoder
sequenceLength	128	256	Maximum length of input sequence
dFF	256	2048	Dimension of the feed-forward layer inside the encoder and decoder
N	6	6	Number of transformer encoder, decoder layers (stacked).
head	8	8	Number of attention heads in multi-head attention layer
srcVocab	15700	15700	Size of source vocabulary (number of unique tokens).
tgtVocab	22470	22470	Size of target vocabulary

The model architecture are based on the [Attention Is All You Needed](#) paper

Performance Measurement

- Timers are inserted into each layers
- The Models were trained on English-Italian machine translation task, the dataset were taken from [opus_book](#)
- The time of each iteration of each was gathered, trimming 10% fastest and slowest iterations
- The average and standard deviation were calculated

The Hugging Face link for opus_book: https://huggingface.co/datasets/Helsinki-NLP/opus_books
The model was executed for 500 and 40 iterations on Machines A and B, respectively.

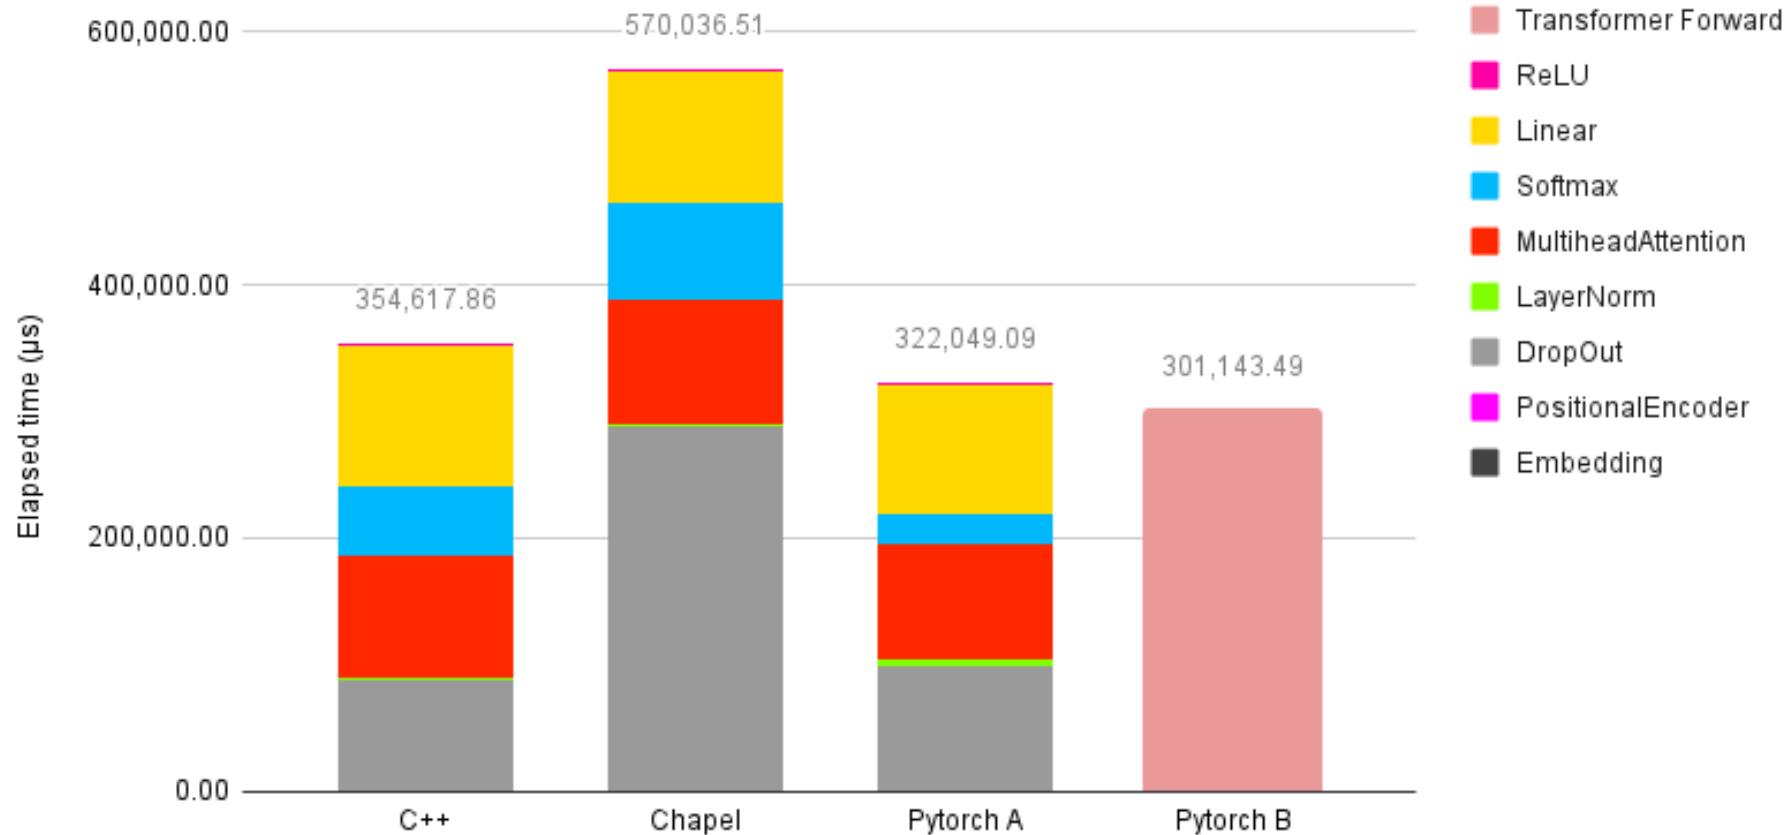
Small-Size Model on Single Thread



Description

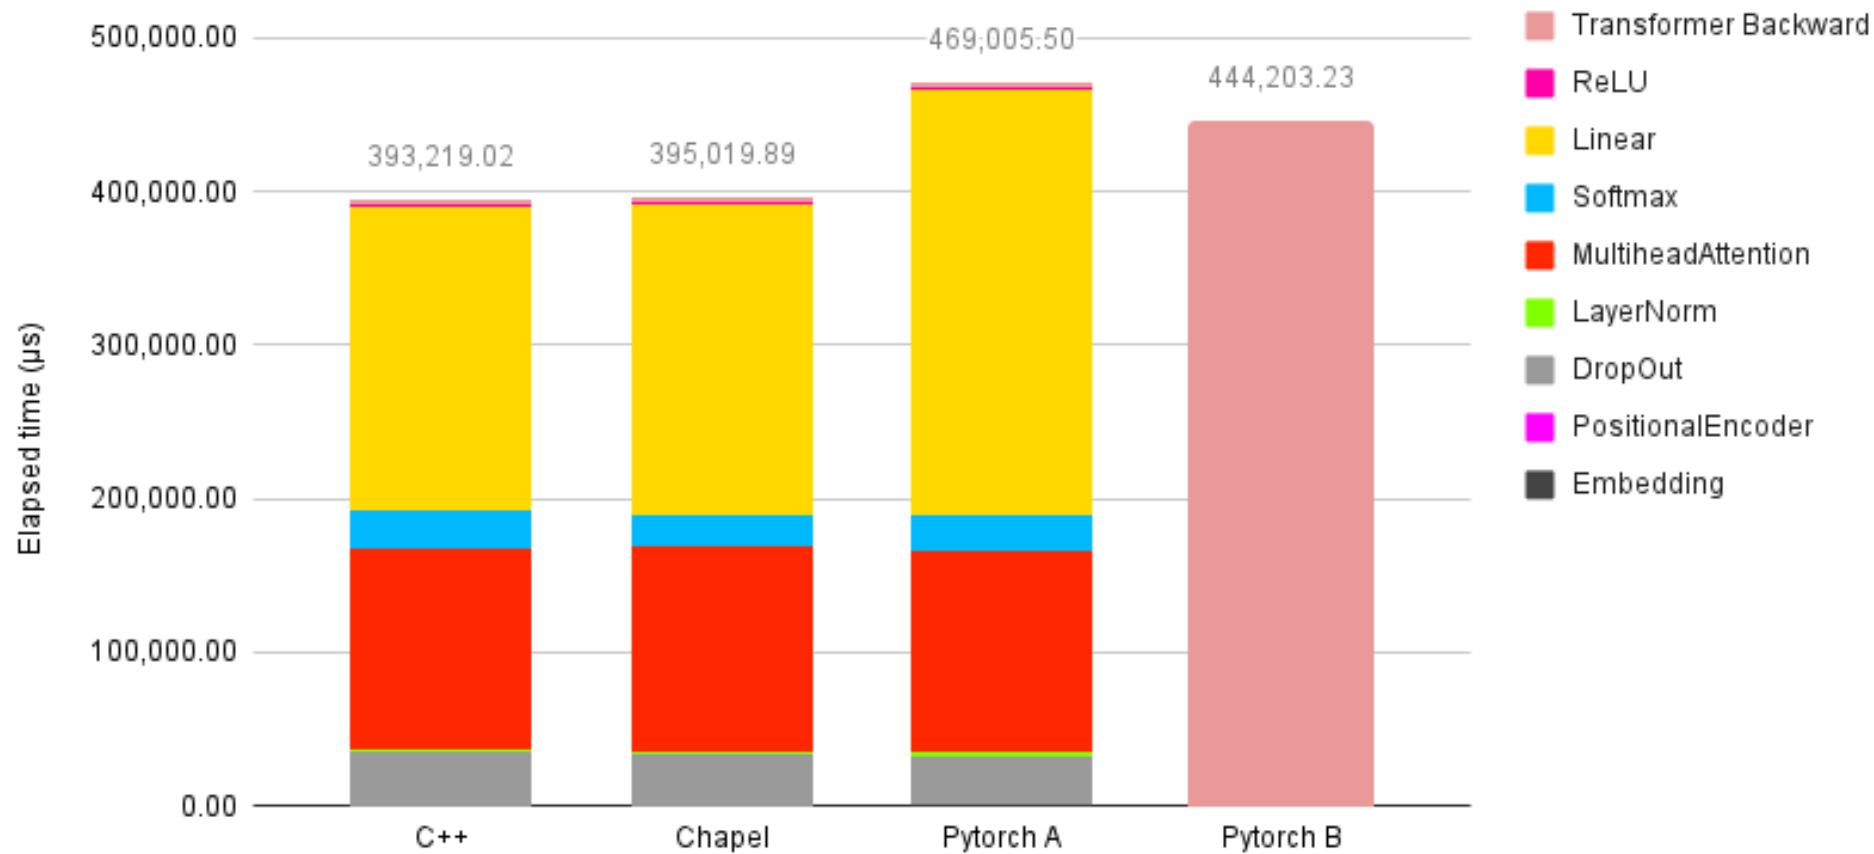
- Tested small model on Machine A
- The models were run for 500 iterations
- GitHub link for all single-thread code:
<https://github.com/markthitrin/Transformer/tree/SingleThread>
- Google Spreadsheet for detailed results:
<https://docs.google.com/spreadsheets/d/1aHkE9Ckl0-waxVwu-f4dlJ0peM6jIUQv3IU1-bFa0p0/edit?gid=2029252533#gid=2029252533>

Result of Forward Pass



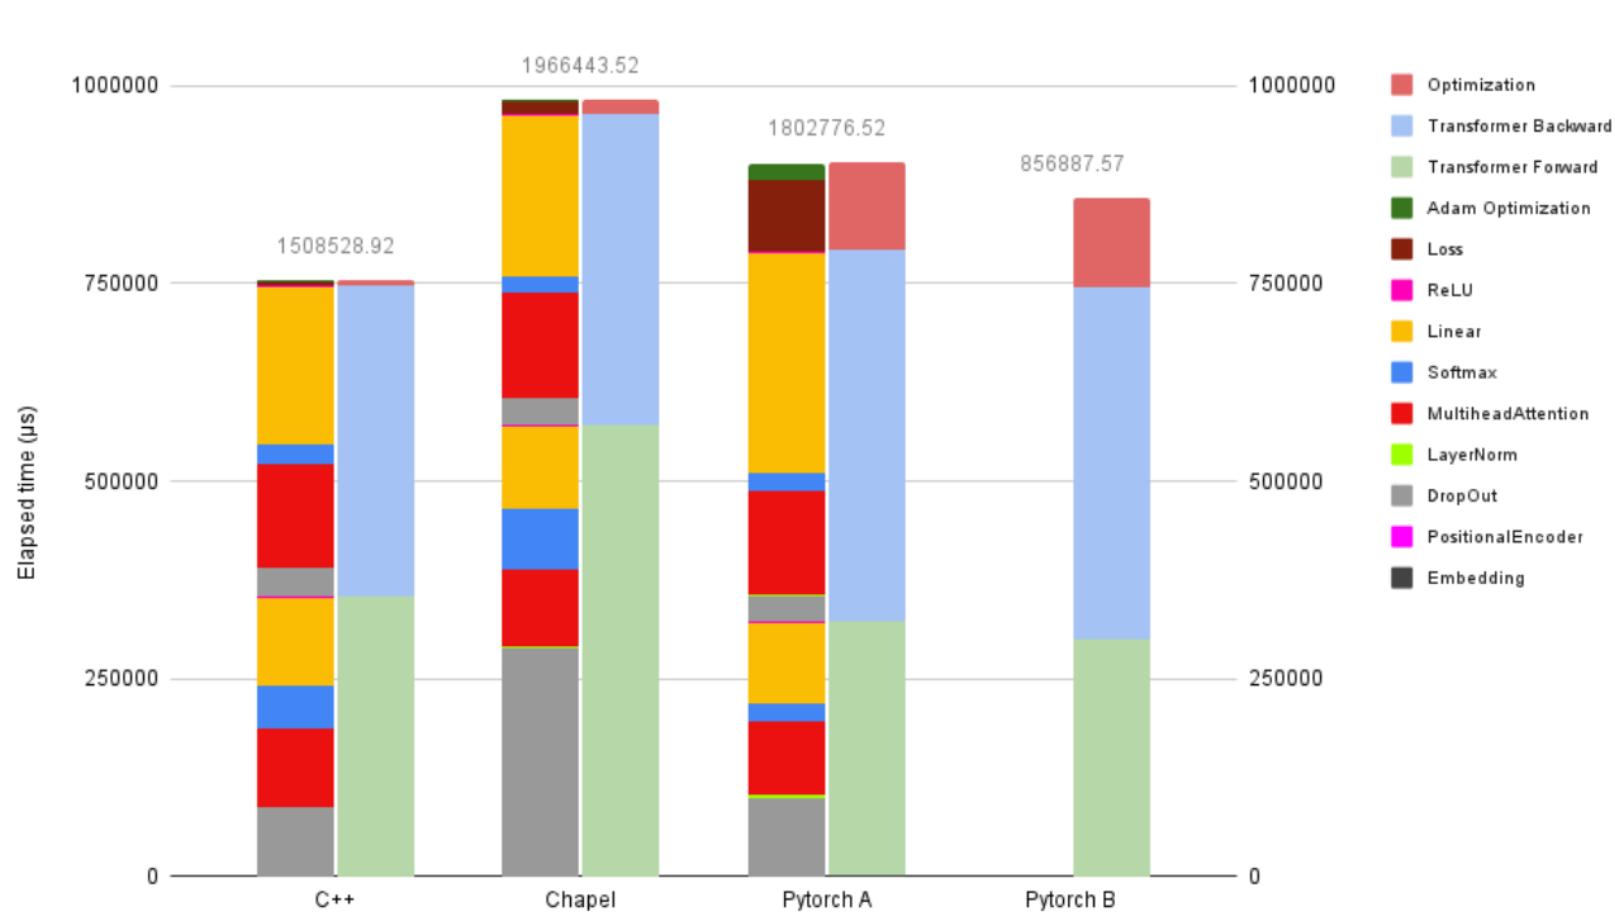
Time spent on each layer (in microseconds) during a single forward-pass training iteration for each model, tested on Machine A (single-threaded) using the small model configuration.

Result of Backward Pass



Time spent on each layer (in microseconds) during a single backward-pass training iteration for each model, tested on Machine A (single-threaded) using the small model configuration.

Overall Result



Time spent on each layer (in microseconds) per training iteration (including forward, backward, and update) for each model, tested on Machine A (single-threaded) using the small model configuration.

Matrix Representation

In C++, The TensorView is for capturing a portion of the Tensor.

```
class Tensor {  
    Tensor(int row, int column) {data = new float[row * column];}  
    ~Tensor() {delete[] data;}  
    float* data;  
};  
class TensorView {  
    TensorView(Tensor& t) {data = t.data;}  
    ~TensorView() {/*do nothing*/}  
    float* data;  
};
```

But in Chapel, ref is not allowed in a class or record

```
class TensorView {  
    // ref data; error  
}
```

Matrix Representation

- 1D array
- No multi-dimensional array
 - Slow when doing `for (a,b) in zip(A,B)`
 - No loop optimization (no loop unrolling, no vectorization)
 - Huge overhead from `advance_chpl`
 - Can be avoided with `for i in A.domain`
 - Mentioned in the [Chapel website](#) as performance concern
- Nest array, `var A: [0..#N][0..#N] real(32)`, is better but not best
 - Basically 1D array of 1D arrays
 - Non-continuous array

Matrix Multiplication

- Chapel can do better than C++ at some specific size of matrix, and worse at other size.
- Even though the compiler-generated code looks the same.

Matrix Operation

- Element-wise addition, multiplication, reduction, etc.
- Design matters a lot than expected

```
proc PlusReduce1(ref A: [?D] real(32), out output: real(32)) : void {
    output = 0.0;
    for i in D {
        output += A[i];
    }
}

proc PlusReduce2(D: domain(1), ref A: [] real(32)) : void {
    output = 0.0;
    for i in D {
        output += A[i];
    }
}
```

Matrix Operation

```
proc PlusReduce3(in start: int, in count: int, ref A: [] real(32), out output: real(32))
: void {
    output = 0.0;
    for i in start..#count {
        output += A[i];
    }
}

proc PlusReduce4(ref A: [?D] real(32), out output real(32)) : void {
    output = + reduce(A);
}
```

Matrix Operation

Design	Optimization
PlusReduce1	No
PlusReduce2	Unrolling
PlusReduce3	Unrolling + vectorize
PlusReduce4	No, create task

To prevent future problem, PlusReduce3's design is used

**This is hard to reproduce; it happens only on some specific code structures*

**I have tried operator overloading too. It gives the same performance as PlusReduce1.*

```
operator +=(ref sum: real(32), ref A: [] real(32)) {
    var output: real(32) = 0.0;
    for i in A.domain {
        output += A[i];
    }
    sum = output;
}
```

Softmax

- Slow compared to C++
- Chapel don't exponential vectorization (`_ZGVdN8v_expf_avx2`) while Clang enables vectorization using `-fveclib=libmvec`
- Chapel refuses to vectorize exponential function, even with:
 - Simple for loop iterating over the array's domain
 - Simple for loop iterating over the array's elements
 - Switching from `real(32)` to `real(64)`
 - Direct assignment `B = exp(A)`
 - Using foreach loops.
 - Passing the same flags used in Clang via `--ccflag`
 - Using `--no-ieee-float`

Process	Performance (μs)
Softmax Forward	C++: 53,759.49 Chapel: 75,521.40
Softmax Backward	C++: 25,531.68 Chapel: 20,907.40
Softmax Total	C++: 79,291.17 Chapel: 96,428.80

Dropout

- Use `randomStream.fill()`
 - Need `CHPL_RT_NUM_THREADS_PER_LOCALE=1` when do single thread experiment
 - Use integer random.
-

Process	Performance (μs)
Dropout Forward	C++: 87,045.49 Chapel: 288,983.16
Dropout Backward	C++: 35,605.02 Chapel: 34,410.46
Dropout Total	C++: 122,650.51 Chapel: 323,393.62

Multihead Attention

- The forward pass works fine
- The issue is in the weight gradient and the next layer's gradient computation during the backward pass

Transformer/Chapel/MultiheadAttention.chpl

```
147     for i in 0..#batch {  
148         MatMulPlusAB(dModel, sequenceLength, dModel, QTGradient[(i * block)..#block], inputQ[(i * block)..#block], WQOpt.gradient);  
149         MatMulPlusAB(dModel, sequenceLength, dModel, KTGradient[(i * block)..#block], inputK[(i * block)..#block], WKOpt.gradient);  
150         MatMulPlusAB(dModel, sequenceLength, dModel, VTGradient[(i * block)..#block], inputV[(i * block)..#block], WVOpt.gradient);  
151     }  
152     for i in 0..#batch {  
153         MatMulPlusATB(sequenceLength, dModel, dModel, QTGradient[(i * block)..#block], WQ, inputGradientQ[(i * block)..#block]);  
154         MatMulPlusATB(sequenceLength, dModel, dModel, KTGradient[(i * block)..#block], WK, inputGradientK[(i * block)..#block]);  
155         MatMulPlusATB(sequenceLength, dModel, dModel, VTGradient[(i * block)..#block], WV, inputGradientV[(i * block)..#block]);  
156     }  
157 }
```

The loop was heavily unrolled but no vectorization

Multihead Attention

- The Problem was fixed by changing `param` to `var` in `config.chpl`

Transformer/Chapel/Config.chpl

```
7 config var dModel: int = 512;
8 config var head: int = 8;
9 config var dFF: int = 2048;
10 config var dropoutRate: real(32) = 0.1;
11 config var N: int = 6;
12
```

This is quite a tricky solution

ReLU

- One problem is that the backward pass needs to be divided into two sections

Transformer/Chapel/ReLU.chpl (old)

```
22     for i in D {  
23         inputGradient[i] = if input[i] >= 0 then outputGradient[i] else 0.0:real(32);  
24     }
```

Transformer/Chapel/ReLU.chpl (new)

```
24     for i in 0..#(batch * sequenceLength * dFF) {  
25         outputGradient[i] = if input[i] >= 0 then outputGradient[i] else 0.0:real(32);  
26     }  
27     Copy(0,0,batch * sequenceLength * dFF,outputGradient,inputGradient);
```

This allows optimization to take place

ReLU

- Another mystery is that when tested on the small-size model, Chapel is slightly faster in the forward pass. But when tested on the full-size model, it becomes much slower

Process	Performance (μ s)
ReLU Forward (Small Size)	C++: 2,003.26 Chapel: 1,170.46
ReLU Forward (Full Size)	C++: 42,211.31 Chapel: 239,258.80

- The only difference I found in the compiler-generated code is that Chapel and C++ took different approaches

```
// Chapel
load mem -> res
max 0,res -> res
store res -> mem
```

```
// C++
max 0,mem -> res
store res -> mem
```

**Both version got same degree of vectorized and loop unrolling*

- This effect can also be seen in the backward pass of LayerNorm.

Other Layers

- Other layers are working fine
- The parameter updating (Adam optimization) in C++ and Chapel is much faster than in PyTorch

Full-Size Model on Single and Multiple Threads

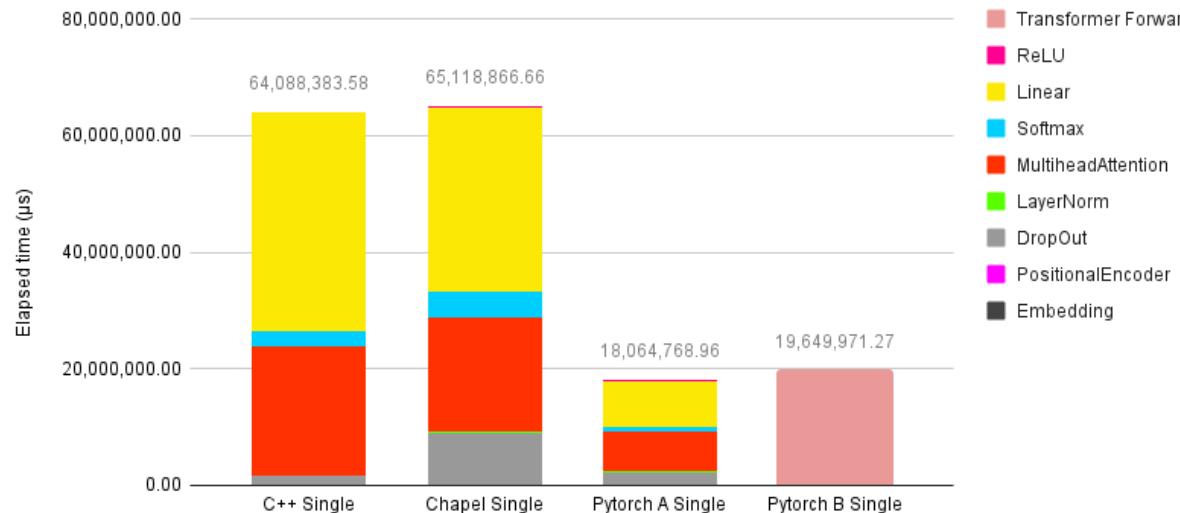


Description

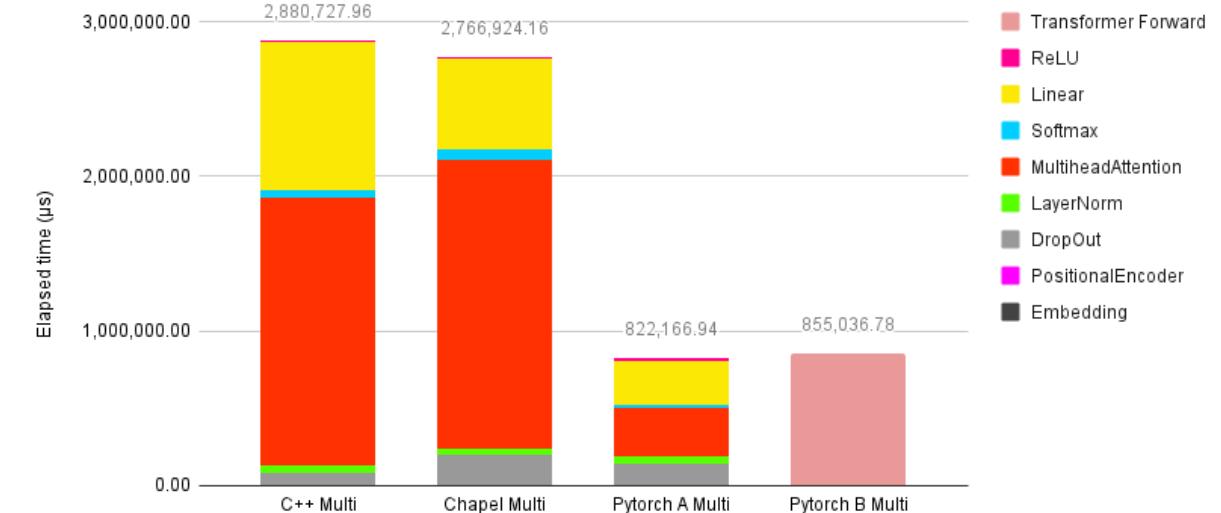
- Tested full-size model on the Machine B
- The models were run for 40 iterations
- C++ uses OpenMP
- Chapel uses `forall`, `coforall`, and custom iterators
- The degree of parallelism is estimated for each layer/operation
- The degree of parallelism is the same in both C++ and Chapel
- GitHub link for all single-thread code:
<https://github.com/markthitrin/Transformer/tree/MultiThread>
- Google Spreadsheet for detailed results:
<https://docs.google.com/spreadsheets/d/1aHkE9Ckl0-waxVwu-f4dlJ0peM6jIUQv3IU1-bFa0p0/edit?gid=2029252533#gid=2029252533>

Result of Forward Pass

Forward Pass Time Analysis



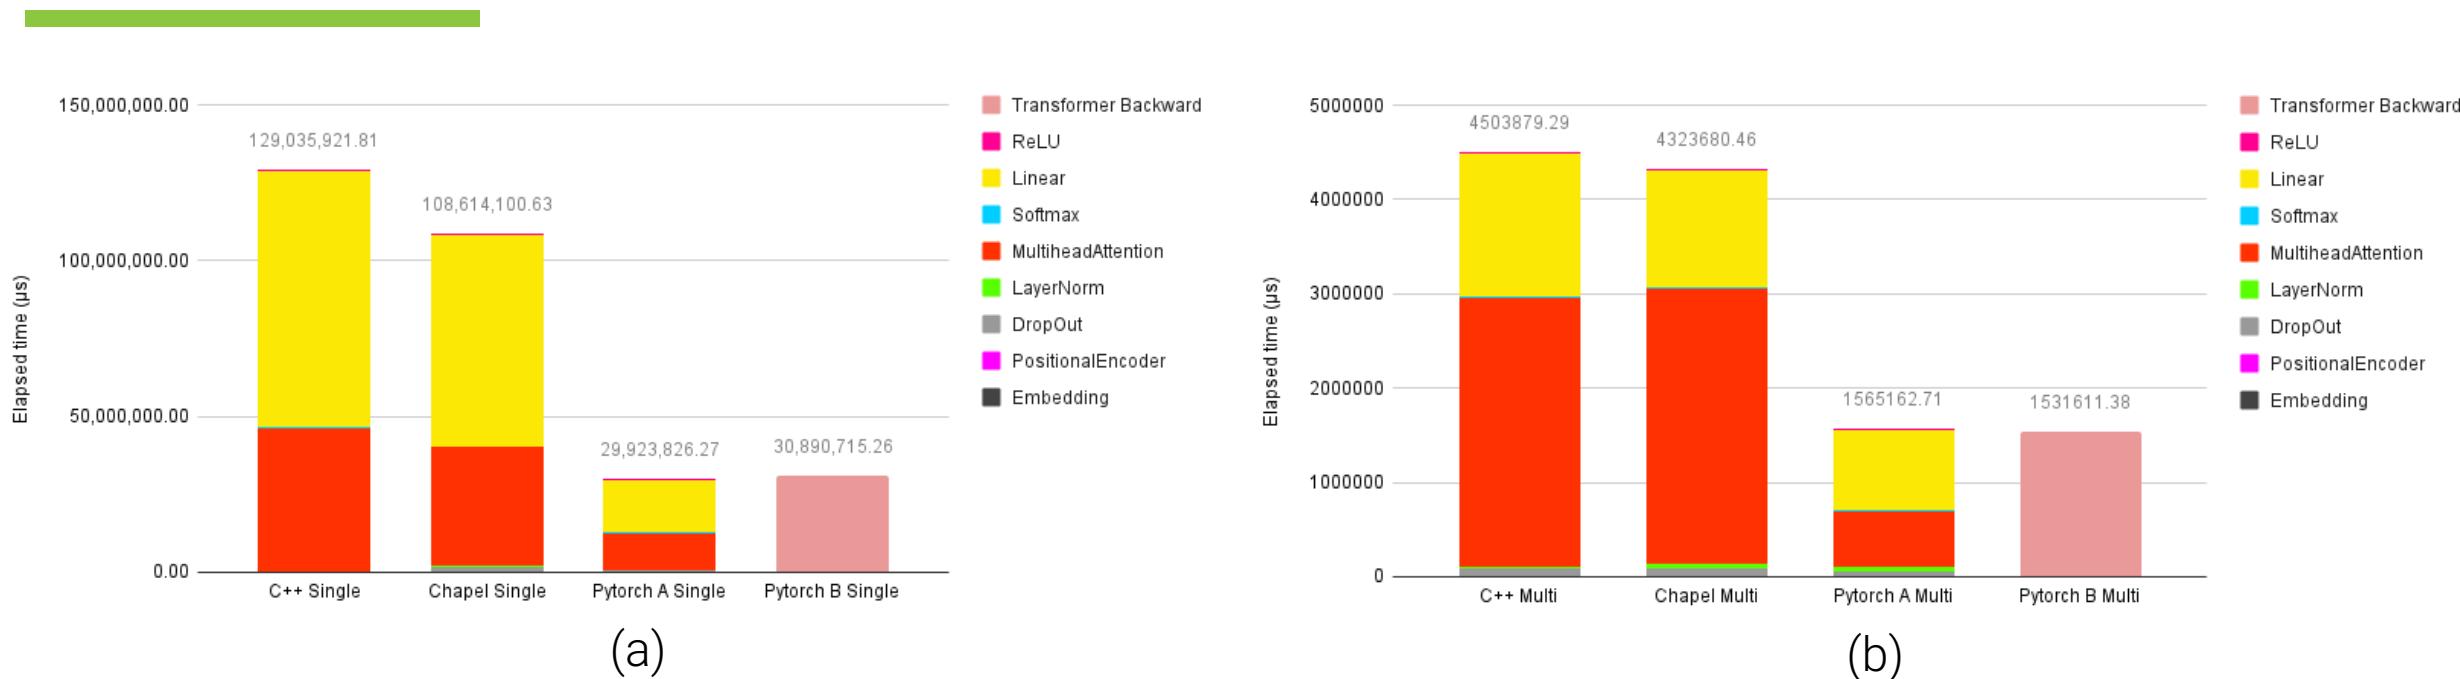
(a)



(b)

Time spent on each layer (in microseconds) during a single forward-pass training iteration for each model, measured on Machine B (single-threaded (a), multi-threaded (b)) using the full-size model configuration.

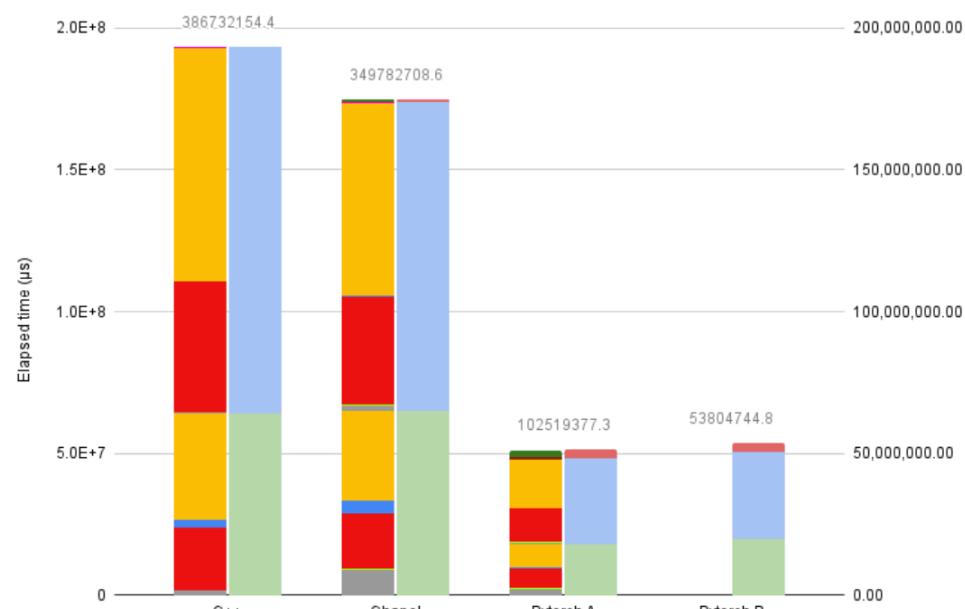
Result of Backward Pass



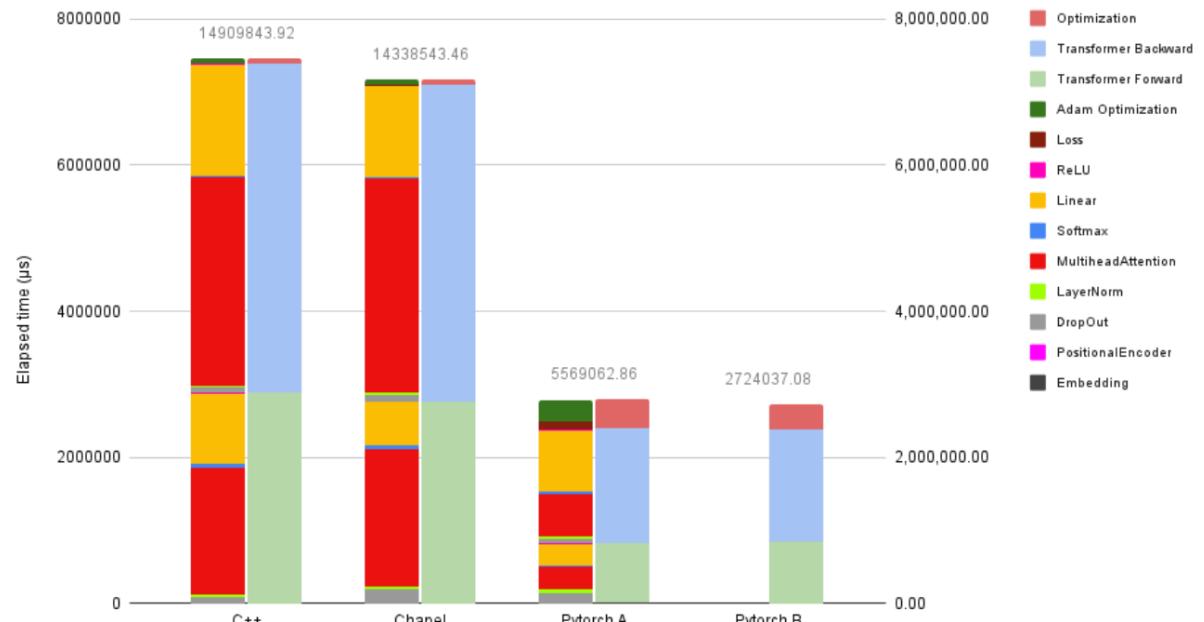
Time spent on each layer (in microseconds) during a single backward-pass training iteration for each model, measured on Machine B (single-threaded (a), multi-threaded (b)) using the full-size model configuration.

Overall Result

Performance Comparison



(a)



(b)

Time spent on each layer (in microseconds) per training iteration (including forward, backward, and update) for each model tested on Machine B (single-threaded (a), multi-threaded (b)) using the full-size model configuration.

Softmax

- A buffer is needed to store exponential values
- Unlike C++, Chapel can not perform stack allocation

```
// C++
#pragma omp parallel
for(int i = 0;i < row;i++) {
    float buffer[size];// stack memory
    // do something
}
```

The buffer is stored in stack memory

- This can be solved by moving the buffer declaration outside the loop

```
// Chapel
forall i in Par(start, end, numThread) {
    var buffer: [0..&N] real(32);
    // do something
}
```

The buffer is allocated and deallocated in every iteration

Other Layer

- Many layer perform as well as C++, even if they are slower in single-threaded. This is likely due to being bound by memory bandwidth
- Parameter Updates (Adam optimization) in C++ and Chapel are significantly faster than in PyTorch

Productivity



Productivity

- Things I like:
 - Easy to learn, similar to Python
 - Simple parallel programming through `for` loops
 - Requires type declaration of variables
 - Object memory management
 - Memory management across threads
 - Easier to run programs on multiple locales
- A Few Drawbacks I Noticed:
 - Long compilation time
 - All the performance issues that needed tricky solution I mentioned
 - Type casting between number types (e.g. `real(32)` from/to `real(64)`)
 - Generative AI support is limited

Conclusion



Conclusion

- This project compares four transformer models, implemented in C++, Chapel, Python
 - The achieve performance is reasonable
 - Chapel outperforms C++ in some parts
 - Performance issues were found, required tricky solutions
- Limitation in this project
 - No GPU and multi-locales
 - Not the most optimal code

For suggestions, advice, comments, or questions?, please contact me here:
thitrin.sastarasadhit@gmail.com

Thank you
