



**Hewlett Packard  
Enterprise**

# **Chapel's Batteries-Included Approach for Portable Parallel Programming**

---


Engin Kayraklioglu

Advances in Applied Computer Science Invited Speaker Series  
Los Alamos National Laboratory  
June 18, 2025

# Computational Science

## You start to learn programming...


- You have a processor (singular) and some memory
- You store your data in *the* memory
- You crunch your numbers in *the* processor
- Python is typically good enough



- You "just" learn some combination of
- C, C++, Fortran
  - MPI
  - OpenMP
  - CUDA, HIP, SYCL, Kokkos
  - SLURM

## One day, you need to use a supercomputer...

- You have millions of processors
  - Some of them are GPUs
- Your memory is now distributed
- Have to pay attention where your data is



**"Because there is no batteries-included programming abstraction to enable you to program both"**

**... or is there?**

# What is Chapel?

---

**Chapel:** A modern parallel programming language

- portable & scalable
- open-source & collaborative



## Goals:

- Support general parallel programming
- Make parallel programming at scale far more productive

[chapel-lang.org](http://chapel-lang.org)



# **What does Chapel code look like?**


---



# Programming with Chapel: Fundamentals

CPU Core    GPU Core

Memory



```
var A: [1..10] int;  
for elem in A do  
    elem += 1;
```


Local, non-distributed array allocation

Sequential iteration over the array

# Programming with Chapel: Basic Data Parallelism

CPU Core    GPU Core

Memory




```
var A: [1..10] int;
```

```
forall elem in A do ← Parallel iteration over the array  
    elem += 1;
```

# Programming with Chapel: Basic Data Parallelism

CPU Core    GPU Core

Memory



```
use BlockDist;  
var Arr = blockDist.createArray(1..10, int);  
  
forall elem in Arr do  
    elem += 1;
```


Block-distributed array allocation

Distributed, parallel iteration over  
the array

# Programming with Chapel: Basic Data Parallelism

CPU Core    GPU Core


Memory



```
var A: [1..10] int;  
  
forall elem in A do  
    elem += 1;
```

# Programming with Chapel: Locality as a First-Class Citizen

 CPU Core    GPU Core



## Predefined Variables for Locality

`here`

The current execution locale

`Locales`


An array storing all locales

`locale.gpus`

An array storing GPU sublocales

# Programming with Chapel: Locality as a First-Class Citizen

 CPU Core    GPU Core



## Predefined Variables for Locality

`here`

The current execution locale

`Locales`

An array storing all locales


`locale.gpus`

An array storing GPU sublocales

# Programming with Chapel: Basic Data Parallelism + Locality

CPU Core    GPU Core

Memory




```
var A: [1..10] int;  
  
forall elem in A do  
    elem += 1;
```

# Programming with Chapel: Basic Data Parallelism + Locality

CPU Core    GPU Core

Memory




```
on Locales[1] {  
    var A: [1..10] int;  
  
    forall elem in A do  
        elem += 1;  
}
```

The 'on' statement moves the execution to a remote locale

# Programming with Chapel: Basic Data Parallelism + Locality

CPU Core    GPU Core

Memory




```
on Locales[1].gpus[0] {
    var A: [1..10] int;
    forall elem in A do
        elem += 1;
}
```

Each locale object has a  
'gpus' array that store GPU sublocales

# Programming with Chapel: Basic Data Parallelism + Locality

CPU Core   GPU Core

Memory




```
on here.gpus[0] {
    var A: [1..10] int;
    forall elem in A do
        elem += 1;
}
```

'here' is a built-in representing  
the current execution locale

# Programming with Chapel: Data + Task Par. + Locality

CPU Core   GPU Core

Memory




```
coforall gpu in here.gpus {  
    on gpu {  
        var A: [1..10] int;  
  
        forall elem in A do  
            elem += 1;  
    }  
}
```

'coforall' loops run each iteration  
in a parallel task

# Programming with Chapel: Data + Task Par. + Locality

CPU Core   GPU Core

Memory



'do' can be used instead of curly braces  
for single-statement blocks  
(this is just a matter of style)


```
coforall gpu in here.gpus do on gpu {
    var A: [1..10] int;
}

forall elem in A do
    elem += 1;
}
```

# Programming with Chapel: Data + Task Par. + Locality

CPU Core   GPU Core

Memory




```
coforall loc in Locales do on loc {  
    coforall gpu in here.gpus do on gpu {  
        var A: [1..10] int;  
  
        forall elem in A do  
            elem += 1;  
    }  
}
```

A very similar 'coforall' + 'on' for multilocal parallelism

# Programming with Chapel: Data + Task Par. + Locality + Data Movement

CPU Core   GPU Core

Memory



```
var CpuArr: [1..10] int;  
coforall loc in Locales do on loc {  
    coforall gpu in here.gpus do on gpu {  
        const myChunk = start..end; //math omitted  
        var A = CpuArr[myChunk];  
  
        forall elem in A do  
            elem += 1;  
    }  
}
```

Arrays or slices can be copied across any locales,  
including GPU sublocale

# Programming with Chapel: Honorable Mentions

## 'begin' statement

- Starts an asynchronous task

## 'cobegin' statement

- Every statement in the block becomes a parallel task

## 'sync' statement

- Synchronizes all asynchronous tasks at the end of the block

## 'atomic' variables

- All operations on the variable is performed atomically, potentially across the network

## 'sync' variables

- Atomic variables with empty/full semantics: e.g. it can't be read twice in a row



## **How does Chapel perform?**

---



# HPCC Stream Triad and RA in C + MPI + OpenMP vs. Chapel

# **STREAM TRIAD: C + MPI + OPENMP**

```
use BlockDist;



config const n = 1_000_000,
          alpha = 0.01;
const Dom = blockDist.createDomain({1..n});
var A, B, C: [Dom] real;

B = 2.0;
C = 1.0;

A = B + alpha * C;
```

## **HPCC RA: MPI KERNEL**

```
...  
forall (_ , r) in zip(Updates, RASTream()) do  
    T[r & indexMask].xor(r);
```



# Bale IG in Chapel vs. SHMEM on HPE Cray EX (Slingshot-11)

## Chapel (Simple / Auto-Aggregated version)


```
forall (d, i) in zip(Dst, Inds) do  
    d = Src[i];
```

## SHMEM (Exstack version)

```
i=0;  
while( exstack_proceed(ex, (i==l_num_req)) ) {  
    i0 = i;  
    while(i < l_num_req) {  
        l_idx = pckindx[i] >> 16;  
        pe = pckindx[i] & 0xffff;  
        if(!exstack_push(ex, &l_idx, pe))  
            break;  
        i++;  
    }  
  
    exstack_exchange(ex);  
  
    while(exstack_pop(ex, &idx , &fromth)) {  
        idx = ltable[idx];  
        exstack_push(ex, &idx, fromth);  
    }  
    lgp_barrier();  
    exstack_exchange(ex);  
  
    for(j=i0; j<i; j++) {  
        fromth = pckindx[j] & 0xffff;  
        exstack_pop_thread(ex, &idx, (uint64_t)fromth);  
        tgt[j] = idx;  
    }  
    lgp_barrier();  
}
```



## SHMEM (Conveyors version)

```
i = 0;  
while (more = convey_advance(requests, (i == l_num_req)),  
      more | convey_advance(replies, !more)) {  
  
    for (; i < l_num_req; i++) {  
        pkg.idx = i;  
        pkg.val = pckindx[i] >> 16;  
        pe = pckindx[i] & 0xffff;  
        if (!convey_push(requests, &pkg, pe))  
            break;  
    }  
  
    while (convey_pull(requests, ptr, &from) == convey_OK) {  
        pkg.idx = ptr->idx;  
        pkg.val = ltable[ptr->val];  
        if (!convey_push(replies, &pkg, from)) {  
            convey_unpull(requests);  
            break;  
        }  
    }  
  
    while (convey_pull(replies, ptr, NULL) == convey_OK)  
        tgt[ptr->idx] = ptr->val;  
}
```



# Distributed Sorting Performance and Productivity Survey

- Using distributed sort as proxy to compare
  - Chapel, MPI, SHMEM and others
- All implementations are available
  - [github.com/mppf/distributed-lsb](https://github.com/mppf/distributed-lsb)



# Further Reading on Sorting Capabilities of Standard Libraries

- Chapel's standard library leverages parallelism out-of-the-box

Read more on Chapel blog

[chapel-lang.org/blog/posts/std-sort-performance/](https://chapel-lang.org/blog/posts/std-sort-performance/)



Chapel Language Blog

About Chapel Website Featured Series Tags Authors All Posts

## Comparing Standard Library Sorts: The Impact of Parallelism


Posted on January 30, 2024.

Tags: [Sorting](#) [Performance](#) [Language Comparison](#)



By: [Michael Ferguson](#)

Computing hardware is parallel. Everything from the Raspberry Pi to a supercomputer uses parallelism. The Chapel language and standard library make it easy to use that parallel hardware effectively.

The Chapel standard library `sort` routine is at least **10 times** faster than any other standard library `sort` measured in this benchmarking experiment. The reason: Chapel's standard library `sort` routine is parallel but the others are not. Chapel is designed for parallel computing and its standard library is built to leverage that parallelism.



# Other Comparisons



Diehl et al. "Benchmarking the Parallel 1D Heat Equation Solver in Chapel, Charm++, C++, HPX, Go, Julia, Python, Rust, Swift, and Java"







Fig. 2: miniBUDE results for small deck bm1

Milthorpe et al. "Performance Portability of the Chapel Language on Heterogeneous Architectures"



Slaughter et al. "Task Bench: A Parameterized Benchmark for Evaluating Parallel Runtime Performance"




Carneiro et al. "Investigating Portability in Chapel for Tree-Based Optimization on GPU-Powered Clusters"

**All great... but what  
about real-world usage?**

---




# Productivity Across Diverse Application Scales (code and system size)



**Computation:** Aircraft simulation / CFD  
**Code size:** 100,000+ lines  
**Systems:** Desktops, HPC systems



**Computation:** Coral reef image analysis  
**Code size:** ~300 lines  
**Systems:** Desktops, HPC systems w/ GPUs



**Computation:** Atmospheric data analysis  
**Code size:** 5000+ lines  
**Systems:** Desktops, sometimes w/ GPUs



## 7 Questions for Éric Laurendeau: Computing Aircraft Aerodynamics in Chapel

Posted on September 17, 2024.

Tags: Computational Fluid Dynamics, User Experiences, Interviews  
By: [Engin Kayraklıoglu](#), [Brad Chamberlain](#)

*"Chapel worked as intended: the code maintenance is very much reduced, and its readability is astonishing. This enables undergraduate students to contribute, something almost impossible to think of when using very complex software."*



## 7 Questions for Scott Bachman: Analyzing Coral Reefs with Chapel

Posted on October 1, 2024.

Tags: Earth Sciences, Image Analysis, GPU Programming, User Experiences, Interviews  
By: [Brad Chamberlain](#), [Engin Kayraklıoglu](#)

In this second installment of our [Seven Questions for Chapel Users](#) series, we're looking at a recent success story in which Scott Bachman used Chapel to unlock new scales of biodiversity analysis in coral reefs to study ocean health using satellite image processing. This is work that

*"With the coral reef program, I was able to speed it up by a factor of 10,000. Some of that was algorithmic, but Chapel had the features that allowed me to do it."*



## 7 Questions for Nelson Luís Dias: Atmospheric Turbulence in Chapel

Posted on October 15, 2024.

Tags: User Experiences, Interviews, Data Analysis, Computational Fluid Dynamics  
By: [Engin Kayraklıoglu](#), [Brad Chamberlain](#)

In this edition of our [Seven Questions for Chapel Users](#) series, we turn to Dr. Nelson Luis Dias from Brazil who is using Chapel to analyze data generated by the [Amazon Tall Tower Observatory \(ATTO\)](#), a project dedicated to long-term, 24/7 monitoring of greenhouse gas fluctuations. Read on

*"Chapel allows me to use the available CPU and GPU power efficiently without low-level programming of data synchronization, managing threads, etc."*

[read this interview series at: <https://chapel-lang.org/blog/series/7-questions-for-chapel-users/>]

# More on CHAMPS & CFD: Previous Talk at NASA Ames is Available

Check out

[www.nas.nasa.gov/pubs/ams/2025/02-20-25.html](http://www.nas.nasa.gov/pubs/ams/2025/02-20-25.html)



NASA Advanced Supercomputing (NAS) Division

HOME ABOUT ▾ TECHNICAL AREAS ▾ SUPERCOMPUTING AT NAS ▾ SOFTWARE ▾

[AMS Seminar Series](#)

**High-Performance, Productive Programming using Chapel with Examples from the CFD Solver CHAMPS**

Speakers: Engin Kayraklioglu, Hewlett Packard Enterprise  
Eric Laurendeau, Polytechnique Montreal  
M. Karim Mohamad Zayni, Ph.D. Student, Polytechnique Montreal  
February 20, 2025

**Presentation**

**Advanced Modeling & Simulation (AMS)  
Seminar Series**

The screenshot shows a slide titled "Advanced Modeling & Simulation (AMS) Seminar Series". The slide features a blue background with various 3D simulation visualizations, including flow fields and structural models. A video player interface at the bottom indicates a duration of 58:26. The NASA logo is visible in the bottom right corner.

[Seminar Slide Deck \(PDF-11.8MB\)](#)


# **Arkouda**

---




# What is Arkouda?

## **Q: “What is Arkouda?”**



# What is Arkouda?

Q: “What is Arkouda?”



A: “A scalable version of NumPy / Pandas for data scientists”

# Performance and Productivity: Sorting with Arkouda

## HPE Cray EX


- Slingshot-11 network (200 Gb/s)
- 8192 compute nodes
- 256 TiB of 8-byte values
- ~8500 GiB/s (~31 seconds)

## HPE Cray EX

- Slingshot-11 network (200 Gb/s)
- 896 compute nodes
- 28 TiB of 8-byte values
- ~1200 GiB/s (~24 seconds)

## HPE Apollo

- HDR-100 InfiniBand network (100 Gb/s)
- 576 compute nodes
- 72 TiB of 8-byte values
- ~480 GiB/s (~150 seconds)




**Implemented using ~100 lines of Chapel**



# Performance and Productivity: Telemetry Analysis with Arkouda


- ~500 GB of server telemetry data
  - Stored in Parquet files
    - Loaded in dataframes
  - Measured time includes:
    - IO
    - Histogram, mean, max, covariance

**Arkouda performs significantly better than Dask and Spark**



# What is Arkouda?

**Q:** “What is Arkouda?”



**A:** “A scalable version of NumPy / Pandas for data scientists”

**A’:** “An extensible framework for arbitrary HPC computations”

**A”:** “A way to drive HPC systems interactively from Python on a laptop”

# Arkouda Resources

**Website:** <https://arkouda-www.github.io/>

The Arkouda website homepage features a dark header with the Arkouda logo, navigation links for GitHub, documentation, and Gitter, and a search bar. The main content area includes a section titled "Arkouda is..." with three boxes: "Fast" (Arkouda is powered by Chapel, a programming language built from the ground up to support parallelism and distributed computing), "Interactive" (By distributing your data across multiple nodes, Arkouda allows you to rapidly transform and wrangle datasets in real time that are simply intractable for a laptop or desktop), and "Extensible" (One can expand on Arkouda's capabilities, thus enabling arbitrary scalable computations to be performed from Python). Below this is a "Powered by Chapel" section featuring the Chapel logo and a brief description of its implementation. A "Try it Out" button is located at the bottom left. A "Arkouda v2024.12.06 released!" section highlights the new release with a quote from Tess Hayes, Bytoa, and another from Jake Trockman, Erias.

**GitHub:** <https://github.com/Bears-R-Us/arkouda>

The Arkouda GitHub repository page shows the README and License tabs. The README contains a hand-drawn illustration of a bear with the text "arkoúða" and "massive scale data science". Below the illustration, there is a section titled "Arkouda (άρκούðα) Interactive Data Analytics at Supercomputing Scale" with links to CI status, documentation, and performance charts. The page also lists "Online Documentation" (link to Arkouda docs at GitHub Pages), "Nightly Arkouda Performance Charts" (link to Arkouda nightly performance charts), "Gitter channels" (links to Arkouda Gitter channel and Chapel Gitter channel), and "Talks on Arkouda" (links to Mike Merrill's SIAM PP-22 Talk and Arkouda Hack-a-thon videos).

# Arkouda Interview

**Blog:** Interview with founding co-developer, Bill Reus: <https://chapel-lang.org/blog/posts/7qs-reus/>

The screenshot shows a blog post on the Chapel Language Blog. The post title is "7 Questions for Bill Reus: Interactive Supercomputing with Chapel for Cybersecurity". It was posted on February 12, 2025, by Engin Kayraklioglu and Brad Chamberlain. The post discusses the 2025 edition of the Seven Questions for Chapel Users series, featuring an interview with Bill Reus. Bill is one of the co-creators of Arkouda, a Chapel's flagship application for interactive data analysis at massive scales. The post includes a table of contents with seven questions.

**Table of Contents**

1. Who are you?
2. What do you do? What problems are you trying to solve?
3. How does Chapel help you with these problems?
4. What initially drew you to Chapel?
5. What are your biggest successes that Chapel has helped achieve?
6. If you could improve Chapel with a finger snap, what would you do?
7. Anything else you'd like people to know?

*"I was on the verge of resigning myself to learning MPI when I first encountered Chapel. After writing my first Chapel program, I knew I had found something much more appealing."*

*"Chapel's separation of concerns immediately felt like the most natural way to think about large-scale computing. I would highly encourage anyone wanting to get into HPC programming to start with Chapel."*

## **How do I learn more about Chapel?**

---



# Ways to Engage with the Chapel Community

## “Live” Virtual Events

- [ChapelCon](#) (formerly CHIUW), annually
- [Project Meetings](#), weekly
- [Deep Dive / Demo Sessions](#), weekly timeslot

## Community / User Forums

- [Discord](#)
- [Discourse](#)  
chapel+qs@discoursemail.com
- Email Contact Alias
- [GitHub Issues](#)
- [Gitter](#)
- [Reddit](#)
- [Stack Overflow](#)



chapel+qs@discoursemail.com



## Electronic Communications

- [Chapel Blog](#), ~biweekly
- [Community Newsletter](#), quarterly
- [Announcement Emails](#), around big events

## Social Media

- [Bluesky](#)
- [Facebook](#)
- [LinkedIn](#)
- [Mastodon](#)
- [X / Twitter](#)
- [YouTube](#)



# Chapel has been accepted to HPSF

## Timeline:

- **May 2024:** HPSF launched at ISC
- **September 2024:** Began accepting applications for member projects
- **January 2025:** Chapel accepted to HPSF at the “established” project level
- **May 2025:** First-ever [HPSFcon](#)

## Resources:

- **Website:** <https://hpsf.io/>
- **Blog:** <https://hpsf.io/blog/>
- **YouTube channel:** <https://www.youtube.com/@HPSF-community>
- **GitHub org:** <https://github.com/hpsfoundation>



**HPSF**

HIGH PERFORMANCE  
SOFTWARE FOUNDATION



# Closing Remarks

---

## **Chapel allows programmers to leverage most common parallel hardware**

- Multicore, multinode, including cloud resources
- NVIDIA and AMD GPUs are supported with vendor-neutral code

## **Same set of programming abstractions are used to achieve this portability**

- No need to add things on, Chapel comes batteries-included
- No need to paradigm-shift when going from a single node to scaling on a supercomputer

## **Chapel is being used in many different fields, and in a wide range of institutions**

- Some application fields are CFD, data analytics, graph processing, ecological research, astrophysics
- Have been used by academia, industry, and government
- From desktops to supercomputers



# Thank you

---

