



Hewlett Packard
Enterprise

The Secret Sauce of Vendor-Neutral GPU Programming (in Chapel)

Jade Abraham

May 7th, 2025

What is the Secret Sauce?

- What does it take to implement a programming language for performant and portable GPU code?
 - Modern programming language
 - Not another C/C++ library
 - First-class parallel programming features
 - A compiler that can target multiple GPU vendors
 - A portable runtime
- Does something exist today that fills this gap?
 - Yes!



What is Chapel?

Chapel: A modern parallel programming language

- portable & scalable
- open-source & collaborative



Goals:

- Support general parallel programming
- Make parallel programming at scale far more productive


chapel-lang.org



First-Class Parallel Programming – By Example

CPU Core GPU Core

Memory



```
on Locales[0].gpus[0]  
var B: [1..10] int;
```

Local CPU array allocation

```
on Locales[0].gpus[1]  
var C: [1..10] int;
```

Local GPU array
allocation

```
forall elem in A do  
    elem += 1;
```

Compute on all CPUs in
parallel

```
on Locales[0].gpus[0] do  
forall elem in B do
```

Launch a kernel on a single GPU

```
    elem += 1;
```


```
on Locales[0].gpus[1] do  
    C += 1;
```

Launch a kernel on a single GPU
(implicitly parallel)

Hello, GPUs!

CPU Core GPU Core

Memory




```
for gpu in Locales[0].gpus do on gpu {  
    var B: [1..10] int;  
    B += 1;  
}  
  
coforall gpu in Locales[0].gpus do on gpu {  
    var B: [1..10] int;  
    B += 1;  
}
```

Execute on all GPUs in Serial

Execute on all GPUs in Parallel

Portable LLVM-based Compiler



Extensible Runtime Architecture


Interface for:

- Compiler-injected calls
 - e.g. kernel prep and launch
- Extern calls from modules
 - e.g. memory management, data movement

Interacts with the rest of the runtime to:

- Maintain task-private data
 - e.g. GPU streams
- Make host-based allocations
- Move data across locales
- Trigger diagnostics


- Thin layer for primitive GPU operations
 - e.g. call a kernel, initialize driver, query info
- Wraps around drivers



Bonus! Hello, Distributed GPUs!

CPU Core GPU Core

Memory



```
coforall loc in Locales do on loc {  
    coforall gpu in loc.gpus do on gpu {  
        var B: [1..10] int;  
        B += 1;  
    }  
}
```

More about Chapel + GPUs

- How Does Chapel's GPU Support Work?
 - A more in-depth look at Chapel's GPU internals
 - <https://www.youtube.com/watch?v=J0av4VJbS4o>
- Chapel Runtime Overview
 - How the rest of Chapel's runtime handles threading, remote communication, memory management, and more
 - <https://www.youtube.com/watch?v=rC4Oz654bsU>
- The Game of Life: A multi-GPU implementation in Chapel
 - A larger example of programming GPUs in Chapel
 - This video is part of a GPU series with other coding examples
 - <https://www.youtube.com/watch?v=U96mA84Klqo>



Ways to Engage with the Chapel Community

Live/Virtual Events

- ChapelCon (formerly CHIUW), annually
- Chapel project meeting, weekly

Community / User Forums

- Discord
- Discourse
chapel+qs@discoursemail.com
- Email Contact Alias
- GitHub Issues
- Gitter
- Reddit
- Stack Overflow



Electronic Broadcasts

- Chapel Blog, ~biweekly
- Community Newsletter, quarterly
- Announcement Emails, around big events

Social Media

- Bluesky
- Facebook
- LinkedIn
- Mastodon
- X / Twitter
- YouTube



Thank you

<https://chapel-lang.org>
@ChapelLanguage

